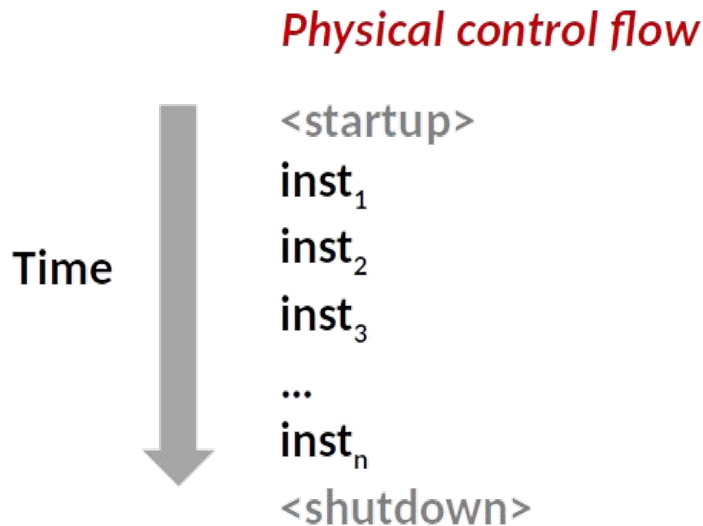


02. Exceptional Control Flow

CS 4352 Operating Systems

Control Flow

- Processors do only one thing:
 - From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time
 - This sequence is the CPU's control flow (or flow of control)



Altering the Control Flow

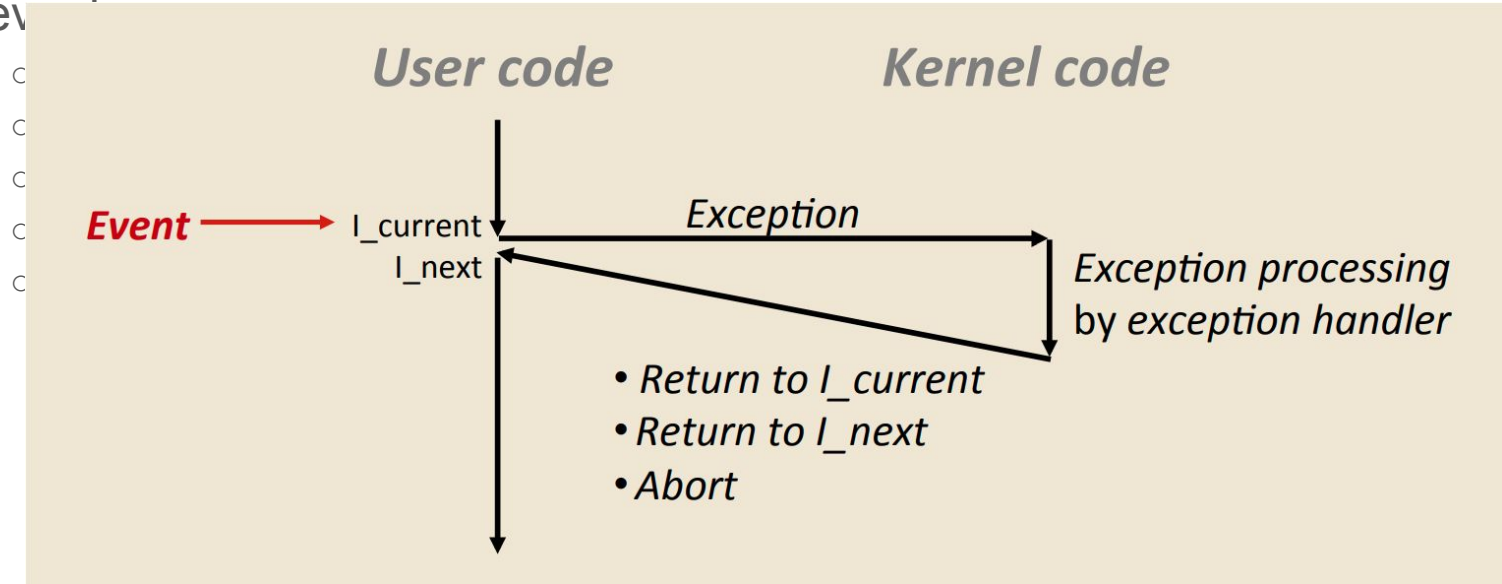
- Two mechanisms for changing control flow in **program state**:
 - Jumps and branches
 - Call and return
- Insufficient for an OS (difficult to react to changes in **system state**):
 - Data arrives from a disk or a network adapter
 - Instruction divides by zero
 - User hits Ctrl-C at the keyboard
 - System timer expires
- OS needs mechanisms for “exceptional control flow”

Exceptional Control Flow

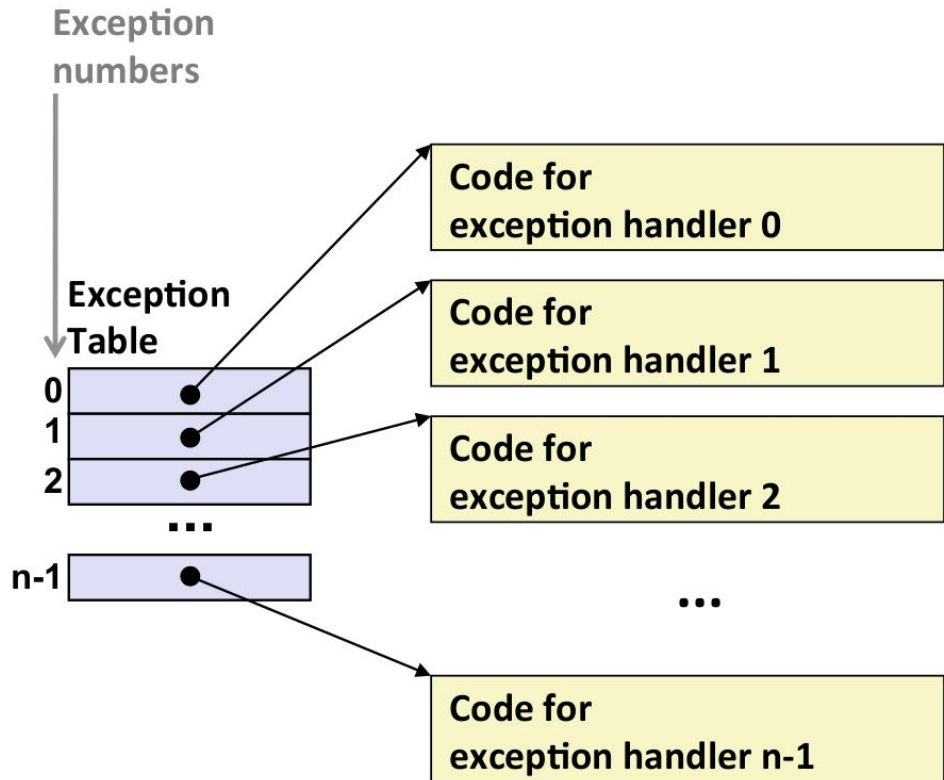
- Exists at all levels of a computer system
- Low level mechanisms
 - 1. Exceptions
 - Change in control flow in response to a system event (i.e., change in system state)
 - Implemented using combination of hardware and OS
- Higher level mechanisms
 - 2. Process context switch
 - Implemented by OS and hardware timer
 - 3. Signals
 - Implemented by OS
 - 4. Nonlocal jumps: `setjmp()` and `longjmp()`
 - Implemented by C runtime library

Exceptions

- An exception is a transfer of control to the OS kernel in response to some event.



Exception Tables

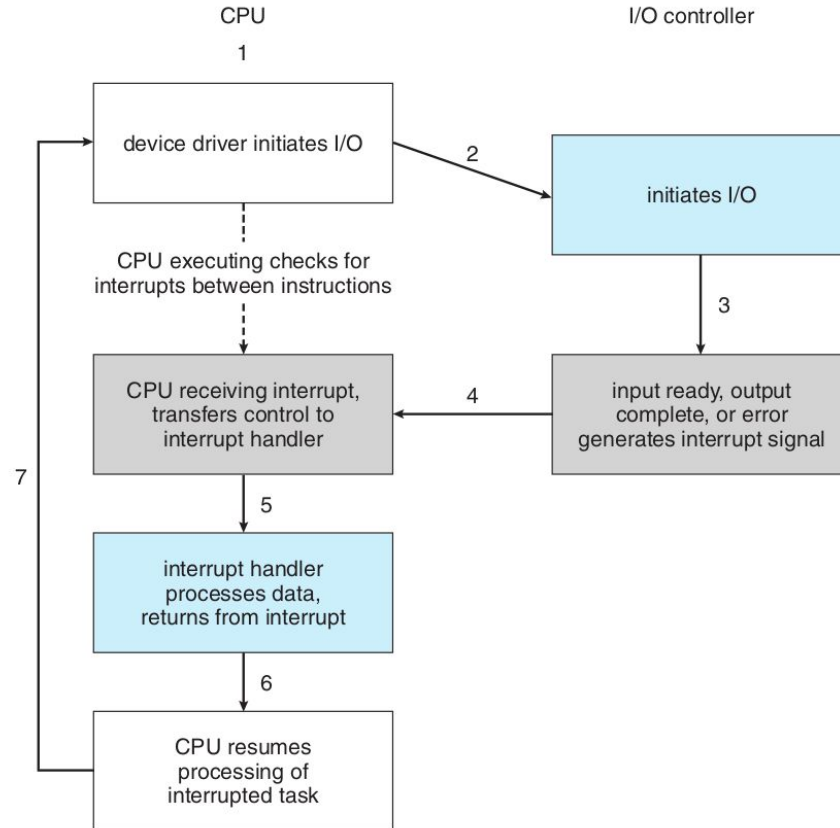


- On x86, they are called interrupt vector tables
 - Although not all are interrupts
- Each type of event has a unique exception number k
- k = index into exception table (a.k.a. interrupt vector)
- Exception handler k is called each time exception k occurs

Asynchronous Exceptions (Interrupts)

- Caused by events **external** to the processor
 - Indicated by setting the processor's **interrupt pin**
 - Handler returns to “next” instruction
- Examples:
 - Timer interrupt
 - Every few ms, an external timer chip triggers an interrupt
 - Used by the kernel to take back control from user programs
 - I/O interrupt from external device
 - Hitting Ctrl-C at the keyboard
 - Arrival of a packet from a network
 - Arrival of data from a disk

Interrupt-Driven I/O Cycle



Synchronous Exceptions

- Caused by events that occur **as a result of executing an instruction** (also called software-generated interrupts):
 - Traps
 - Intentional
 - Examples: system calls, breakpoint traps, special instructions
 - Returns control to “next” instruction
 - Faults
 - Unintentional but possibly recoverable
 - Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
 - Either re-executes faulting (“current”) instruction or aborts
 - Aborts
 - Unintentional and unrecoverable
 - Examples: illegal instruction, parity error, machine check
 - Aborts current program

System Calls

- System

- What



- Each sy

- The

<i>Number</i>	<i>Name</i>	<i>Description</i>
0	read	Read file
1	write	Write file
2	open	Open file
3	close	Close file
4	stat	Get info about file
57	fork	Create process
59	execve	Execute a program
60	_exit	Terminate process
62	kill	Send signal to process

umbers

Invoking a System Call

- A user program runs in user mode, and a system call runs in kernel mode
 - When making a system call in user mode, a special **trap** instruction is executed to switch from user mode to kernel mode
 - The trap handler examines the system call number and dispatches to the corresponding system call handler
- Some system calls have one or more parameters
 - Parameters can be passed in registers
 - Parameters can be stored in a block, and the address of the block is passed in a register
 - Parameters can be passed through the stack
 - Block and stack methods do not limit the number or length of parameters being passed

x86 Example

```
movl $20, %eax # Get PID of current process
int $0x80 # Invoke system call!
# Now %eax holds the PID of the current process
```

- X86 (32 bit and 64 bit) has several methods for system call
 - **int \$0x80** is the classical way to make the system call
 - 0x80 of the system's exception table gives the trap handler
 - Parameters for Linux system call are passed using registers
 - %eax is for system call number, and %ebx, %ecx, %edx, %esi, %edi, %ebp are used for passing 6 parameters to the system call
 - After the system call, register %rax contains the result (return value)
- There is a faster way to make 32-bit system calls: using the **sysenter**
 - Parameter passing is the same as for int \$0x80
- On x86-64, we normally use the **syscall** instruction for fast system calls
 - System call number → %rax
 - Return value → %rax
 - 6 parameters → %rdi, %rsi, %rdx, %r10, %r8, %r9

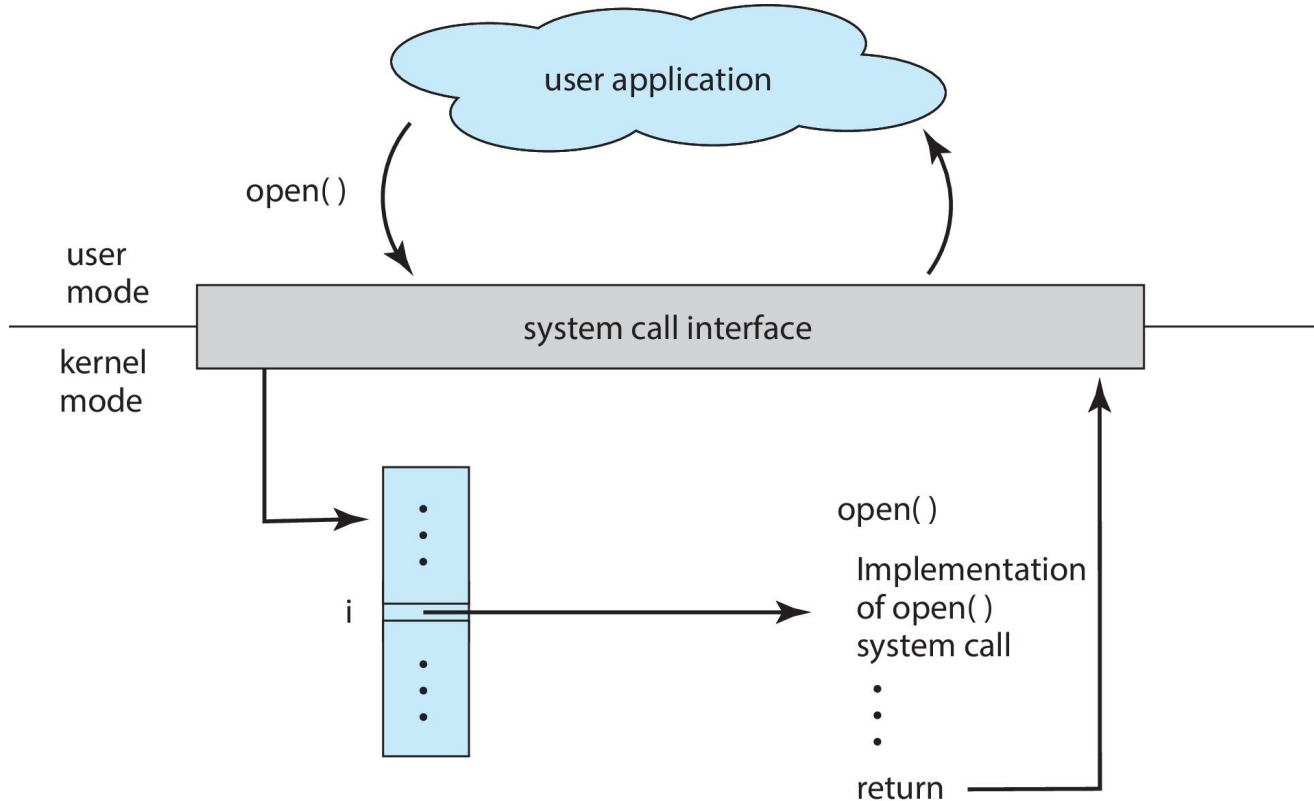
Different Architectures

Arch/ABI	Instruction	System call #	Ret val	Ret val2	Error
alpha	callsys	v0	v0	a4	a3
arc	trap0	r8	r0	-	-
arm/OABI	swi NR	-	r0	-	-
arm/EABI	swi 0x0	r7	r0	r1	-
arm64	svc #0	w8	x0	x1	-
blackfin	excpt 0x0	P0	R0	-	-
i386	int \$0x80	eax	eax	edx	-
ia64	break 0x100000	r15	r8	r9	r10
m68k	trap #0	d0	d0	-	-
microblaze	brki r14,8	r12	r3	-	-
mips	syscall	v0	v0	v1	a3
nios2	trap	r2	r2	-	r7
parisc	ble 0x100(%sr2, %r0)	r20	r28	-	-
powerpc	sc	r0	r3	-	r0
powerpc64	sc	r0	r3	-	cr0.S0
riscv	ecall	a7	a0	a1	-
s390	svc 0	r1	r2	r3	-
s390x	svc 0	r1	r2	r3	-
superh	trap #0x17	r3	r0	r1	-
sparc/32	t 0x10	g1	o0	o1	psr/csr
sparc/64	t 0x6d	g1	o0	o1	psr/csr
tile	swint1	R10	R00	-	R01
x86-64	syscall	rax	rax	rdx	-
x32	syscall	rax	rax	rdx	-
xtensa	syscall	a2	a2	-	-

Wrappers

- The program invokes a wrapper function in the C library
 - If you've ever noticed files like "libc.so", that's the C library
 - But, why?
- The wrapper function puts any arguments to the system call in the registers **expected by the OS kernel**
- The wrapper function executes a trap instruction to invoke the system call (as mentioned above -- system call number is in %rax)
- The wrapper function checks if the service returned an error, and if so, sets a global variable named ***errno*** with this error value
- The wrapper function then returns to the caller and provides an integer return value to indicate success or failure

API – System Call – OS Relationship



System Call Example: Opening File

- User calls: **open(filename, options)**
- Calls **__open()** function, which invokes system call instruction **syscall**
 - %rax contains syscall number -- 0x2
 - Return value in %rax
 - Negative value is an error corresponding to negative errno

```
0000000000e5d70 <__open>:
```

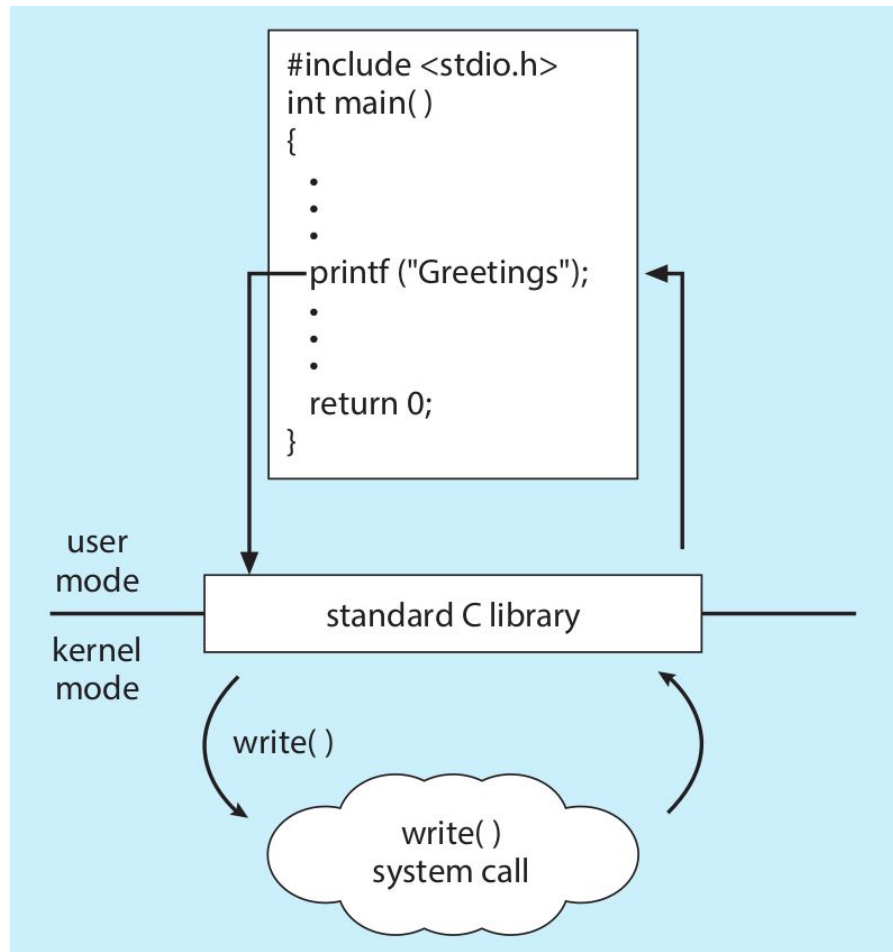
```
...  
e5d79:  b8 02 00 00 00      mov  $0x2,%eax  # open is syscall #2  
e5d7e:  0f 05               syscall          # Return value in %rax  
e5d80:  48 3d 01 f0 ff ff    cmp  $0xffffffffffffffff001,%rax  
...  
e5dfa:  c3                 retq
```


Tracing System Calls

- The strace command lets you see the system calls that are made by a process
 - Example: type “strace ls” at a terminal to see all the system calls that the ls (which lists files in the current directory) makes
 - You might see calls like:
 - execve -- this loads a new program and starts running it
 - open -- open a file
 - read -- read from a file
 - close -- close a file
 - fstat -- get information about a file
- See details by typing “man strace”

Standard C Library Example

- C program invoking printf() library call, which calls write() system call

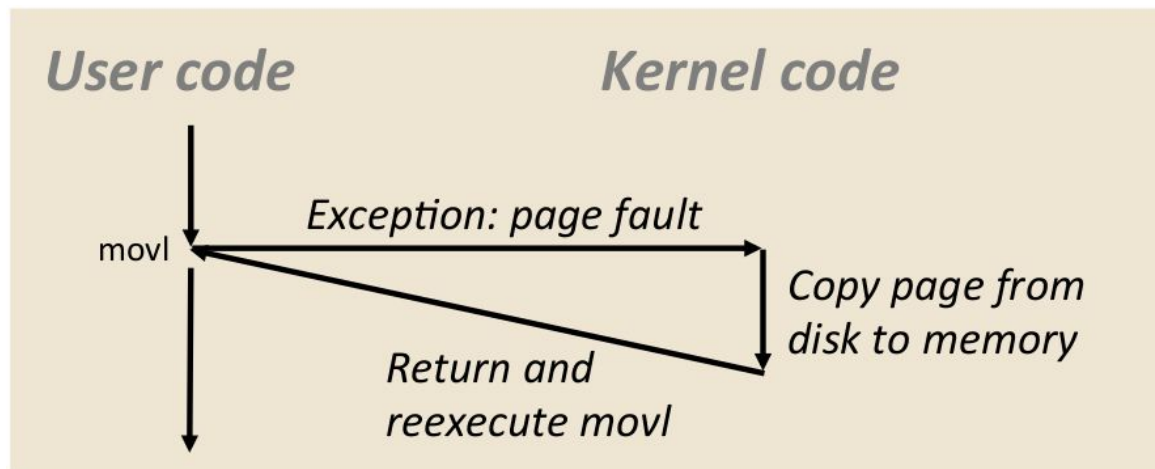


Fault Example: Page Fault

- User writes to memory location
- That portion (page) of user's memory is currently on disk


```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```

```
80483b7:  c7 05 10 9d 04 08 0d  movl    $0xd,0x8049d10
```



Fault Example: Invalid Memory Reference

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```



```
80483b7: c7 05 60 e3 04 08 0d movl $0xd,0x804e360
```


- Sends **SIGSEGV** signal to user process
- User process exits with “segmentation fault”

Homework

- Read Chapter 3
- Mini-project 1 is due on Sep. 11th

Next Lecture

- We will start learning process management



Credit: <http://xkcdsw.com/873>