

# 03. Processes

CS 4352 Operating Systems

# Process Concept


- A process is an instance of a running program
  - One of the most profound ideas in computer science
  - Sometimes also called a job or a task
  - Program is **passive** entity stored on disk (executable file); process is **active**
  - One program can be several processes
    - Consider multiple users executing the same program
  - Current state including the contents in CPU registers and memory
- Process provides each program with two key abstractions:
  - Logical control flow
    - Each program seems to have exclusive use of the CPU
    - Provided by kernel mechanism called **context switching**
  - Private address space
    - Each program seems to have exclusive use of main memory.
    - Provided by kernel mechanism called **virtual memory**

# Process in Memory

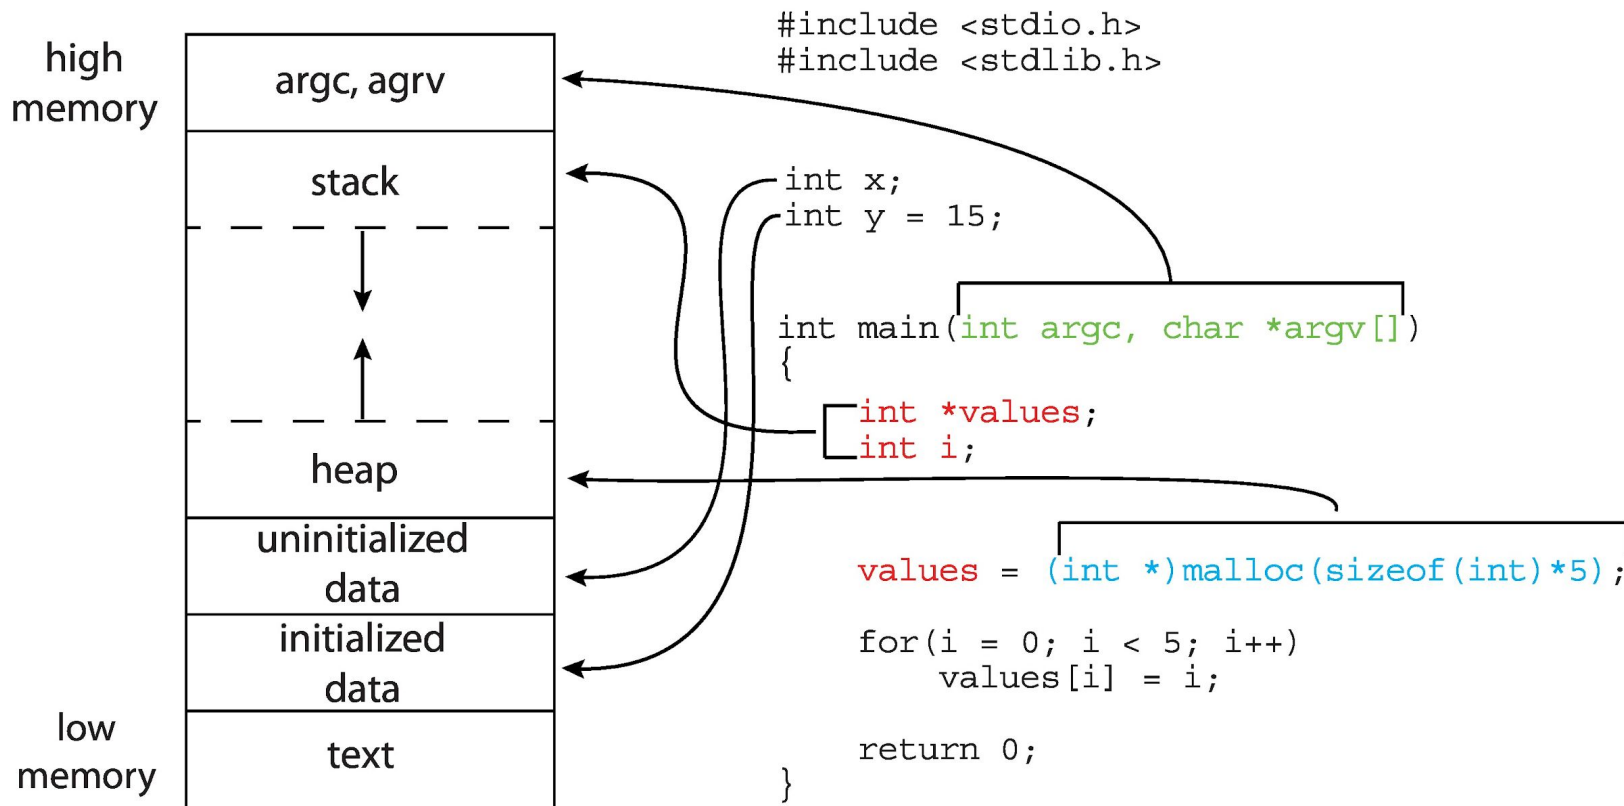
- The program code, also called **text** section
  - **Program counter** contains the next instruction to execute
- **Data** section containing global variables
  - There is also a BSS (Block Started by Symbol) for uninitialized global variables
  - Better Save Space → don't occupy space in the object files
- **Stack** containing temporary data
  - Function parameters, return addresses, local variables
- Heap containing memory dynamically allocated during run time

max

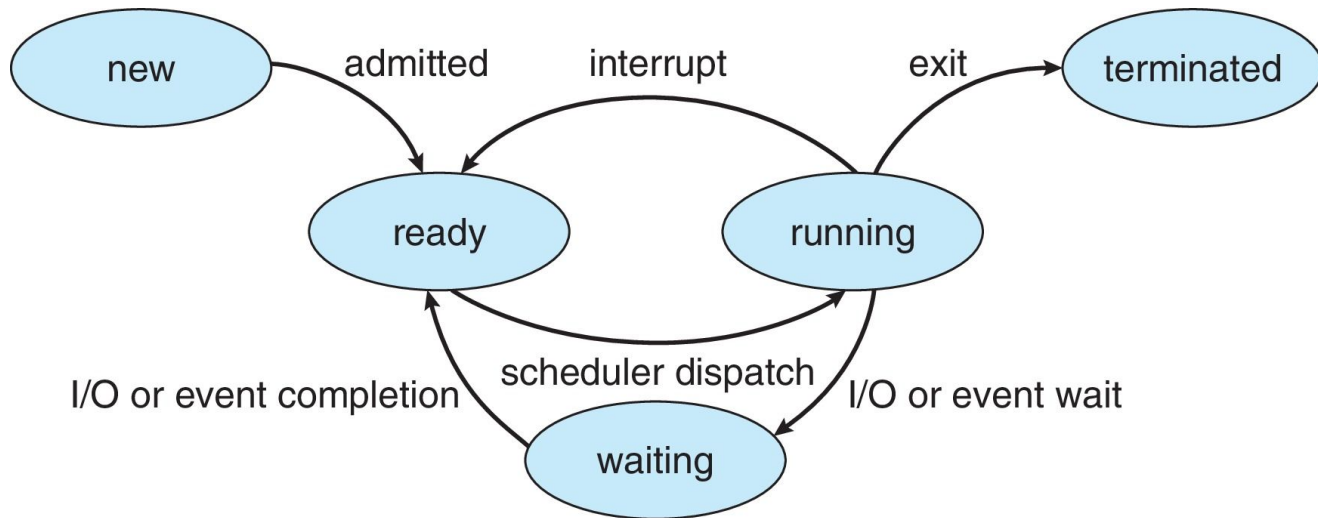
0



# Memory Layout of a C Program



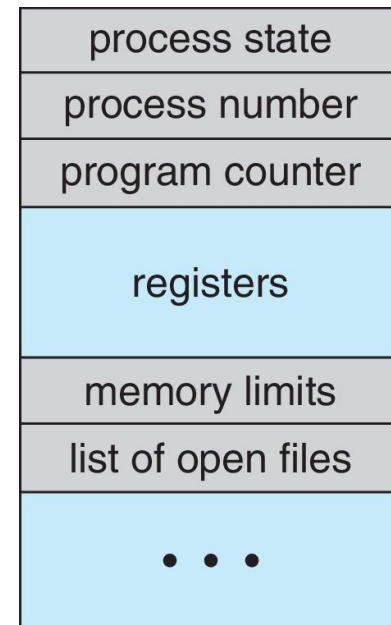
# Process State



- As a process executes, it changes state
  - New: The process is being created
  - Running: Instructions are being executed
  - Waiting: The process is waiting for some event to occur
  - Ready: The process is waiting to be assigned to a processor
  - Terminated: The process has finished execution

# Process Control Block (PCB)

- Information associated with each process(also called task control block)
  - Process state – running, waiting, etc.
  - Program counter – location of instruction to next execute
  - CPU registers – contents of all process-centric registers
  - CPU scheduling information – priorities, scheduling queue pointers
  - Memory-management information – memory allocated to the process
  - Accounting information – CPU used, elapsed clock time, time limits
  - I/O status information – I/O devices allocated to process, list of open files



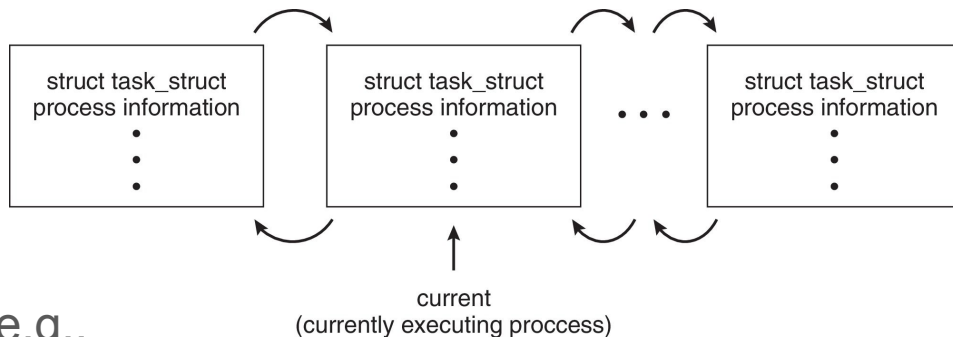
# PCB in Linux

- The **struct task\_struct**


- Defined in include/linux/sched.h

- It has a large number of members, e.g.,

- pid t\_pid; /\* process identifier \*/
- long state; /\* state of the process \*/
- unsigned int time\_slice; /\* scheduling information \*/
- struct task\_struct \*parent; /\* this process's parent \*/
- struct list\_head children; /\* this process's children \*/
- struct files\_struct \*files; /\* list of open files \*/
- struct mm\_struct \*mm; /\* address space of this process \*/



# The Illusion



- Computer runs many processes simultaneously
  - Applications for one or more users
    - Web browsers, email clients, editors, ...
  - Background tasks
    - Monitoring network & I/O devices

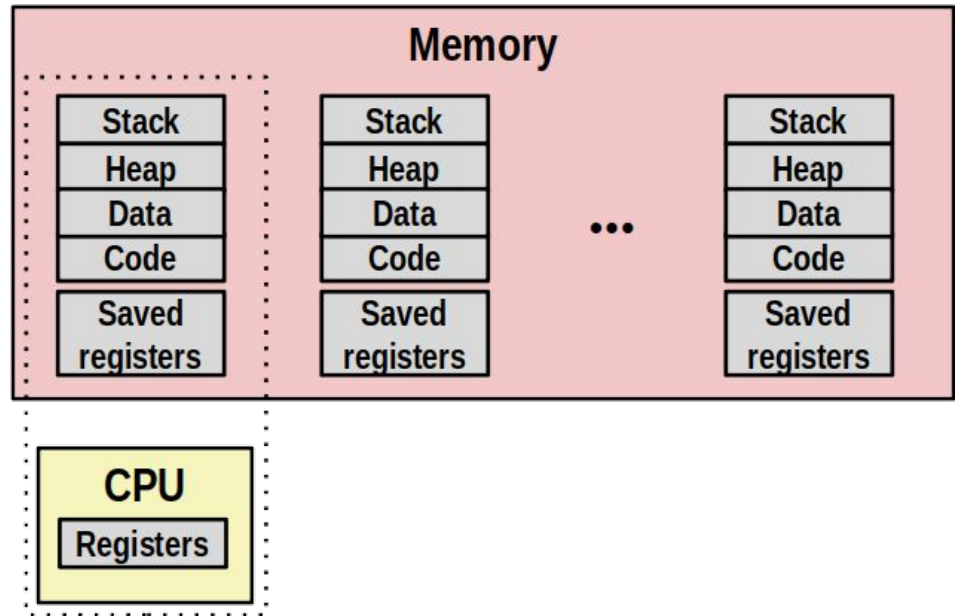


# Top Command

```
zzk@isis: ~  
File Edit View Search Terminal Help  
top - 18:43:28 up 10:36, 1 user, load average: 0.24, 0.38, 0.22  
Tasks: 343 total, 1 running, 268 sleeping, 0 stopped, 0 zombie  
%Cpu(s): 0.3 us, 0.1 sy, 0.0 ni, 99.5 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st  
KiB Mem : 32690432 total, 25633280 free, 2641388 used, 4415764 buff/cache  
KiB Swap: 97654784 total, 97654784 free, 0 used. 28789092 avail Mem
```

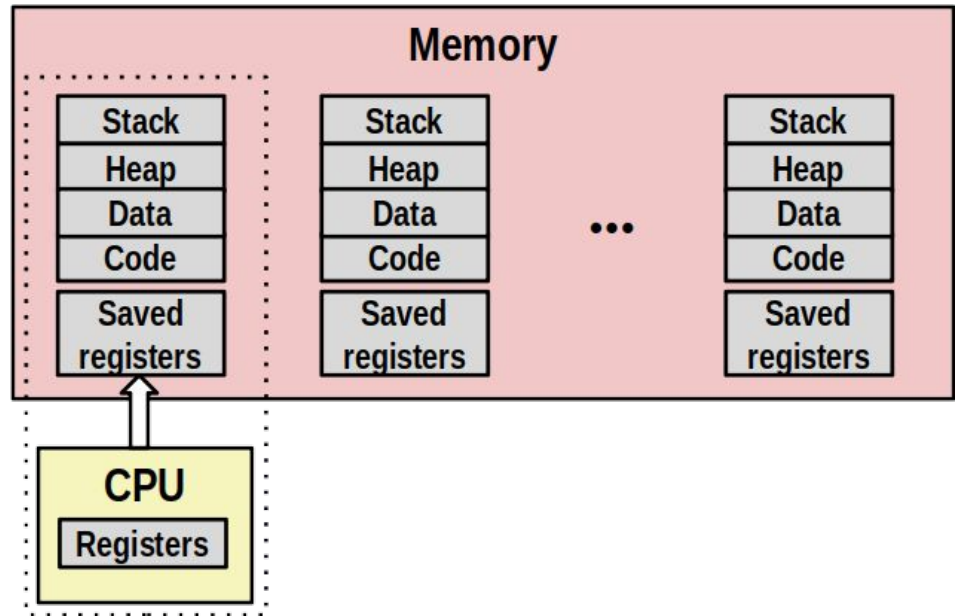
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
14655	zzk	20	0	5059244	196092	116664	S	1.7	0.6	7:39.56	chrome
11717	zzk	20	0	1225520	390404	201136	S	1.0	1.2	8:44.65	chrome
4121	zzk	20	0	4103040	226904	119512	S	0.7	0.7	4:57.40	gnome-shell
4355	zzk	20	0	1539264	112964	69736	S	0.7	0.3	2:54.14	fdm
27678	zzk	20	0	43820	3884	3108	R	0.7	0.0	0:00.25	top
1776	debian-+	20	0	109612	53084	15236	S	0.3	0.2	0:28.22	tor
2505	root	20	0	1394484	11272	9692	S	0.3	0.0	0:36.08	teamviewerd
11758	zzk	20	0	586388	110956	67548	S	0.3	0.3	3:37.56	chrome
26858	root	20	0	0	0	0	I	0.3	0.0	0:00.34	kworker/5:0
27060	zzk	20	0	9171384	252144	96644	S	0.3	0.8	1:04.02	chrome
27108	root	20	0	0	0	0	I	0.3	0.0	0:00.18	kworker/4:2
1	root	20	0	225816	9612	6756	S	0.0	0.0	0:08.71	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.04	kthreadd
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/0:+
6	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	mm_percpu_+
7	root	20	0	0	0	0	S	0.0	0.0	0:00.09	ksoftirqd/0
8	root	20	0	0	0	0	I	0.0	0.0	0:11.88	rcu_sched

# The (Traditional) Reality



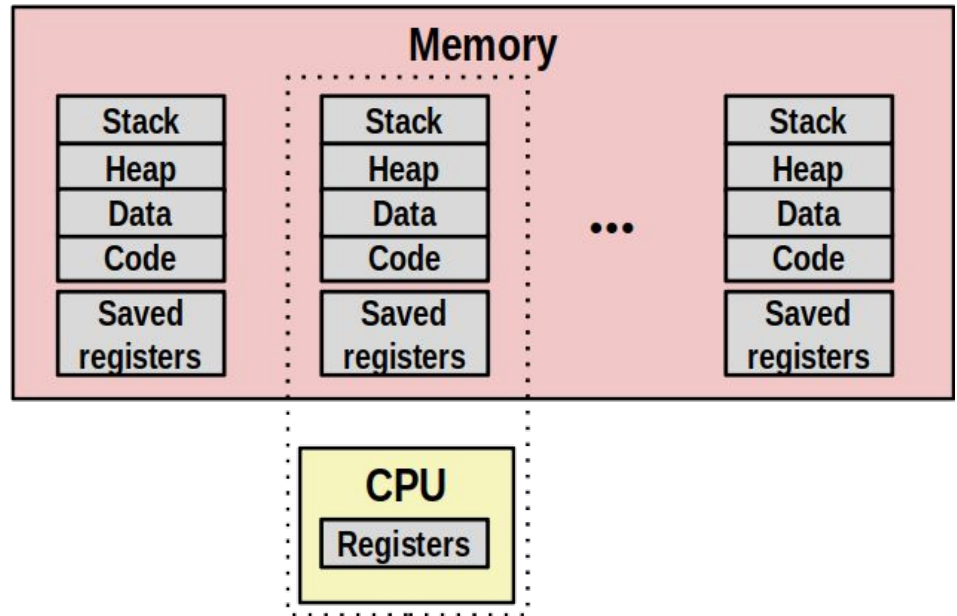
- Single processor executes multiple processes concurrently
  - Process executions interleaved (multitasking)
  - Address spaces managed by virtual memory system (later in course)
  - Register values for non-executing processes saved in memory

# The (Traditional) Reality



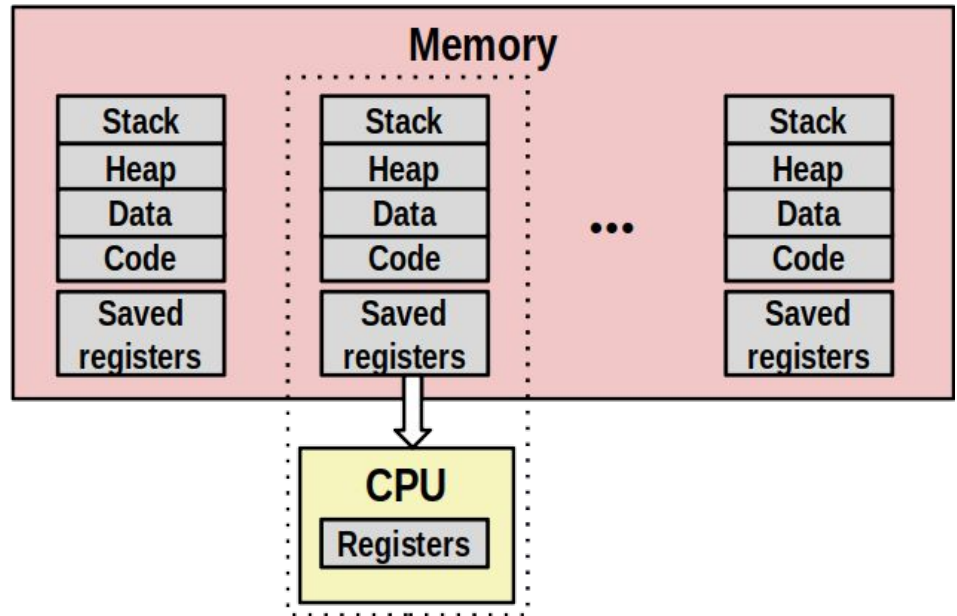
- Save current registers in memory

# The (Traditional) Reality



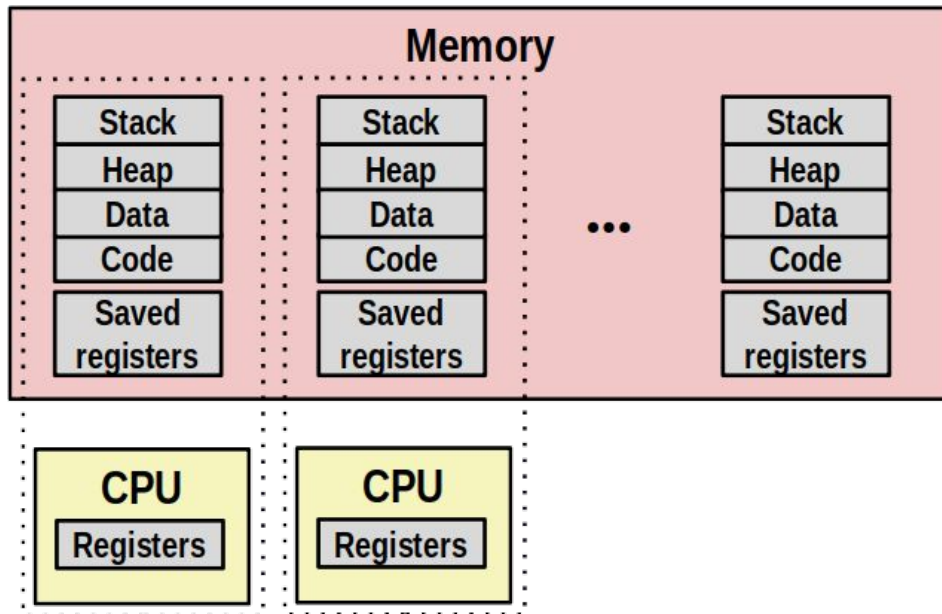
- Schedule next process for execution

# The (Traditional) Reality



- Load saved registers and switch address space (context switch)


# The (Modern) Reality



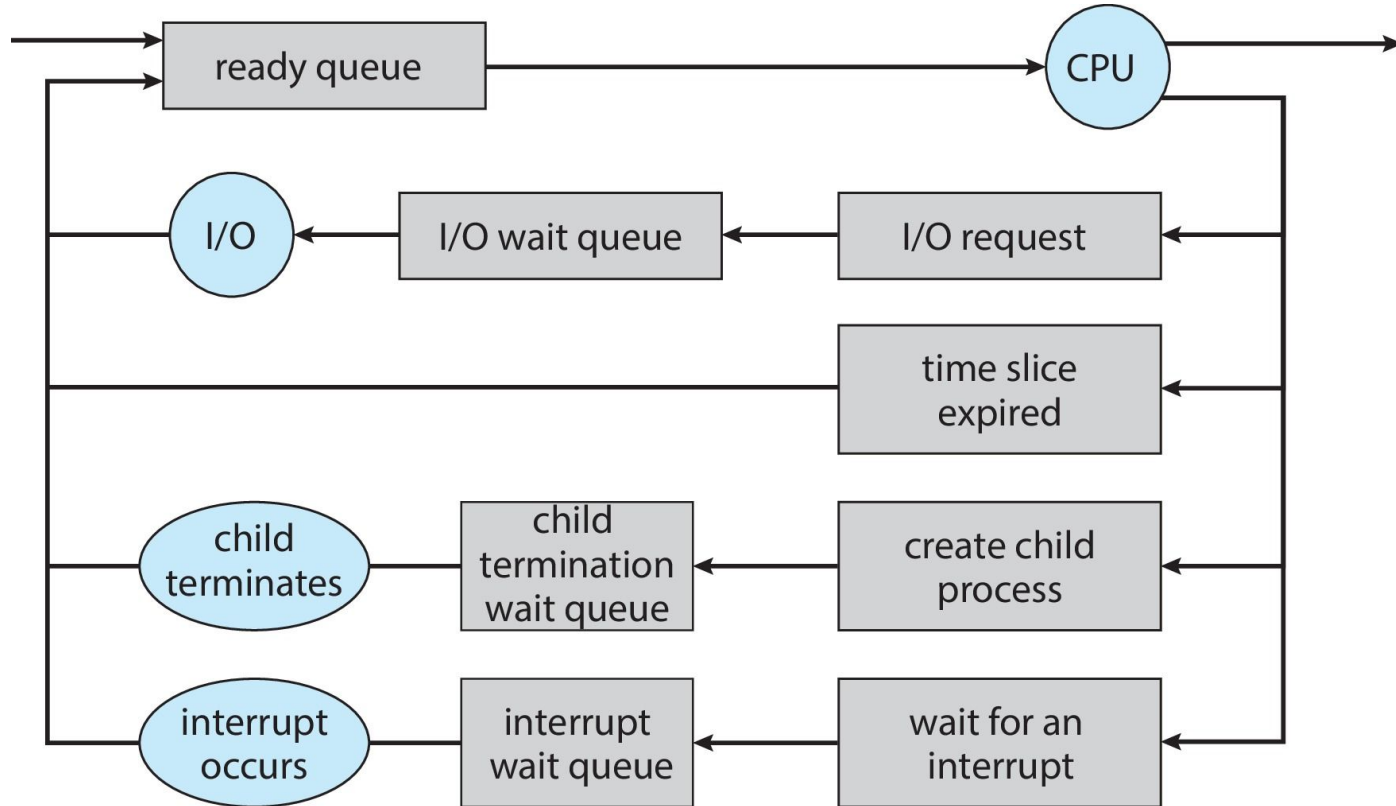
- Multicore processors
  - Multiple CPUs on single chip
  - Share main memory (and some of the caches)
  - Each can execute a separate process
    - Scheduling of processors onto cores done by kernel

# Scheduler

- **Scheduler** selects among available processes for execution on CPU cores
  - Goal -- Maximize CPU use, quickly switch processes onto CPU cores
- Maintains scheduling queues of processes
  - **Ready queue** – set of all processes ready and waiting to execute
  - **Waiting queues** – set of processes waiting for an event (i.e., I/O)
  - Processes migrate among the various queues



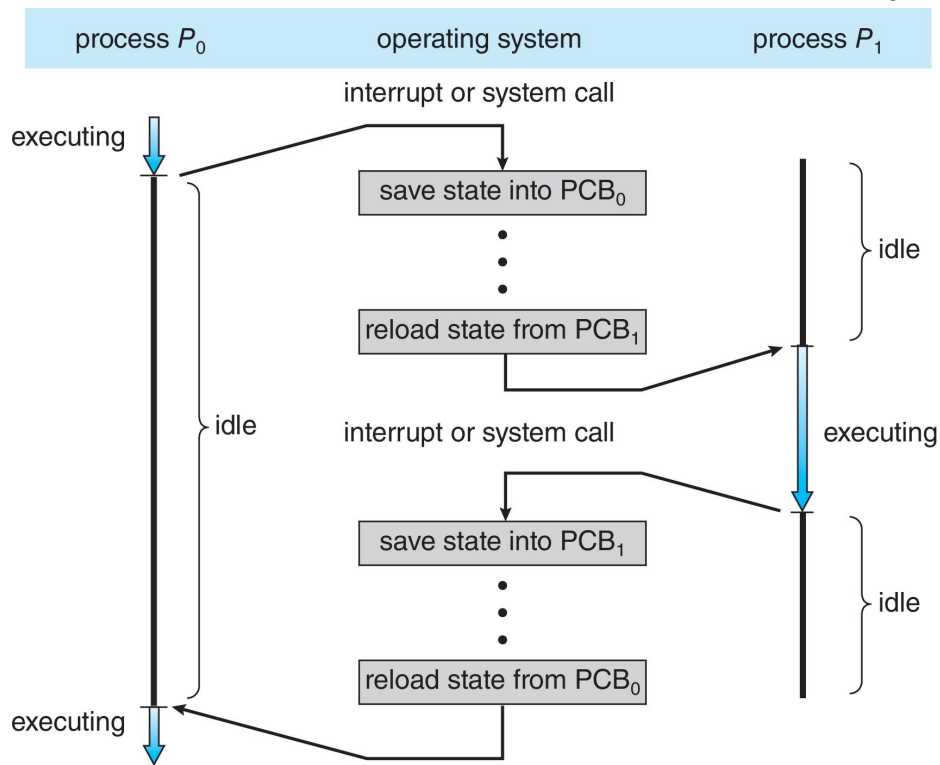
# Representation of Process Scheduling





# CPU Switch From Process to Process

A **context switch** occurs when the CPU switches from one process to another




# Context Switch Details

- Very machine dependent. Typical things include:
  - Save program counter and integer registers (always)
  - Save floating point or other special registers
  - Save condition codes
  - Change virtual address translations
- Non-negligible cost
  - Save/restore floating point registers expensive
    - Optimization: only save if process used floating point
  - May require flushing TLB (memory translation hardware)
    - HW Optimization 1: don't flush kernel's own data from TLB
    - HW Optimization 2: use tag to avoid flushing any data
- Usually causes more cache misses (switch working sets)


# Interprocess Communication

- Processes within a system may be **independent** or **cooperating**
  - Cooperating process can affect or be affected by other processes
  - Reasons for cooperating processes:
    - Information sharing
    - Computation speedup
    - Modularity
    - Convenience
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
  - Shared memory
  - Message passing

# Communications Models



(a)



(b)

# Producer-Consumer Problem

- Paradigm for cooperating processes:
  - Producer process produces information that is consumed by a consumer process
- Two variations:
  - Unbounded-buffer places no practical limit on the size of the buffer:
    - Producer never waits
    - Consumer waits if there is no buffer to consume
  - Bounded-buffer assumes that there is a fixed buffer size
    - Producer must wait if all buffers are full
    - Consumer waits if there is no buffer to consume

# Shared Memory

- An area of memory **shared among the processes** that wish to communicate
- The communication is **under the control of the users processes** not the operating system
- Major issues is to provide mechanism that will allow the user processes to **synchronize** their actions when they access shared memory.
  - **Synchronization** is discussed in great details later on

# Message Passing

- Processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - `send(message)`
  - `receive(message)`
  - The message size is either fixed or variable
- If processes P and Q wish to communicate, they need to:
  - Establish a communication link between them
    - How are links established?
    - Can a link be associated with more than two processes?
    - How many links can there be between every pair of communicating processes?
    - What is the capacity of a link?
    - Is a link unidirectional or bi-directional?
  - Exchange messages via send/receive
    - Is the size of a message that the link can accommodate fixed or variable?

# Implementation of Communication Link

- Physical:
  - Shared memory
  - Hardware bus
  - Network
- Logical:
  - Direct or indirect
  - Synchronous or asynchronous
  - Automatic or explicit buffering



# Direct Communication

- Processes must name each other explicitly:
  - `send(P, message)` – send a message to process P
  - `receive(Q, message)` – receive a message from process Q
- Properties of communication link
  - Links are established automatically
  - A link is associated with exactly one pair of communicating processes
  - Between each pair there exists exactly one link
  - The link may be unidirectional, but is usually bi-directional

# Indirect Communication

- Messages are directed and received from mailboxes
  - Each mailbox has a unique ID
  - Processes can communicate only if they share a mailbox
- Properties of communication link
  - Link established only if processes share a common mailbox
  - A link may be associated with many processes
  - Each pair of processes may share several communication links
  - Link may be unidirectional or bi-directional
- Operations
  - Create and delete a mailbox
  - Send and receive messages through mailbox
- Primitives are defined as:
  - `send(A, message)` – send a message to mailbox A
  - `receive(A, message)` – receive a message from mailbox A

# Examples of IPC Systems - POSIX

- POSIX Shared Memory


- Process first creates shared memory segment
  - `shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`
  - Also used to open an existing segment
- Set the size of the object
  - `ftruncate(shm_fd, 4096);`
- Use **mmap()** to memory-map a file pointer to the shared memory object
- Reading and writing to shared memory is done by using the pointer returned by **mmap()**

# Communications in Client-Server Systems

- Sockets
- Remote Procedure Calls

# Sockets

- A socket is defined as an endpoint for communication
- Concatenation of IP address and port
  - The socket 161.25.19.8:1625 refers to port 1625 on host 161.25.19.8
  - Special IP address 127.0.0.1 (loopback) to refer to system on which process is running
  - All ports below 1024 are well known, used for standard services
- Communication consists between a pair of sockets




# Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems
  - Again uses ports for service differentiation
- **Stubs** – client-side proxy for the actual procedure on the server
  - The client-side stub locates the server and marshalls the parameters
  - The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server
- Data representation handled via External Data Representation (XDL) format to account for different architectures
  - Big-endian and little-endian
- OS typically provides a rendezvous service to connect client and server

# Next Lecture

- We learn process creation and termination (assignment 1 is due tomorrow)



Credit: <https://xkcd.com/557/>