

06. Threads & Concurrency

CS 4352 Operating Systems

Processes

- Recall that a process includes many things
 - An address space (defining all the code and data pages)
 - OS resources (e.g., open files) and accounting information
 - Execution state (PC, SP, regs, etc.)
- Creating a new process is costly
 - Because of all of the data structures that must be allocated and initialized
- Communicating between processes is also costly
 - Because most communication goes through the OS
 - Overhead of system calls and copying data

Concurrent Programs

- A web server (or any parallel program)...
 - Forks off copies of itself to handle multiple simultaneous requests
- To execute these programs we need to
 - Create several processes that execute in parallel
 - Cause each to map to the same address space to share data
 - They are all part of the same computation
 - Have the OS schedule these processes
- This situation is very inefficient
 - Space: PCB, page tables, etc.
 - Time: create data structures, fork and copy addr space, etc.

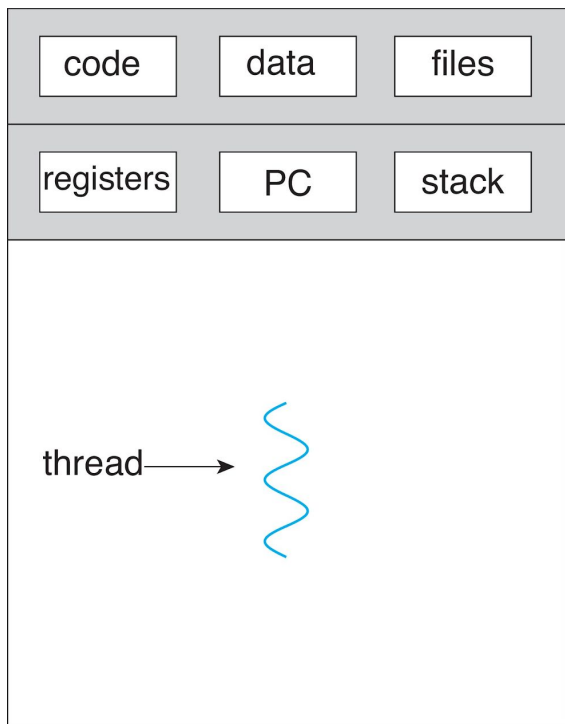
Rethinking Processes

- What is similar in these cooperating processes?
 - They all share the same code and data (address space)
 - They all share the same privileges
 - They all share the same resources (files, sockets, etc.)
- What don't they share?
 - Each has its own execution state: PC, SP, and registers
- Key idea: Why don't we separate the concept of a process from its execution state?
 - Process: address space, privileges, resources, etc.
 - Execution state: PC, SP, registers
- Execution state also called thread of control, or thread

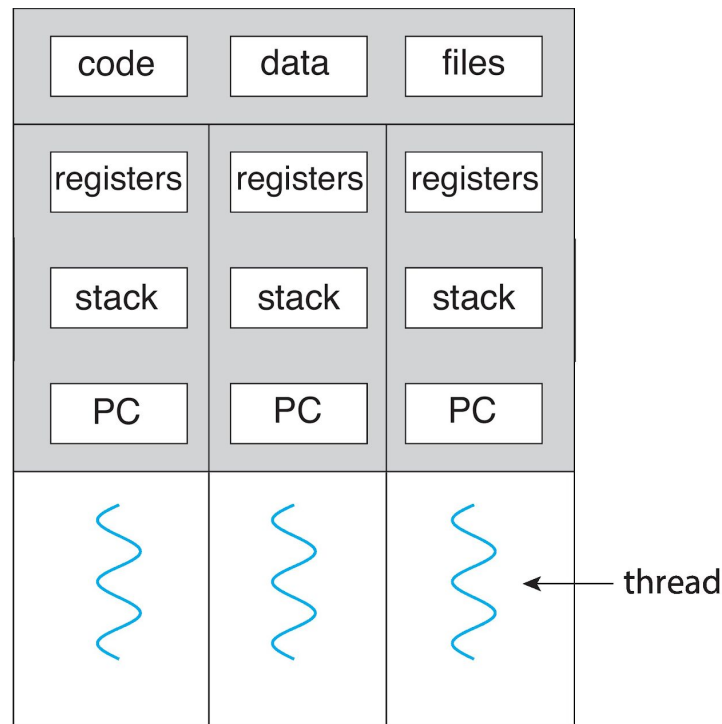
Threads

- Modern OSes separate the concepts of processes and threads
 - The thread defines a sequential execution stream within a process (PC, SP, registers)
 - The process defines the address space and general process attributes (everything but threads of execution)
- A thread is bound to a single process
 - Processes, however, can have **multiple threads**
- Most modern applications are multithreaded
 - Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- OS kernels are actually also multithreaded

Single and Multithreaded Processes

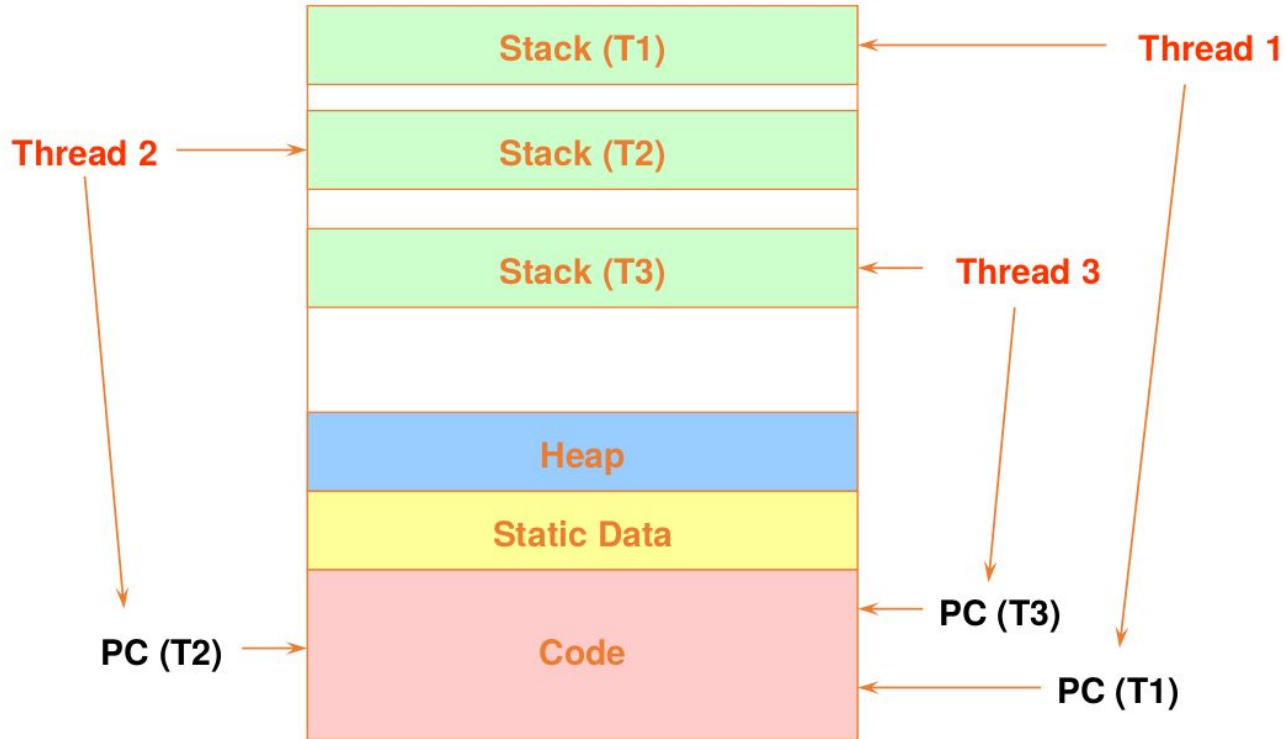


single-threaded process



multithreaded process

Threads in a Process

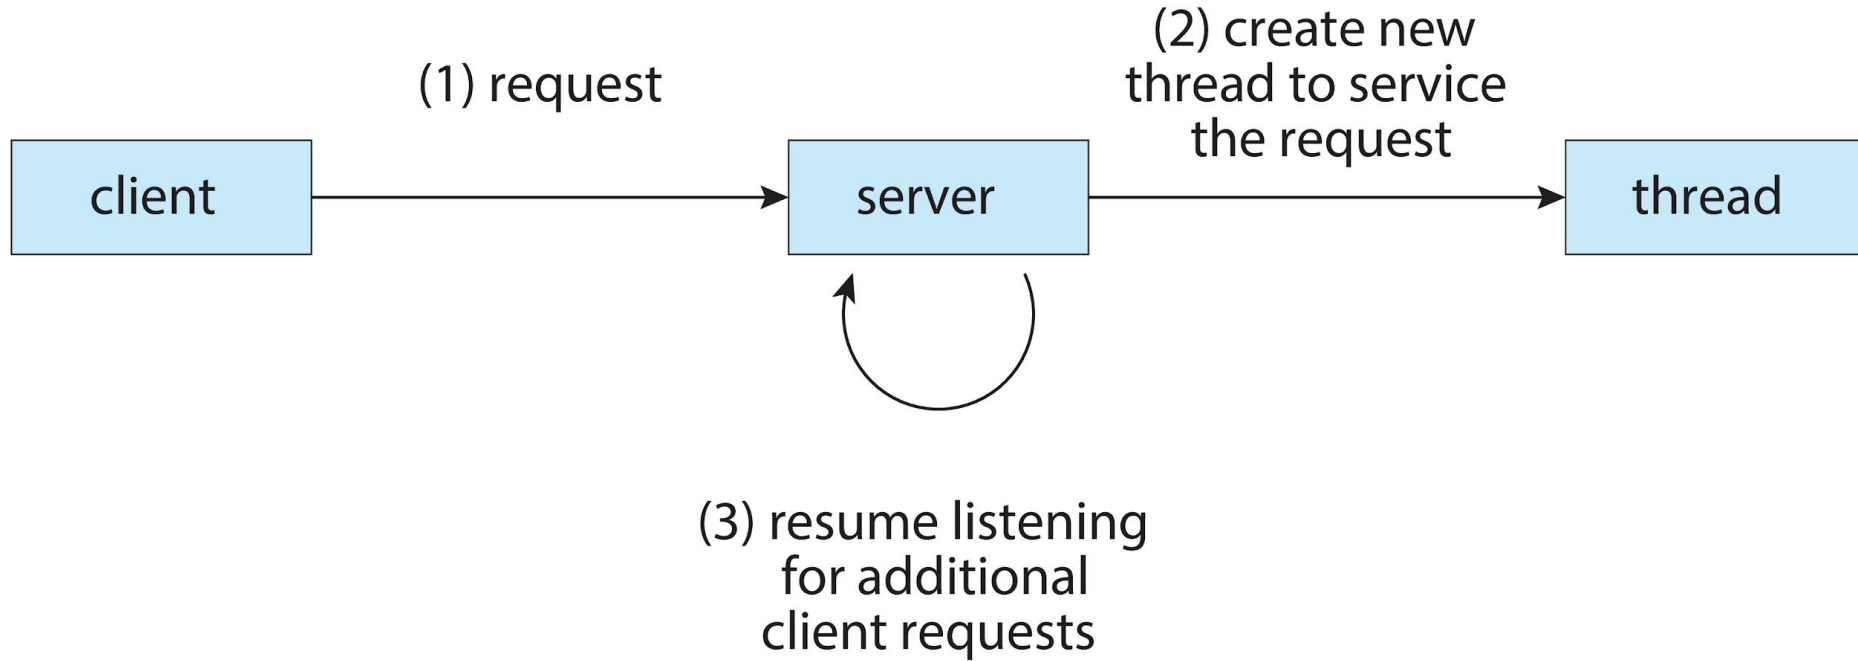


Web Server Example

- Using `fork()` to create new processes to handle requests is an overkill
- Instead, we can create a new thread for each request

```
web_server() {  
    while (1) {  
        int sock = accept();  
        thread_fork(handle_request, sock);  
    }  
}  
  
handle_request(int sock) {  
    Process request  
    close(sock);  
}
```


Multithreaded Server Architecture



Benefits


- Responsiveness – may allow continued execution if part of process is blocked
- Resource Sharing – threads share resources of process
- Economy – cheaper than process creation, thread switching lower overhead than context switching
- Scalability – process can take advantage of multicore architectures

Multicore Programming


- Multicore or multiprocessor systems putting pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
 - Single processor / core, scheduler providing concurrency

Concurrency v.s. Parallelism

- Concurrent execution on single-core system:



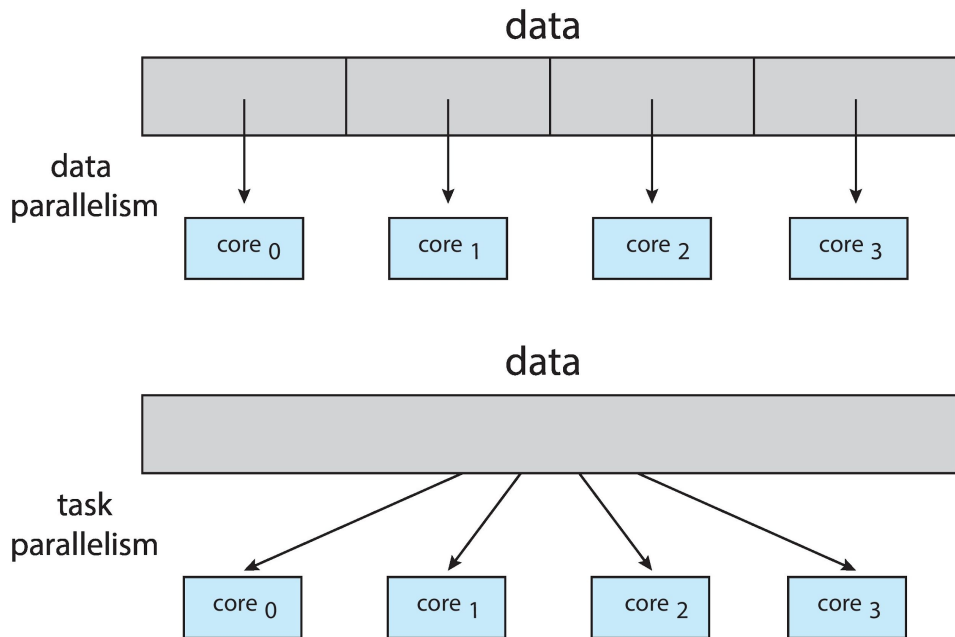
- Parallelism on a multi-core system:



Data and Task Parallelism

- Types of parallelism

- Data parallelism – distributes subsets of the same data across multiple cores, same operation on each
- Task parallelism – distributes threads across cores, each thread performing unique operation



Amda


- It ide
that l

-
-

-

- As N

-




plication

in speedup

by adding


User and Kernel Threads

- User threads - management done by user-level threads library
- Kernel threads - Supported by the kernel




Many-to-One Multithreading Model

- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Few systems currently use this model




One-to-One Multithreading Model

- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead



Many-to-Many Multithreading Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Windows with the ThreadFiber package
- Otherwise not very common



Thread Libraries

- Thread library provides programmer with APIs for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS

Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- **Specification**, not **implementation**
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX-like operating systems (Linux & Mac OS X)

Pthread APIs

- Basic functions:
 - `pthread_create()` -- create a thread
 - `pthread_join()` -- join (“wait for”) a terminated thread
 - `pthread_exit()` -- terminate the calling thread (does not cause the whole program to terminate)
- Other functions:
 - `pthread_attr_init()` -- specify “attributes” for the thread, e.g., stack size, is the thread detached
 - `pthread_self()` -- get a “handle” to a pthread
 - `pthread_cancel()` -- cancel a thread
 - `pthread_kill()` -- send a signal to a thread
 - `pthread_detach()` -- detach a thread: automatically free its resources when it's done
 - `pthread_equal()` -- compare two threads for equality

Pthread Example

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}
```

```
/* The thread will execute in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

Pthread Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

Semantics of fork() and exec()

- Does fork() duplicate only the calling thread or all threads?
 - Some UNIXes have two versions of fork
- exec() usually works as normal – replace the running process including all threads

Linux Threads

- Linux refers to them as **tasks** rather than **threads**
- Thread creation is done through **clone()** system call
- clone() allows a child task to share the address space of the parent task (process)

- Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

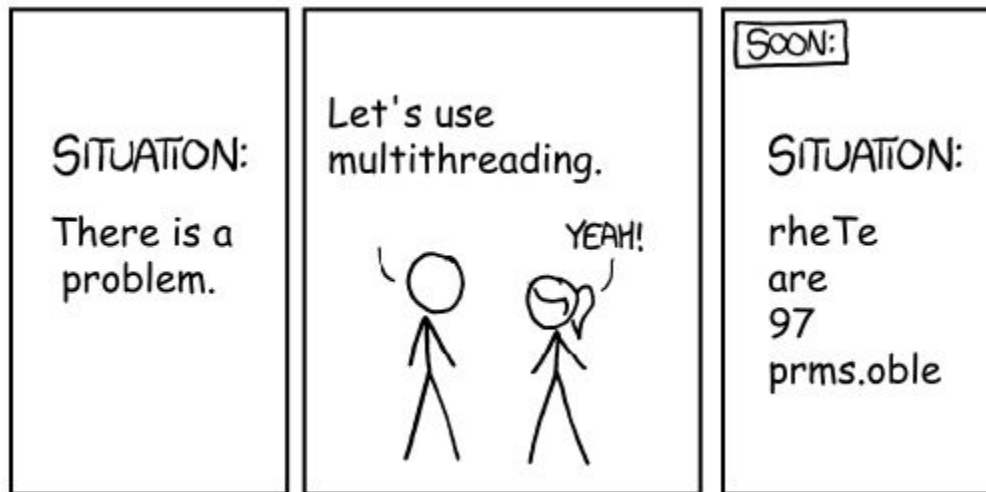
- **struct task_struct** is used for process and threads

Homework

- Read Chapters 6 & 7
- Assignment 2 is due tomorrow

Next Lecture

- We will look at synchronization



Credit: https://www.reddit.com/r/ProgrammerHumor/comments/dtiufv/multithreading_fixing_a_problem/