

08. Synchronization II

CS 4352 Operating Systems

Mutex Locks

- Previous solutions are generally inaccessible to application programmers
- OS designers build software tools to solve critical section problem
 - The simplest is **mutex lock**
 - Boolean variable indicating if lock is available or not
- Protect a critical section by
 - First **acquire()** a lock
 - Then **release()** the lock
- Calls to **acquire()** and **release()** must be atomic
 - Usually implemented via hardware atomic instructions such as compare-and-swap

Solution to CS Problem Using Mutex Locks

```
while (true) {  
    acquire lock  
  
    critical section  
  
    release lock  
  
    remainder section  
}
```

Deadlocks with Mutexes

- Consider the following scenario

Thread A

1. `pthread_mutex_lock(mutex1);`
2. `pthread_mutex_lock(mutex2);`
blocks

Thread B

1. `pthread_mutex_lock(mutex2);`
2. `pthread_mutex_lock(mutex1);`
blocks

- A deadlock may occur!
- Lesson: threads should always acquire locks in the same order

Pthread Mutexes

- Data type:
 - pthread_mutex_t
- Common operations:
 - int pthread_mutex_lock(pthread_mutex_t *mutex);
 - int pthread_mutex_unlock(pthread_mutex_t *mutex);
- Typical use:

```
static pthread_mutex_t mtx = PTHREAD_MUTEX_INITIALIZER;

void *thread_func(void *arg)
{
    pthread_mutex_lock(&mtx); // get the lock
    // access shared data
    pthread_mutex_unlock(&mtx); // release the lock
}
```

Locking the Same Lock?

- What if the same thread tries to obtain the same mutex multiple times?
 - The result depends on how the mutex was initialized
- Types of pthread mutexes:
 - PTHREAD_MUTEX_DEFAULT or PTHREAD_MUTEX_NORMAL
 - Results in a deadlock if the same pthread tries to lock it a second time using the pthread_mutex_lock subroutine without first unlocking it. This is the default type
 - PTHREAD_MUTEX_ERRORCHECK
 - Avoids deadlocks by returning a non-zero value if the same thread attempts to lock the same mutex more than once without first unlocking the mutex
 - PTHREAD_MUTEX_RECURSIVE
 - Allows the same pthread to recursively lock the mutex using the pthread_mutex_lock subroutine without resulting in a deadlock or getting a non-zero return value from pthread_mutex_lock. The same pthread has to call the pthread_mutex_unlock subroutine the same number of times as it called pthread_mutex_lock subroutine in order to unlock the mutex for other pthreads to use

Producer-Consumer Problem

- A canonical example: the producer-consumer problem
 - Producer threads produce elements
 - Consumer threads consume the elements produced by the producer threads
- A lock alone isn't a good solution:
 - It only ensures mutual exclusion
 - Consider the case where a consumer wants to run but there are no elements available:
 - Obtain lock
 - Check for elements
 - Release lock
 - Sleep
- Condition variables to the rescue!

Condition Variables

- A condition variable allows one thread to inform other threads about changes in the state of a shared variable (or other shared resource) and allows the other threads to wait (block) for such notification
- Common operations:
 - `int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex);`
 - `int pthread_cond_signal(pthread_cond_t *cond);`
 - `int pthread_cond_broadcast(pthread_cond_t *cond);`

Solution to Producer-Consumer Problem

```
static pthread_mutex_t mtx = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t cond_full = PTHREAD_COND_INITIALIZER;
pthread_cond_t cond_empty = PTHREAD_COND_INITIALIZER;
```

```
void *producer_func(void *arg)
{
    pthread_mutex_lock(&mtx);
    while (num_avail >= MAX_SIZE)
        pthread_cond_wait(&cond_empty, &mtx);
    num_avail++;
    pthread_mutex_unlock(&mtx);
    pthread_cond_signal(&cond_full);
}
```

```
void *consumer_func(void *arg)
{
    pthread_mutex_lock(&mtx);
    while (num_avail <= 0)
        pthread_cond_wait(&cond_full, &mtx);
    // consumer data and process
    num_avail--;
    pthread_mutex_unlock(&mtx);
    pthread_cond_signal(&cond_empty);
}
```

Mutex + Condition Variable

- The mutex associated with a condition variable is for mutual exclusion
- The condition variable is for signaling
- Important: always check the condition in a while loop!

```
void *consumer_func(void *arg)
{
    pthread_mutex_lock(&mtx);
    while (num_avail <= 0)
        pthread_cond_wait(&cond_full, &mtx);
    // consumer data and process
    num_avail--;
    pthread_mutex_unlock(&mtx);
    pthread_cond_signal(&cond_empty);
}
```

This atomically:

1. Unlocks the mutex
2. Waits on the condition variable

When execution reaches here, you have obtained the mutex, so you must unlock it

Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for processes to synchronize their activities
- Semaphore S – integer variable
 - Can only be accessed via two indivisible (atomic) operations
 - wait() and signal()
 - Originally called P() and V()
 - Dutch: Probeer (try) and Verhoog (increment) in Dijkstra's original paper
 - Counting semaphore – integer value can range over an unrestricted domain
 - Binary semaphore – integer value can range only between 0 and 1
 - Same as a mutex lock

Semaphore Implementation

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;  
  
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}  
  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

Semaphore Usage Example

- With semaphores we can solve various synchronization problems
- Solution to the critical section problem
 - Create a semaphore “mutex” initialized to 1
- Consider P1 and P2 that with two statements S1 and S2 and the requirement that S1 to happen before S2
 - Create a semaphore “synch” initialized to 0

`wait(mutex) ;`

`CS`

`signal(mutex) ;`

P1 :

`S1 ;`

`signal(synch) ;`

P2 :

`wait(synch) ;`

`S2 ;`

POSIX Semaphores (Unnamed)

- Data type:
 - `sem_t` (not `<pthread.h>` but `<semaphore.h>`)
- Common operations:
 - `int sem_init(sem_t *sem, int pshared, unsigned int value);`
 - `int sem_wait(sem_t *sem);`
 - `int sem_post(sem_t *sem);`
 - `int sem_getvalue(sem_t *sem, int *valp);`
 - `int sem_destroy(sem_t *sem);`

Semaphore Example

```
#include <pthread.h>
#include <stdio.h>
#include <semaphore.h>
```

```
sem_t mutex;
```

```
void *
foo(void *data)
{
    sem_wait(&mutex);

    //critical section

    sem_post(&mutex);
}
```

```
int
main()
{
    pthread_t thd;


    sem_init(&mutex, 0, 1);

    pthread_create(&thd, NULL, foo, NULL);
    pthread_join(thd, NULL);

    sem_destroy(&mutex);
    return 0;
}
```

Dining-Philosophers Problem

- N philosophers' sit at a round table with a bowl of rice in the middle
 - They spend their lives alternating thinking and eating
 - They do not interact with their neighbors
 - Occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
 - Need both to eat, then release both when done
 - In the case of 5 philosophers, the shared data
 - Bowl of rice (data set)
 - Semaphore chopstick[5] initialized to 1



What Can Go Wrong?

- What is the problem with the following solution?

```
while (true){
    wait(chopstick[i]);
    wait(chopstick[(i + 1) % 5]);

    /* eat for a while */

    signal(chopstick[i]);
    signal(chopstick[(i + 1) % 5]);

    /* think for a while */
}
```

Linux Synchronization

- Linux:
 - Prior to kernel Version 2.6, disables interrupts to implement short critical sections
 - Version 2.6 and later, fully preemptive
- Linux provides:
 - Semaphores
 - Atomic integers
 - Spinlocks
 - Busy waiting!
- On single-CPU system, spinlocks replaced by enabling and disabling kernel preemption
 - Why?

Atomic Variables

- `atomic_t` is the type for atomic integer
 - `atomic_t counter;`

<i>Atomic Operation</i>	<i>Effect</i>
<code>atomic_set(&counter,5);</code>	<code>counter = 5</code>
<code>atomic_add(10,&counter);</code>	<code>counter = counter + 10</code>
<code>atomic_sub(4,&counter);</code>	<code>counter = counter - 4</code>
<code>atomic_inc(&counter);</code>	<code>counter = counter + 1</code>
<code>value = atomic_read(&counter);</code>	<code>value = 12</code>

Linux Implementation of Mutexes

- Locking:
 - Bit 31: indicates if lock is taken
 - 1: taken
 - 0 free
 - Remaining bits: number of waiters
 - Line 7: don't return from here until we get the lock
 - Lines 8-10: check lock is free and decrement # of waiters if so
 - Decrement because we incremented
 - Line 15: check if lock is taken
 - Line 17: futex system call
 - Put calling process on queue

```
1 void mutex_lock (int *mutex) {
2     int v;
3     /* Bit 31 was clear, we got the mutex (the fastpath) */
4     if (atomic_bit_test_set (mutex, 31) == 0)
5         return;
6     atomic_increment (mutex);
7     while (1) {
8         if (atomic_bit_test_set (mutex, 31) == 0) {
9             atomic_decrement (mutex);
10            return;
11        }
12        /* We have to waitFirst make sure the futex value
13         we are monitoring is truly negative (locked). */
14        v = *mutex;
15        if (v >= 0)
16            continue;
17        futex_wait (mutex, v);
18    }
19 }
20
21 void mutex_unlock (int *mutex) {
22     /* Adding 0x80000000 to counter results in 0 if and
23     only if there are not other interested threads */
24     if (atomic_add_zero (mutex, 0x80000000))
25         return;
26
27     /* There are other threads waiting for this mutex,
28     wake one of them up. */
29     futex_wake (mutex);
30 }
```

Linux Implementation of Mutexes

- Unlocking:
 - Lines 24-25: if `*mutex == 0` after adding `0x80000000`, then nobody was waiting
 - Line 29: invoke the futex system call
 - Argument of `FUTEX_WAKE`
- `futex()` is a multiplexed system call
 - Different arguments change the behavior
 - `FUTEX_WAIT`: go to sleep until mutex is available
 - `FUTEX_WAKE`: wake up someone waiting on this mutex
- Described in paper “Futexes are Tricky” by Ulrich Drepper

```
1 void mutex_lock (int *mutex) {
2     int v;
3     /* Bit 31 was clear, we got the mutex (the fastpath) */
4     if (atomic_bit_test_set (mutex, 31) == 0)
5         return;
6     atomic_increment (mutex);
7     while (1) {
8         if (atomic_bit_test_set (mutex, 31) == 0) {
9             atomic_decrement (mutex);
10            return;
11        }
12        /* We have to waitFirst make sure the futex value
13         we are monitoring is truly negative (locked). */
14        v = *mutex;
15        if (v >= 0)
16            continue;
17        futex_wait (mutex, v);
18    }
19 }
20
21 void mutex_unlock (int *mutex) {
22     /* Adding 0x80000000 to counter results in 0 if and
23      only if there are not other interested threads */
24     if (atomic_add_zero (mutex, 0x80000000))
25         return;
26
27     /* There are other threads waiting for this mutex,
28      wake one of them up. */
29     futex_wake (mutex);
30 }
```

Homework

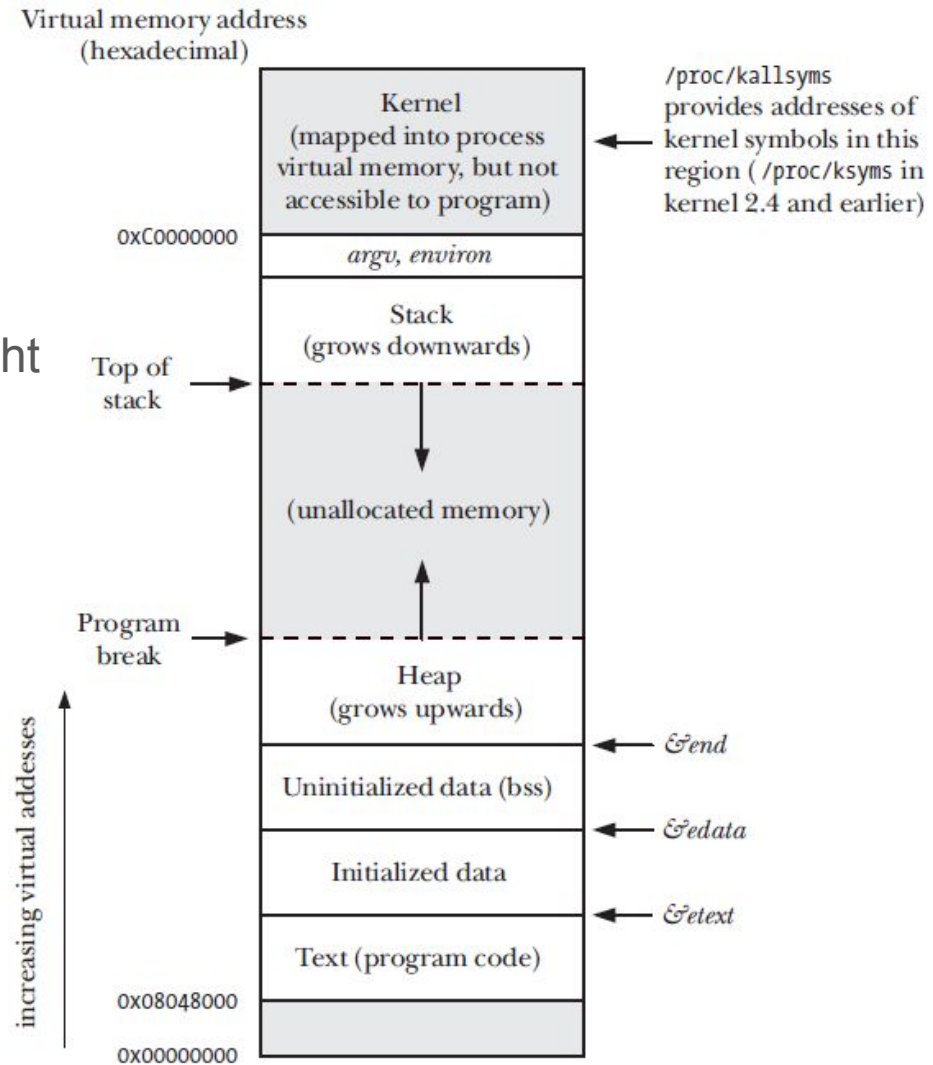
- Start reading Chapters 9 & 10
- Assignment 3 deadline is changed to 10/09

Next Lecture

Let's start looking at memory management!

Miniproject-3 Q&A

- Typical address space layout of a user-level process is shown on the right
- These addresses are all virtual
 - It's the job of the hardware (MMU) + OS to translate these to physical addresses



Check Your Answer

- `/proc/<pid>/maps`
 - It gives each virtual memory address region in a process
 - Address range
 - Permission
 - Offset in a file
 - Where file locate (device, inode, path)
 - You can use `objdump -D` to find all sections
 - Anonymous ones are all 0 in the following fields (no backing files)