

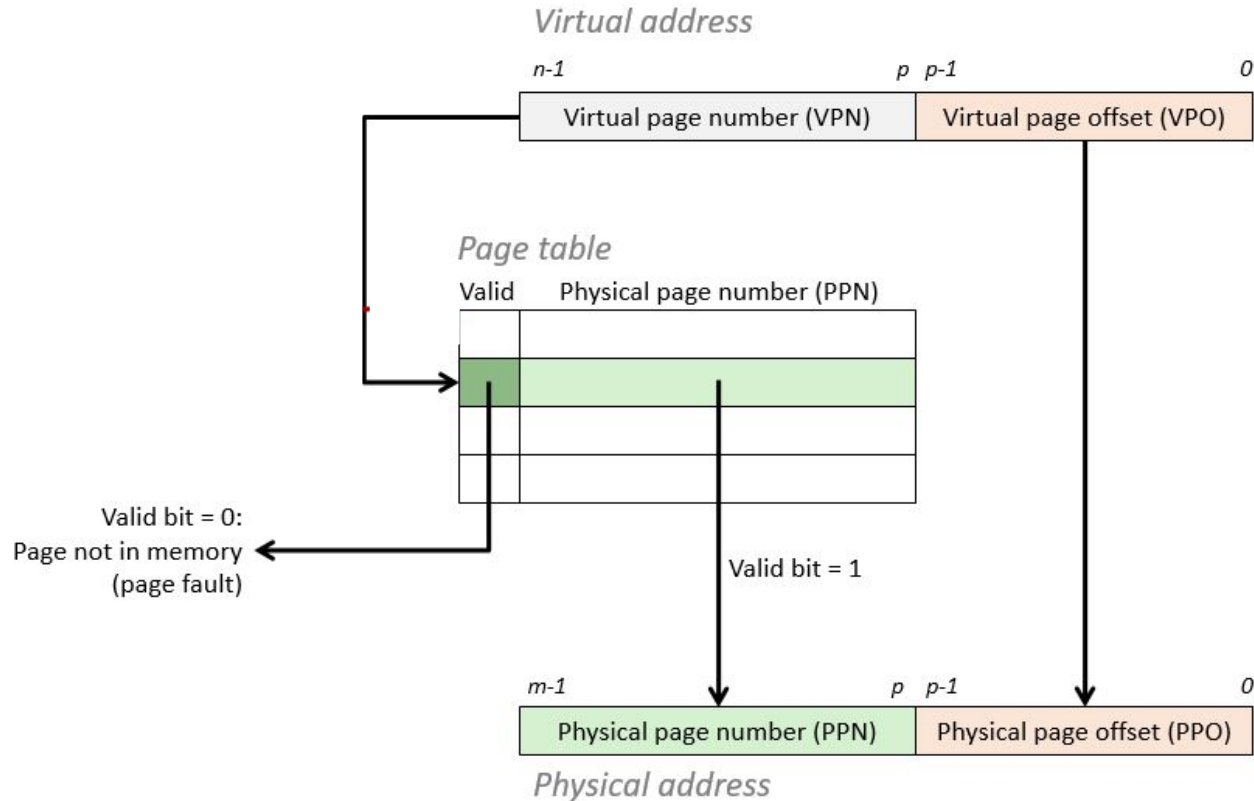
10. More Paging

CS 4352 Operating Systems

Review

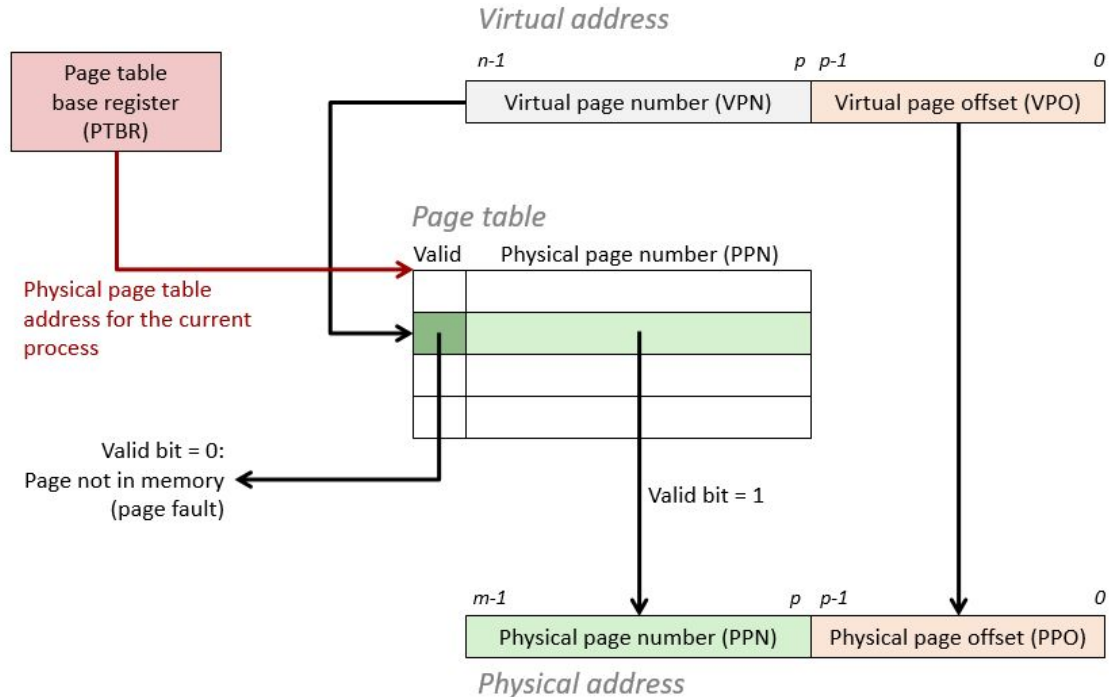
- Simple systems (e.g., elevators) might use physical addressing
 - The addresses used by a program correspond to real, physical addresses
 - Also called single memory, or flat address space
- All modern operating systems use virtual addressing
 - The addresses used by a program are translated to physical addresses
 - This translation is done mostly in hardware; the OS plays a big role in setting up data structures used by the hardware

Basic Address Translation



How to Find Page Table


- Use a special register (page table base register) to hold the physical address of the page table
 - On x86, CR3 register is used
 - Stored in PCB
- Context switch
 - PTBR is loaded a new value
 - Memory mapping



Managing Page Tables

- A single level page table usually takes up too much space
 - 32-bit address space, 4KB page size, 4-byte PTE
 - Need 2^{20} PTEs; each is 4 bytes
 - $2^{22} = 4\text{MB}$ -- not too bad, but still quite a bit
 - How about 48-bit address space, 4KB page size, 8-byte PTE
 - 2^{39} bytes = 512 GB -- way too big!
- How can we reduce this overhead?
 - Observation: Only need to map the portion of the address space actually being used (tiny fraction of entire address space)
- How do we only map what is being used?
 - Can dynamically extend page table...
- Levels of indirection!

Two-Level Page Tables



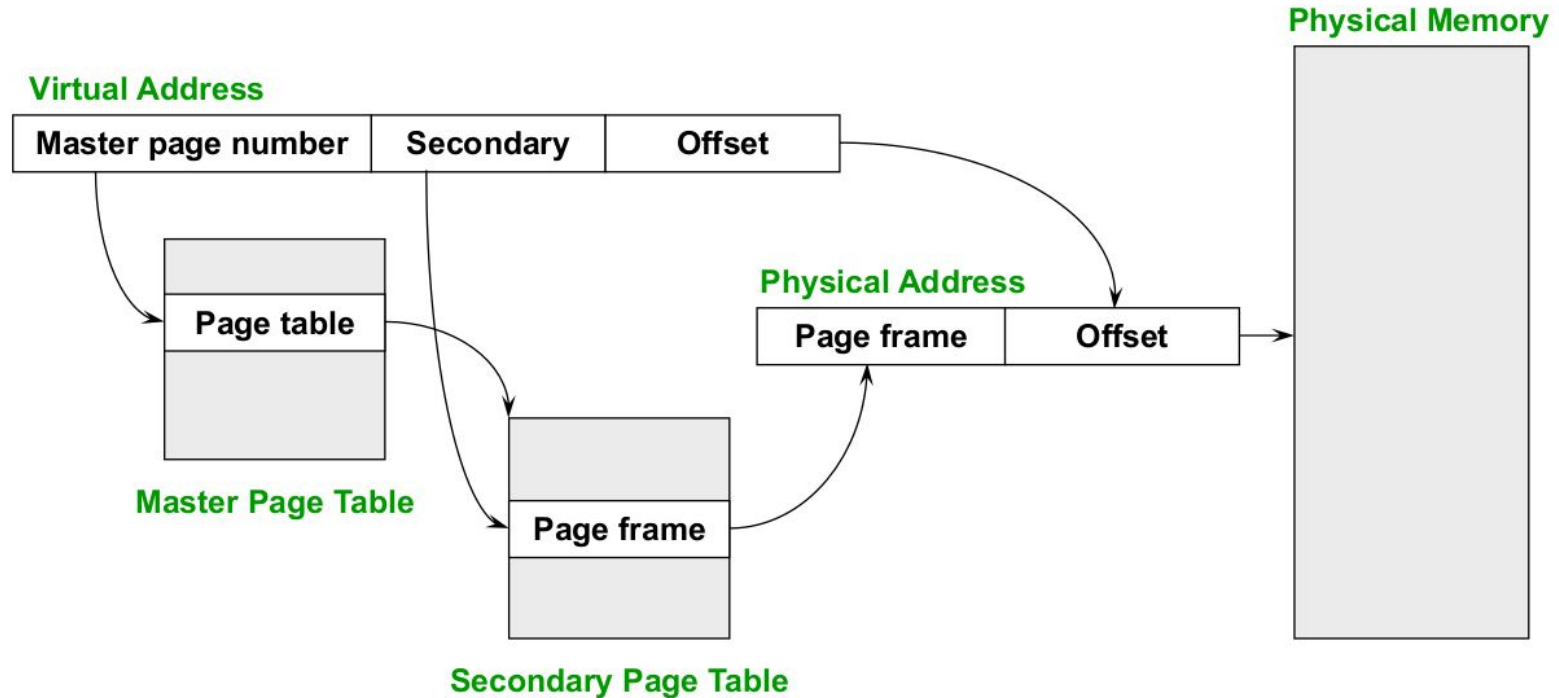
- Two-level page tables

- Virtual address have three parts:
 - Master page number, secondary page number, and offset
- Master page table maps VAs to secondary page table
- Secondary page table maps page number to physical page frame
- Offset indicates where in physical page address is located

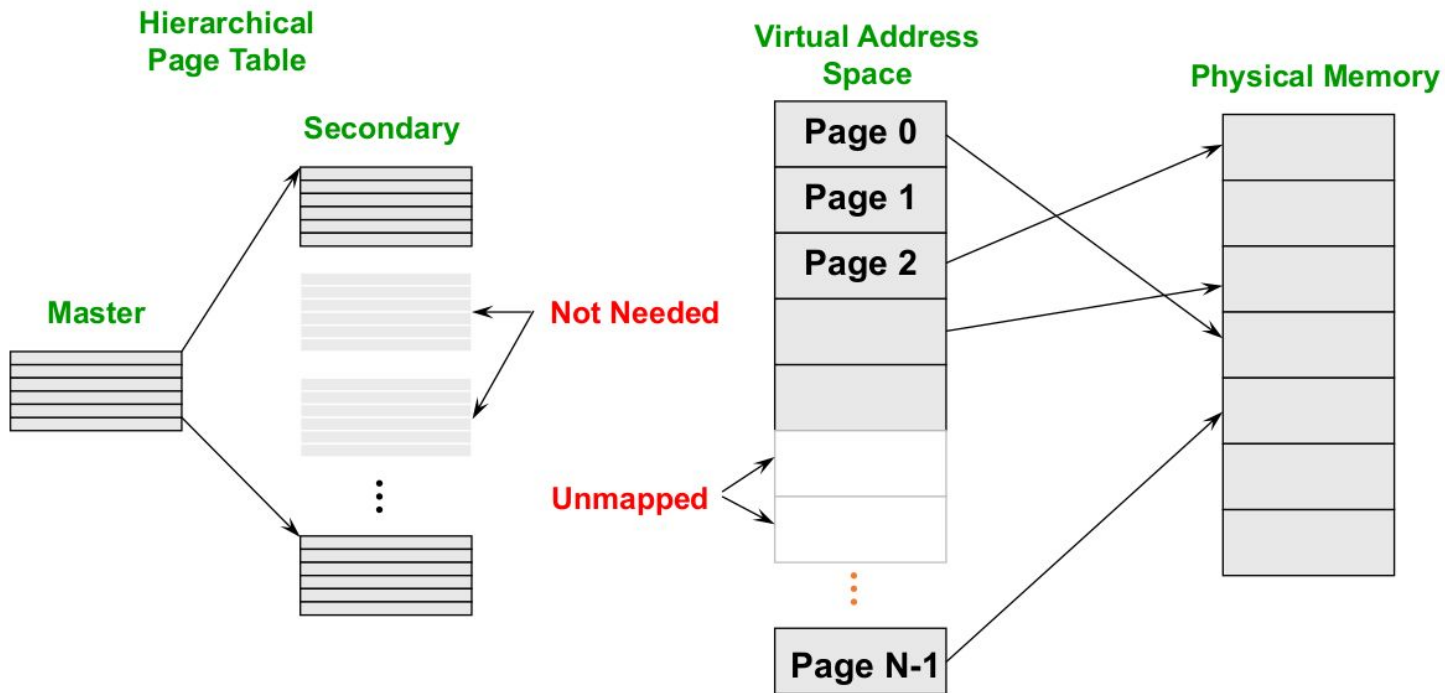
- Example

- 4KB pages, 4 bytes/PTE, how many bits in offset?
 - 12 bits
- Want master page table in one page, how many entries?
 - $4\text{KB} / 4 \text{ bytes} = 1024 \text{ entries}$
- 1024 secondary page tables. How many bits?
 - Master = 10, offset = 12, inner = $32 - 10 - 12 = 10 \text{ bits}$

Two-Level Page Tables

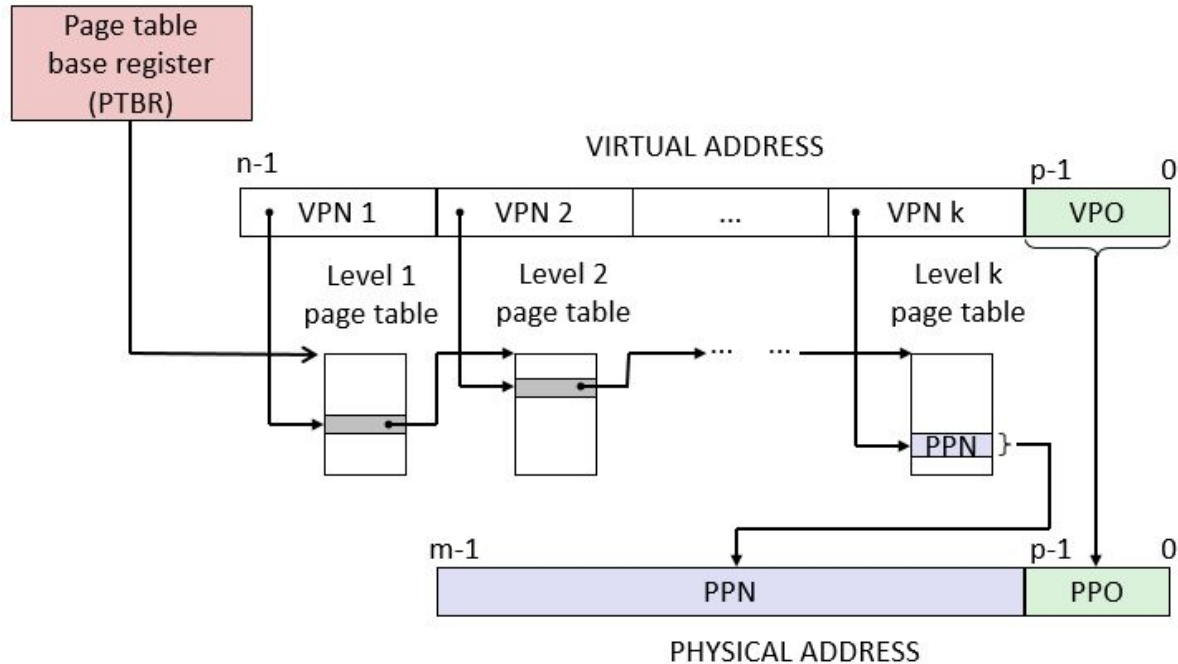


Page Table Evolution



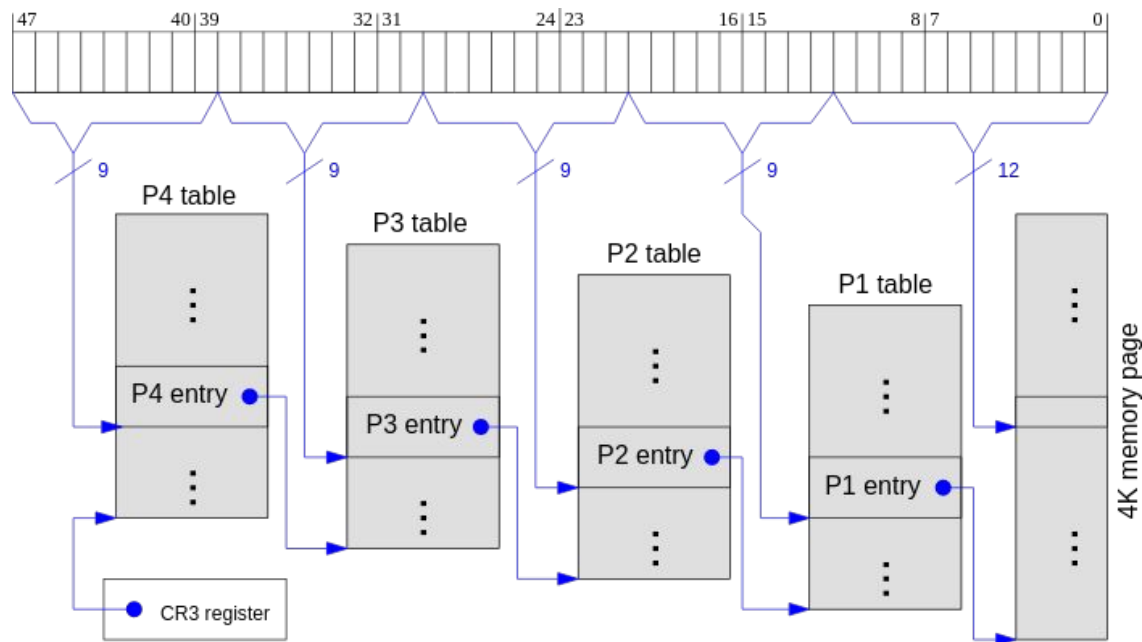
Generic Multi-Level Page Table

- Use PTBR to hold the physical address of the topmost table
 - This table normally occupy exactly one page frame



Modern OS + x86-64 Example

- 64-bit OS uses 48-bit virtual addresses and 4KB pages (normally)
 - 12 bits for offset; 9 bits for each lookup level



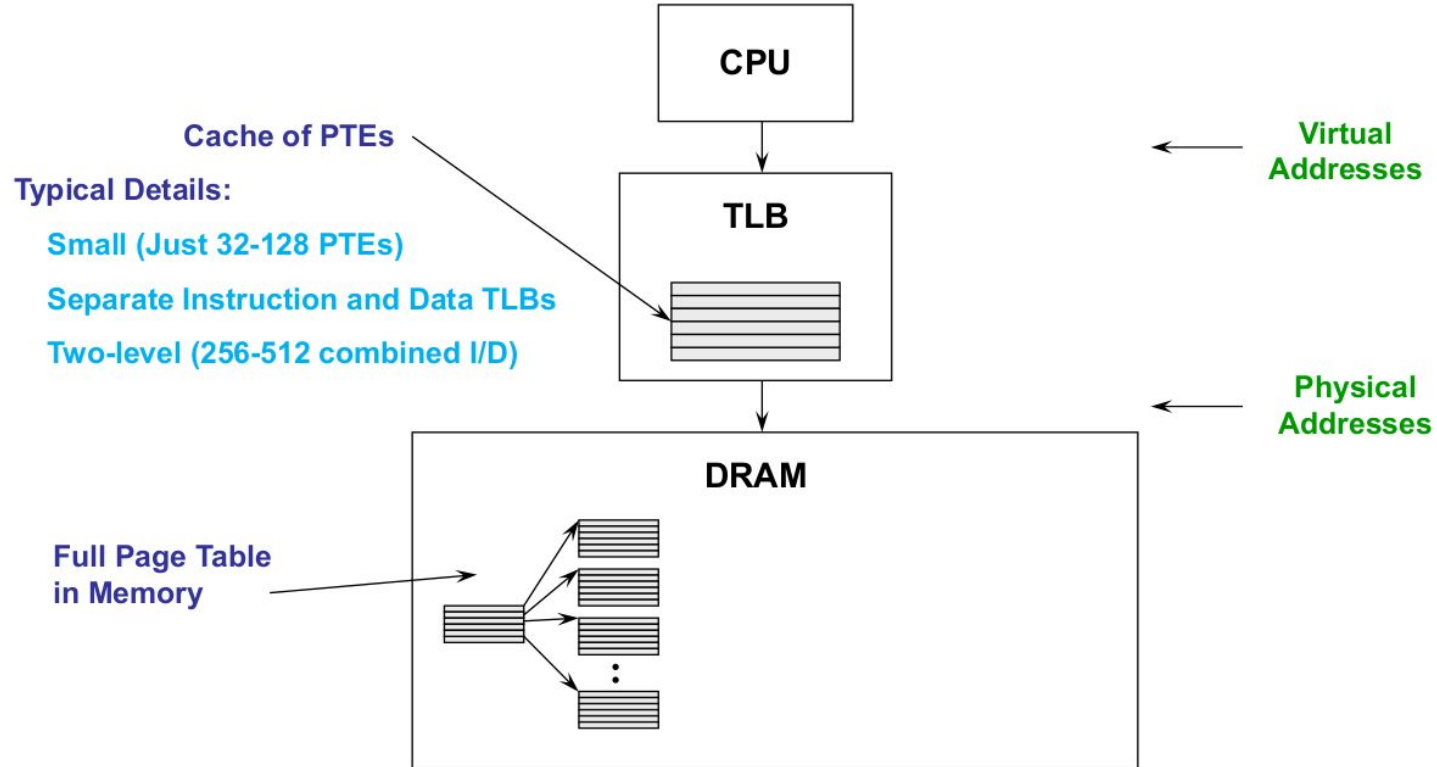
Efficient Translations

- Our original page table already doubled the cost of memory access
 - One lookup into the page table, another to fetch the data
- Now multi-level page tables increase the cost too much!
 - Several lookups into the page tables, then to fetch the data
- How can we use paging but also reduce lookup cost?
 - Cache translations in hardware
 - Translation Lookaside Buffer (TLB)
 - TLB managed by Memory Management Unit (MMU)

TLB

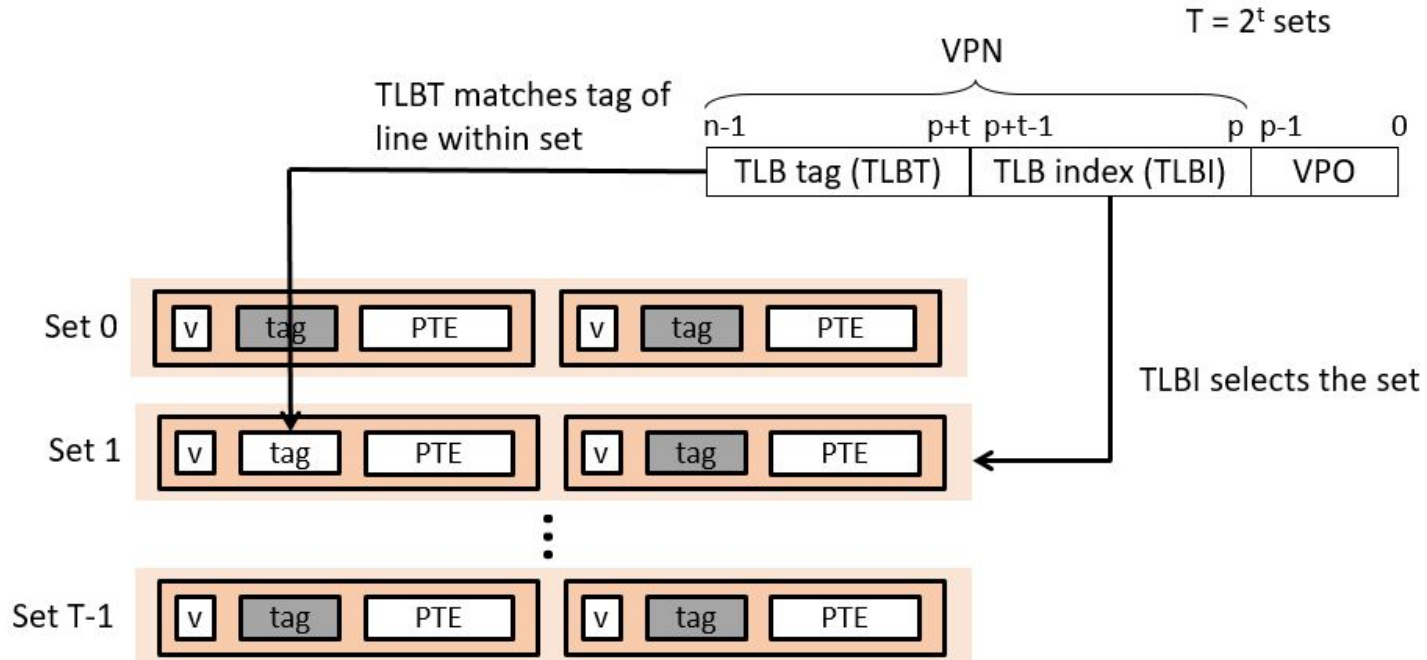
- Translation Lookaside Buffer (TLB)
 - Small set-associative hardware cache in MMU
 - Maps virtual page numbers to physical page frames
- TLBs exploit locality
 - Processes only use a handful of pages at a time (its working set)
- Some architectures have multiple TLBs:
 - Instruction TLB (ITLB)
 - Data TLB (DTLB)
- Can even have multiple levels of TLBs
 - Smaller ones are faster

TLB & Page Table



TLB in More Detail

- MMU uses the VPN portion of the virtual address to access the TLB



Managing TLBs

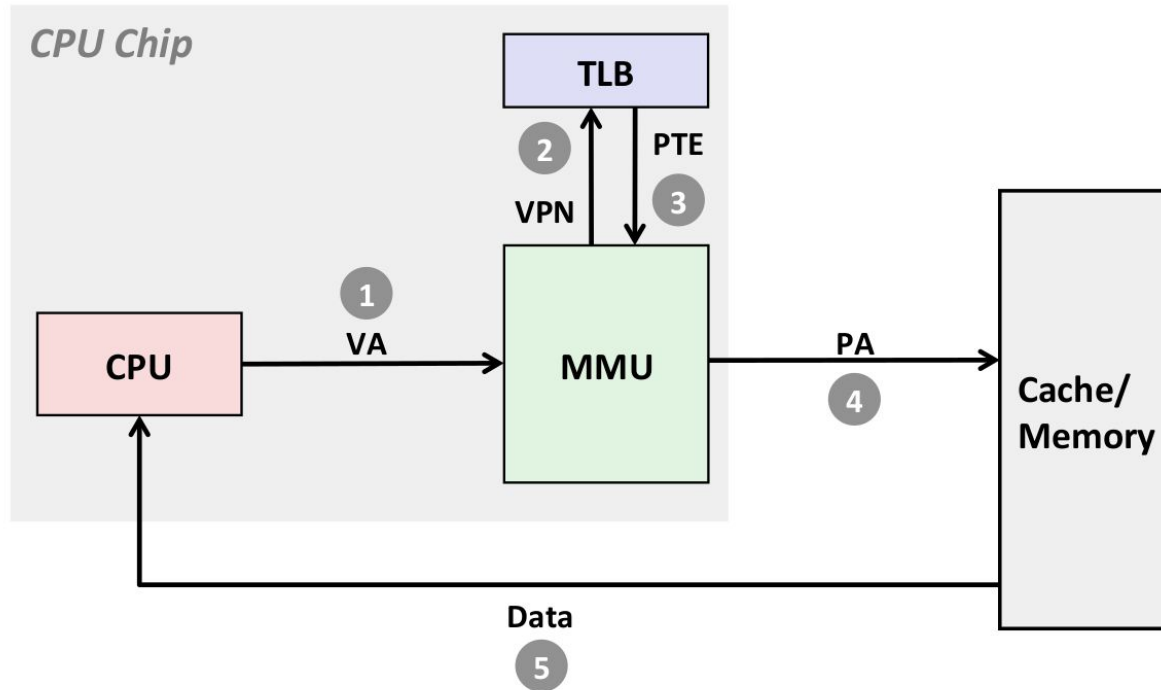
- Address translations for most memory accesses are handled using the TLB
 - But there are misses (TLB miss) ...
- Who places translations into the TLB (loads the TLB)?
 - MMU, e.g., x86
 - Knows where page tables are in main memory
 - OS maintains tables, HW accesses them directly
 - Tables have to be in HW-defined format (inflexible)
 - Software loaded TLB (OS), e.g., MIPS, Alpha, Sparc, PowerPC
 - TLB faults to the OS, OS finds appropriate PTE, loads it in TLB
 - Must be fast (but still 20-200 cycles)
 - CPU ISA has instructions for manipulating TLB
 - Tables can be in any format convenient for OS (flexible)

Managing TLBs (Cont.)

- OS ensures that TLB and page tables are consistent
 - When protection bits of a PTE are changed, it needs to invalidate the PTE if it is in the TLB
- When a context switch happens, the TLB entries are often invalid
 - On the x86, the kernel can flush the entire TLB or individual entries
 - Which is better depends; the kernel can't always know ahead of time
- When the TLB misses and a new PTE has to be loaded, a cached PTE must be evicted
 - Choosing PTE to evict is called the TLB replacement policy
 - Implemented in hardware, often simple (e.g., LRU)

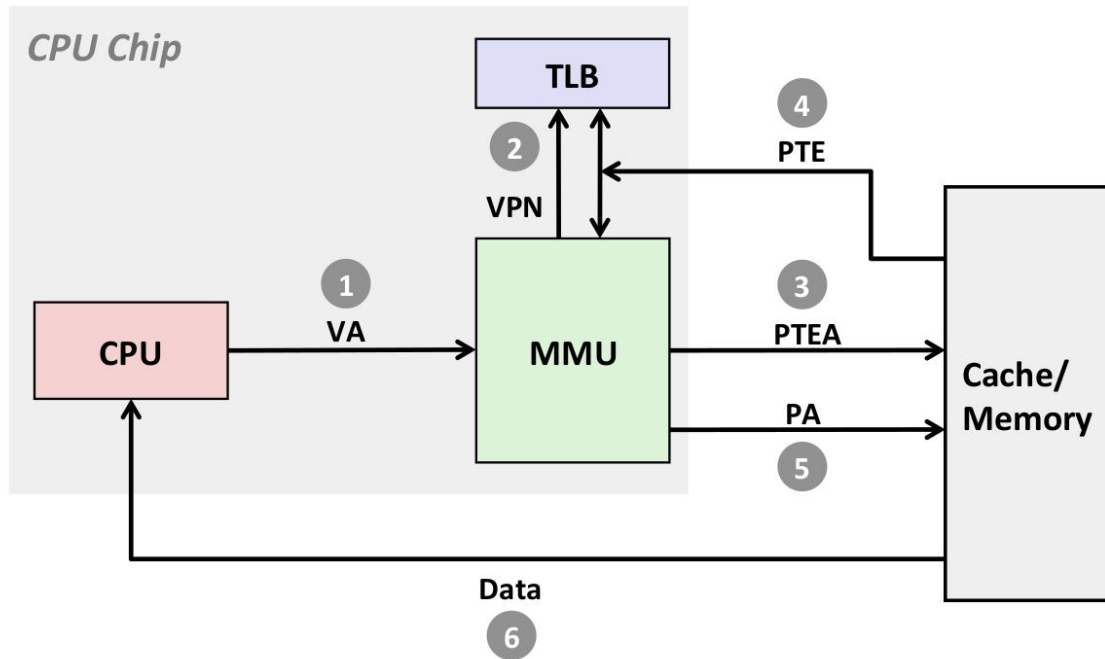
TLB Hit

- A TLB hit eliminates accesses to the page table



TLB Miss


- A TLB miss incurs additional memory accesses (the PTEs)
 - Fortunately, TLB misses are rare, why?



Paged Virtual Memory

- Pages can be moved between memory and disk
 - Use disk to simulate larger virtual than physical mem
 - This process is called paging in/out

- Paging process over time
 - Initially, pages are allocated from memory
 - When memory fills up, allocating a page requires some other page to be evicted
 - Evicted pages go to disk (where? the swap file/backing store)
 - Done by the OS, and transparent to the application



Demand Paging

- Pages are only brought into memory when a reference is made to a location on that page
 - In other words, the page isn't resident until it's accessed (read from or written to)
- Demand paging is really good!
 - The OS avoids wasting time/space loading parts of your program that are never referenced
 - Example: if you never make it past level 1 in a video game, the OS can avoid loading up the code for level 2

Next Lecture

Page faults and replacement policies