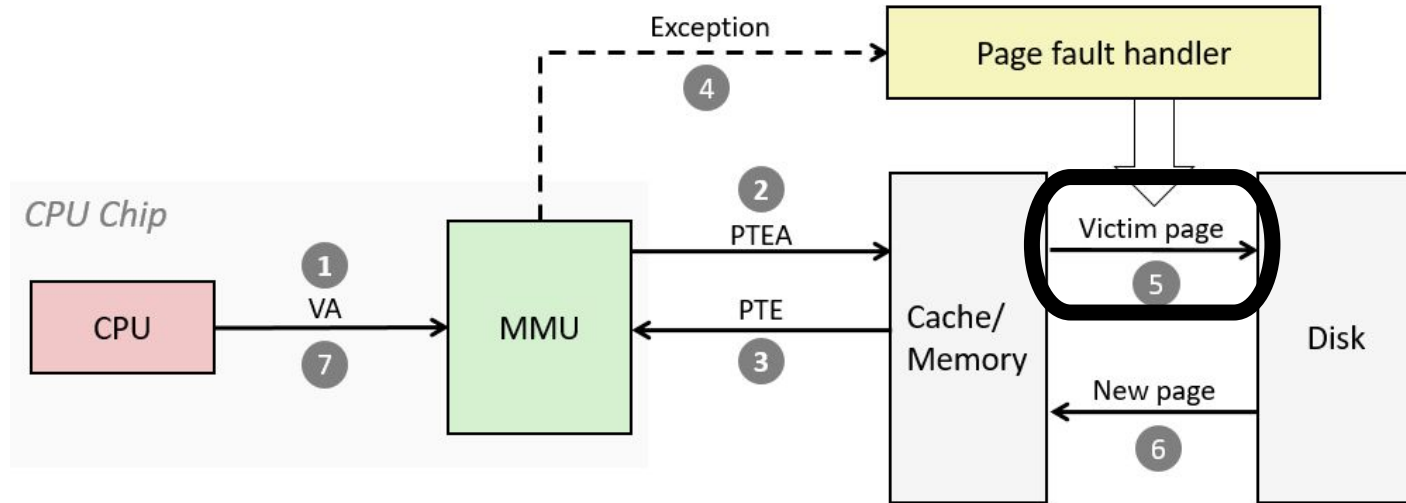


12. Page Replacements

CS 4352 Operating Systems

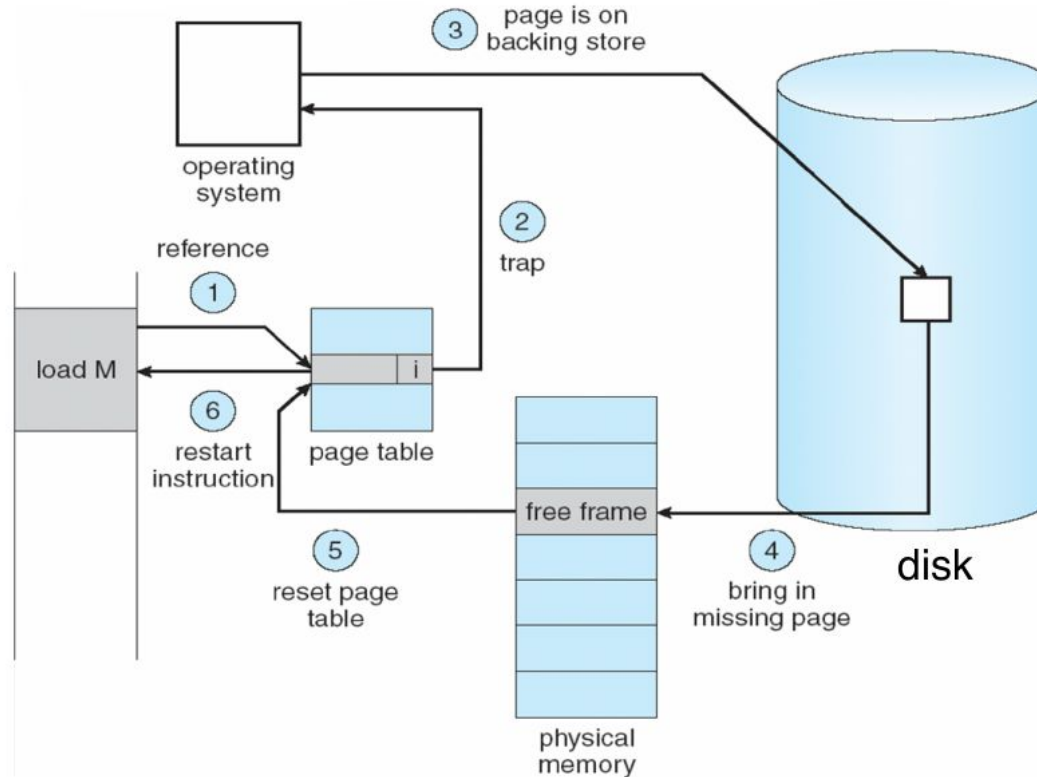
Review



Paging Review

- Paging from the OS perspective:
 - Pages are evicted to disk when memory is full
 - Pages loaded from disk when referenced again
 - References to evicted pages cause a page fault
 - OS allocates a page frame, reads page from disk
 - When I/O completes, the OS fills in PTE, marks it valid, and restarts faulting process
- Dirty vs. clean pages
 - Actually, only dirty pages (modified) need to be written to disk
 - Clean pages do not – but you need to know where on disk to read them from again

Use disk to simulate larger virtual memory




Principle of Locality

- We have mentioned this previously
 - Temporal locality
 - If a memory location is accessed once, it is likely that this memory location will be accessed again in the near future
 - Spatial locality
 - If a memory location is accessed once, it is likely that some nearby memory locations will be accessed in the near future
- Although the cost of paging is high, if it is infrequent enough it is acceptable
 - Processes usually exhibit both kinds of locality during their execution, making paging practical

80/20 Rule



- 20% of memory gets 80% of memory accesses
 - Keep the hot 20% in memory
 - Keep the cold 80% on disk



Page Replacement


- When a page fault occurs, the OS loads the needed page from disk into a page frame of physical memory
- At some point, the process used all of the page frames it is allowed to use
 - This is likely (much) less than all of available memory
- When this happens, the OS must replace a page for each page faulted in
 - It must evict a page to free up a page frame
- The page replacement policies determines how this is done
 - Find some page in memory, but not really in use, page it out
 - E.g., use modify (dirty) bit to reduce overhead of page transfers – only modified pages are written to disk
- Page replacement enables large virtual memory to be provided on a smaller physical memory

Page Replacement



First-In First-Out (FIFO) Algorithm


- Just use a FIFO queue and the victim is the first
- Reference string: 7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1
- 3 frames (3 pages can be in memory at a time per process)



- 15 page faults


FIFO has Belady's Anomaly

- Adding more frames can cause more page faults!
 - Example—reference string 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
 - 3 page frames: 9 page faults
 - 4 page frames: 10 page faults



Optimal Algorithm

- Replace page that will not be used for longest period of time
 - 9 is optimal for the example
- How do you know this?
 - Can't read the future
- Used for measuring how well your algorithm performs



Least Recently Used (LRU) Algorithm

- Use past knowledge rather than future
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used
- But how to implement?

Possible Implementations

- Counter implementation

- Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
- When a page needs to be changed, look at the counters to find smallest value
- Search through table needed

- Stack implementation

- Keep a stack of page numbers in a doubly linked form
- Page referenced:
 - move it to the top
 - requires 6 pointers to be changed
- No search for replacement
- But each update more expensive



LRU Approximation Algorithms

- LRU needs special hardware and still slow
- Reference bit
 - With each page associate a bit, initially = 0
 - When page is referenced bit set to 1
 - Replace any with reference bit = 0 (if one exists)
 - We do not know the order, however

LRU Approximation Algorithms (Cont.)

- Second-chance algorithm

- Generally FIFO, plus hardware-provided reference bit
- Clock replacement
- If page to be replaced has
 - Reference bit = 0 -> replace it
 - reference bit = 1 then:
 - set reference bit 0, leave page in memory
 - replace next page, subject to same rules



Other Replacement Algorithms


- Random eviction
 - Simple to implement
 - Not bad actually
- LFU (least frequently used) eviction
 - Instead of just a reference bit, count # times each page accessed
 - Least frequently accessed may not be very useful
 - Decay usage counts over time
- MFU (most frequently used) algorithm
 - Because page with the smallest count was probably just brought in and has yet to be used
- Neither LFU nor MFU used very commonly

Fixed v.s. Variable Space


- How to determine how much memory to give to each process?
- Fixed space algorithms
 - Each process is given a limit of pages it can use
 - When it reaches the limit, it replaces from its own pages
 - Local replacement
 - Some processes may do well while others suffer
- Variable space algorithms
 - Process' set of pages grows and shrinks dynamically
 - Global replacement
 - One process can ruin it for the rest

Thrashing


- When OS spent most of the time in paging data back and forth from disk
 - Little time spent doing useful work (making progress)
 - In this situation, the system is overcommitted
 - No idea which pages should be in memory to reduce faults
- Reasons for Thrashing
 - Access pattern has no temporal locality




- Hot memory does not fit in physical memory



- Each process fits individually, but too many for system



Thrashing & Multiprogramming

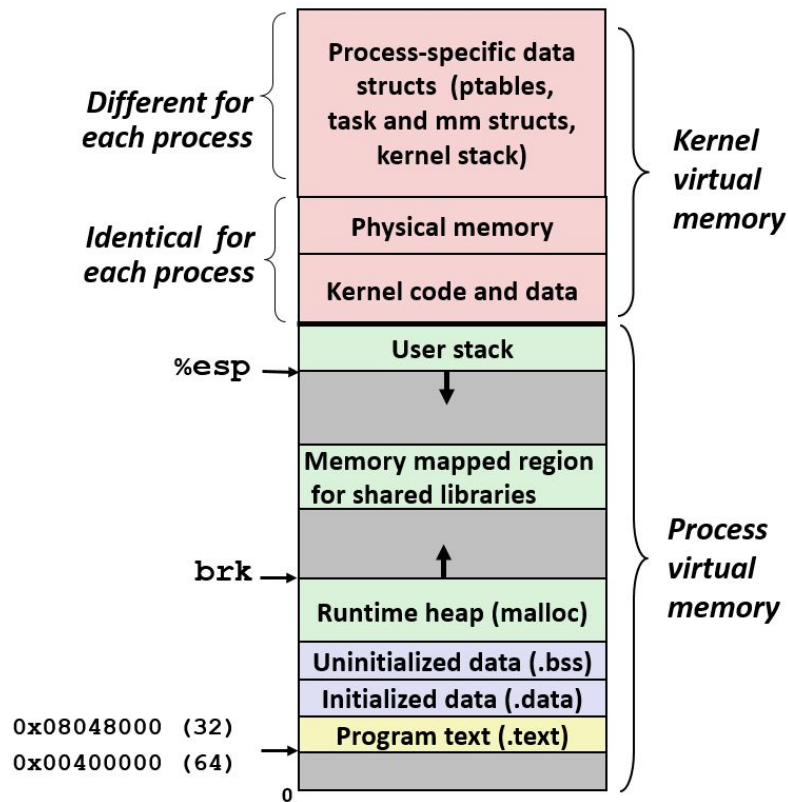


Dealing with Thrashing

- Only run processes if memory requirements can be satisfied
- Write out all pages of a process
- OOM
- Buy more memory...

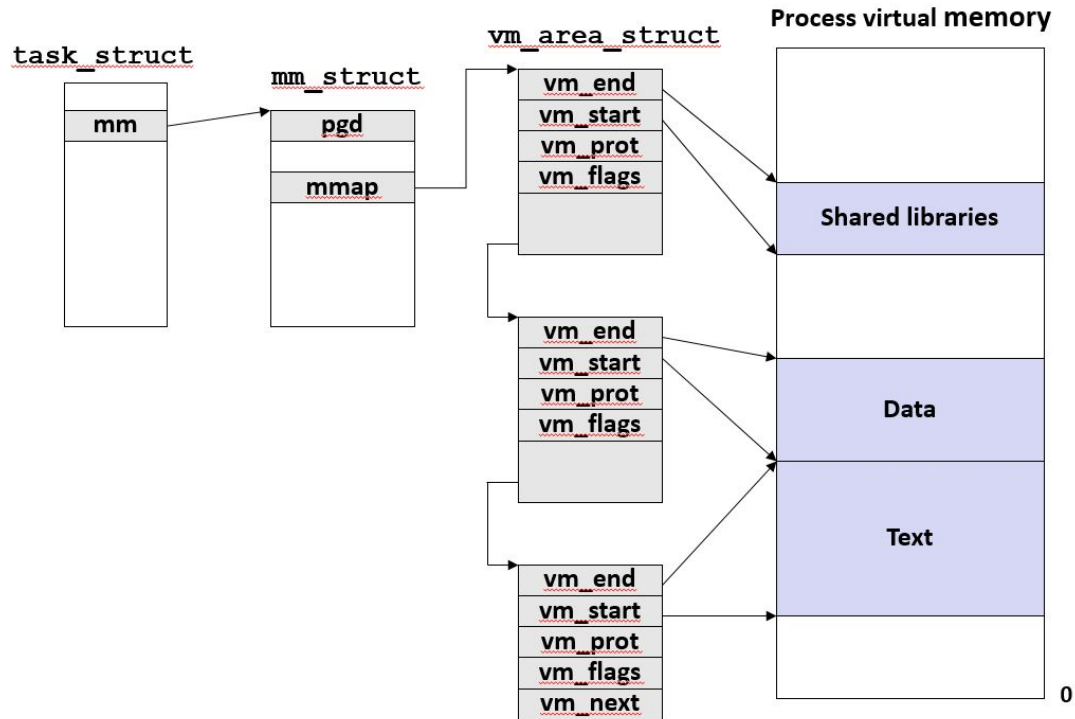
Linux VM revisited

- We're familiar with the address space diagram on the right
- But how does the Linux kernel view this?



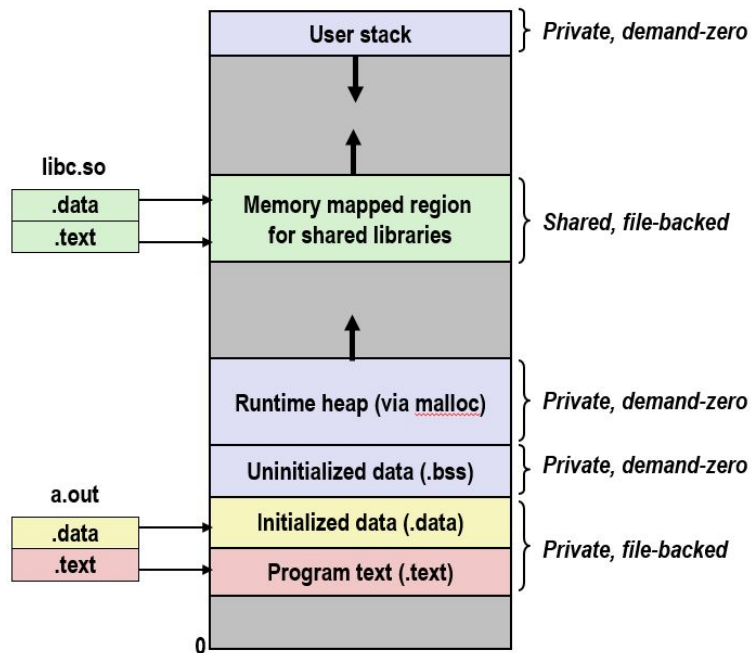
Linux Representation of Memory

- The Linux kernel organizes VM as a collection of “areas”
 - `vm_start` and `vm_end`: point to the beginning and end of the area
 - `vm_prot`: describes read/write permissions
 - `vm_flags`: tells whether pages in this area are private
 - `vm_next`: points to next area `vm_area_struct` in list



exec() in More Detail

- Now we can see how exec works in detail!
- “Blowing away” the address space of the current process means freeing the `vm_area_structs` and page tables for old areas
- New `vm_area_structs` and page tables are created for new areas
 - Programs and initialized data are backed by object files
 - `.bss` and stack backed by anonymous files
- Set the program counter (PC) to the entry point in the `.text` section



Homework

Read Chapters 11 & 12

Next Lecture

We start looking at storage management