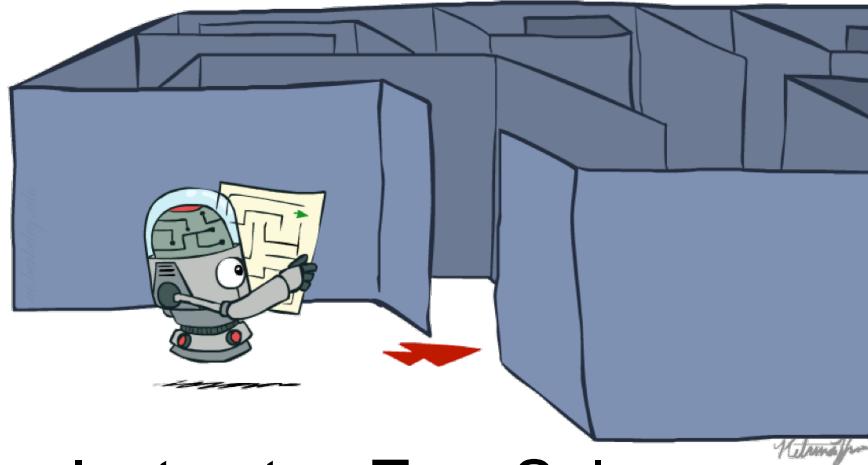


Announcements

- ❑ Project 0: Python Tutorial
 - Due September 7th
 - Submitted via Blackboard
- ❑ Homework 0: Math self-diagnostic
 - Due September 7th. Optional, but important to check your preparedness for second half
 - Submitted via Gradescope
- ❑ Pinned posts on piazza
- ❑ Reminder:
 - Sign up for Gradescope
 - Check Piazza for Announcement
 - Both Course website and Blackboard are for course content (slides, recordings, and other resources)
 - If you see a mismatch, please inform by a Piazza post or email
- ❑ I am still waiting for the TA assignment. Hopefully it will be announced by next week

CS 3568: Intelligent Systems

Search (Part 1)



Instructor: Tara Salman

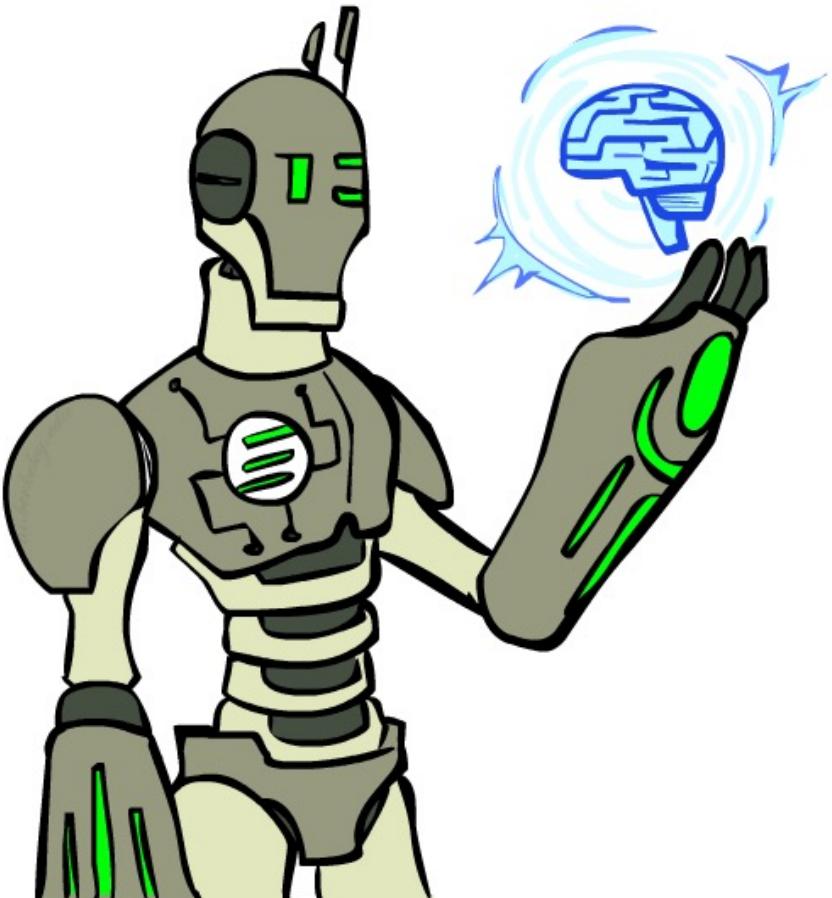
Texas Tech University

Computer Science Department

[These slides were created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley (ai.berkeley.edu).]

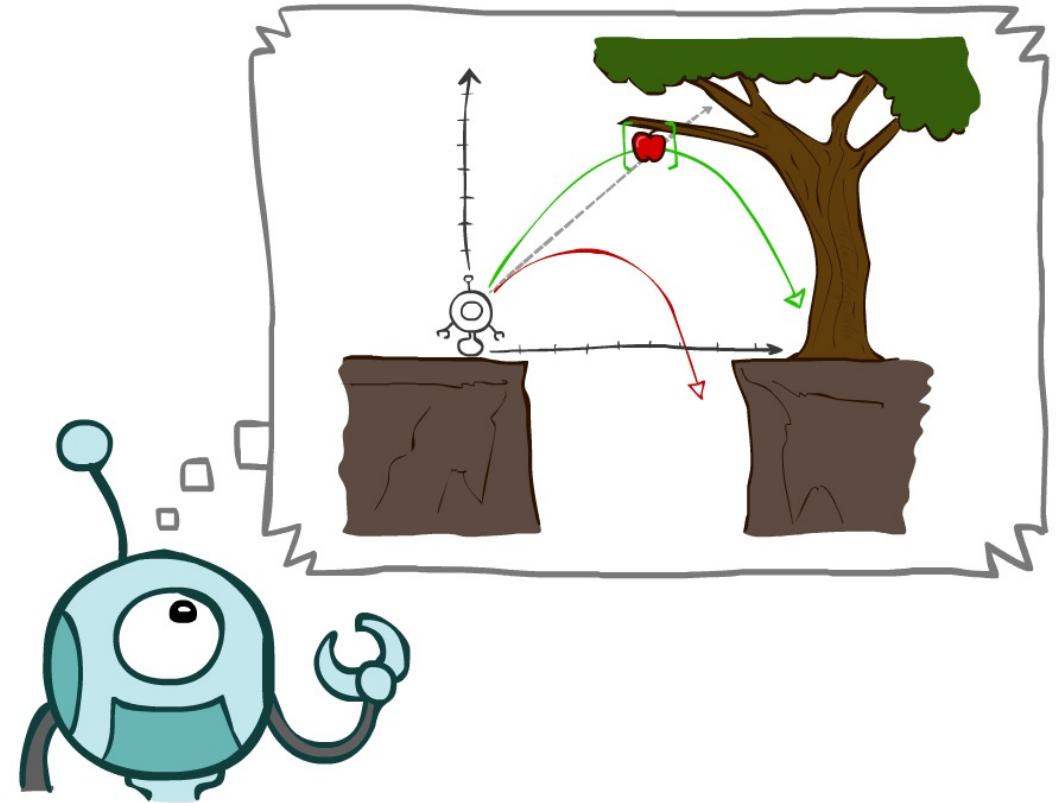
Last Lecture

- Course information
- What is AI?
- History of AI
- What can we currently AI do?
 - Different applications?
- What is this course?

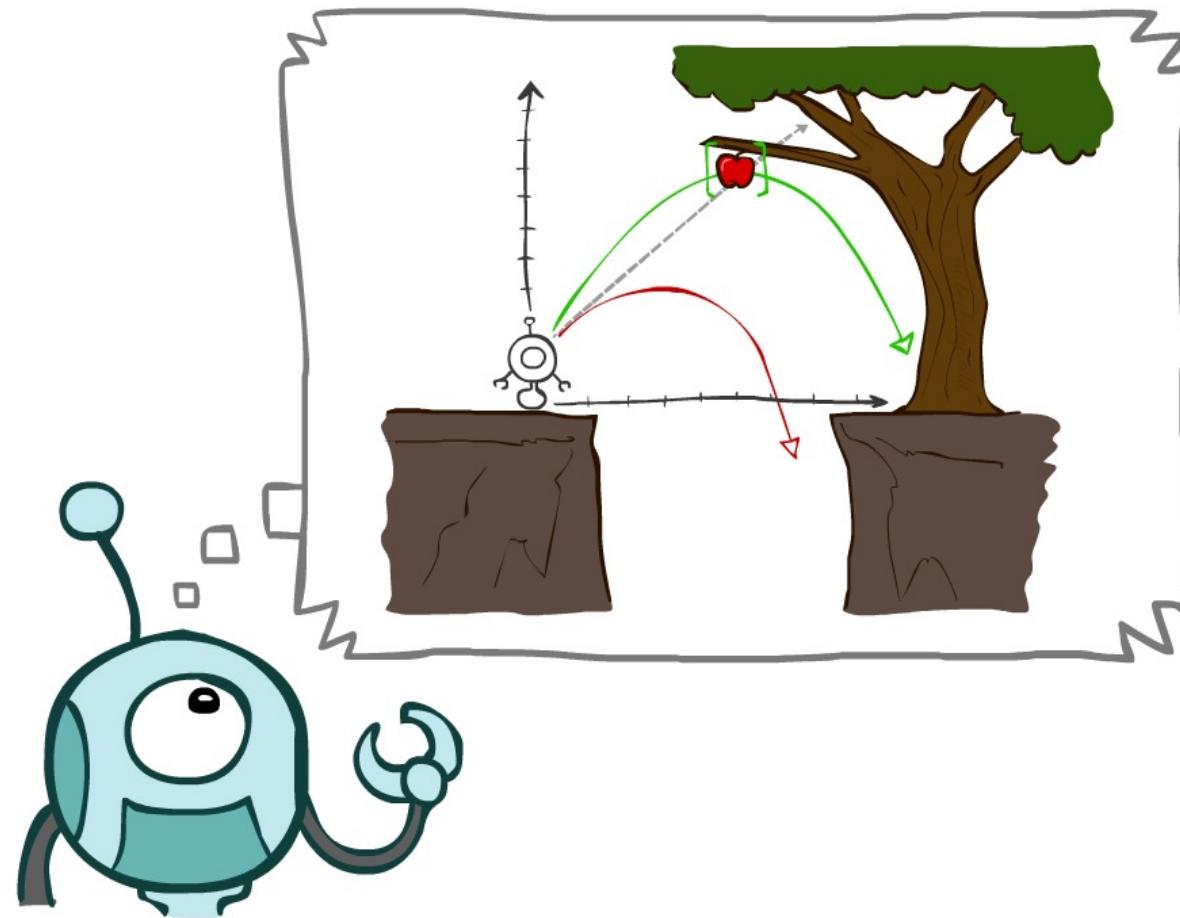


Today's Lecture

- ❑ Agents that Plan Ahead
- ❑ Search Problems
- ❑ Uninformed Search Methods
 - Depth-First Search
 - Breadth-First Search
 - Uniform-Cost Search

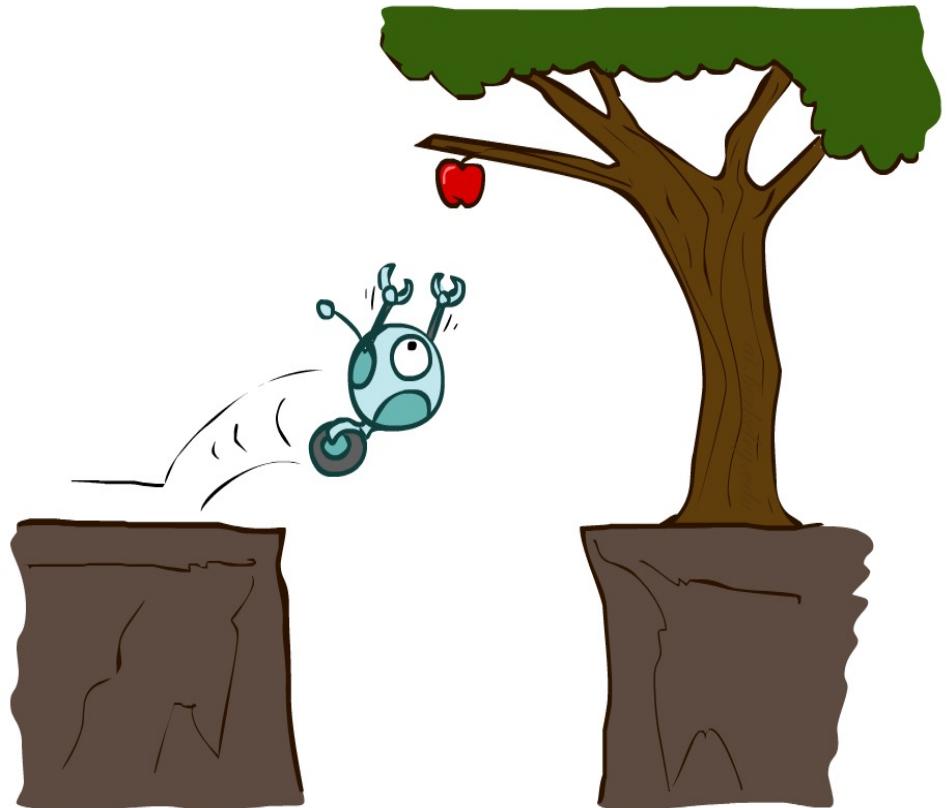


Agents that Plan



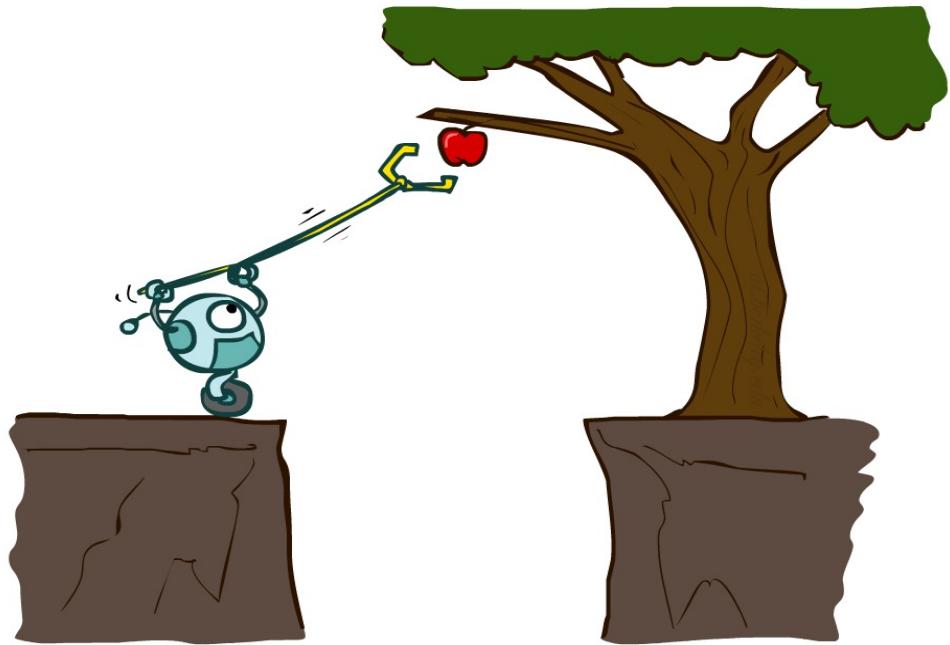
Reflex Agents

- ❑ Reflex agents:
 - Choose action based on current percept (and maybe memory)
 - May have memory or a model of the world's current state
 - Do not consider the future consequences of their actions
 - Consider how the world IS
- ❑ When do you need to be a reflex agent?
- ❑ Can a reflex agent be rational?



Planning Agents

- ❑ Planning agents:
 - Ask “what if”
 - Decisions based on (hypothesized) consequences of actions
 - Must have a model of how the world evolves in response to actions
 - Must formulate a goal (test)
 - Consider how the world **WOULD BE**
- ❑ Optimal vs. complete planning
 - Optimal: achieve goal with minimum co
 - Complete: If there is a solution, it can find it.
- ❑ Planning vs. replanning



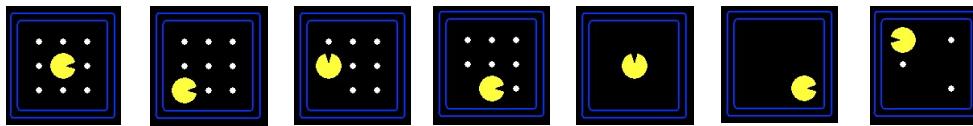
Search Problems



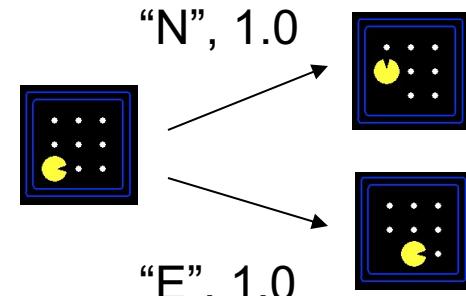
Search Problems

- ❑ A **search problem** consists of:

- A state space



- A successor function
(with actions, costs)



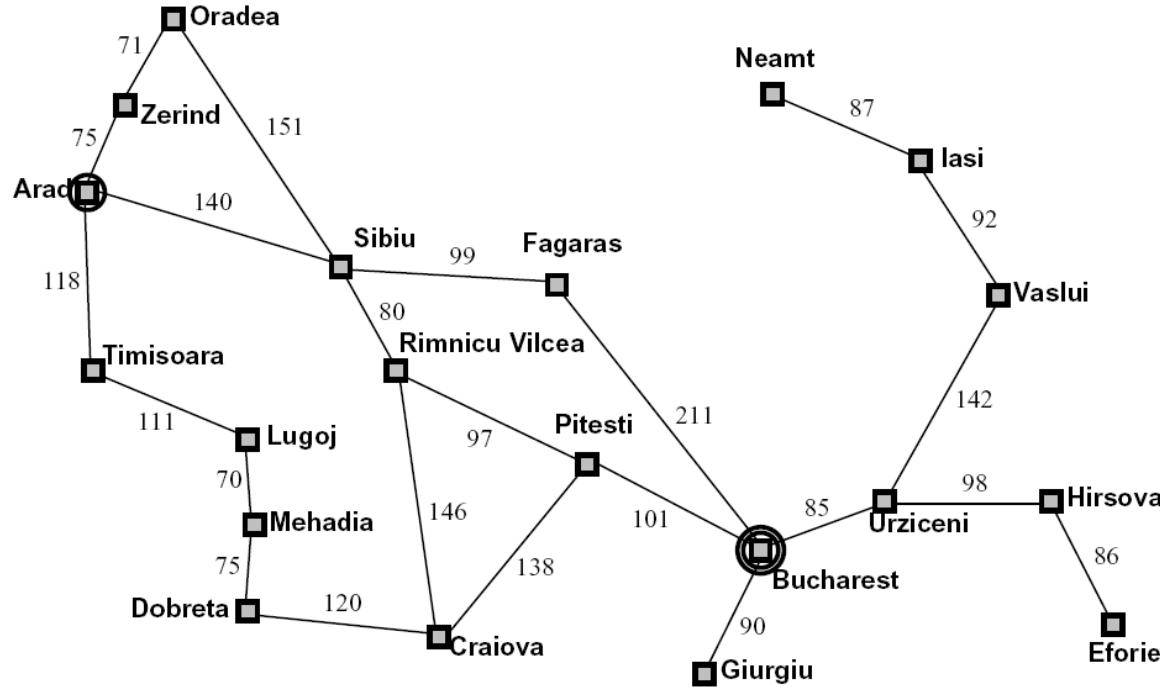
- A start state and a goal test

- ❑ A **solution** is a sequence of actions (a plan)
which transforms the start state to a goal state

Search Problems Are Models



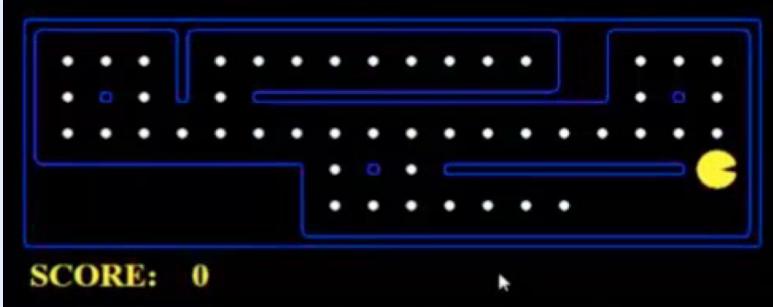
Example: Traveling in Romania



- ❑ State space:
 - Cities
- ❑ Successor function:
 - Roads: Go to adjacent city with cost = distance
- ❑ Start state:
 - Arad
- ❑ Goal test:
 - Is state == Bucharest?
- ❑ Solution?

What's in a State Space?

The **world state** includes every last detail of the environment



A **search state** keeps only the details needed for planning (abstraction)

❑ Problem: Pathing

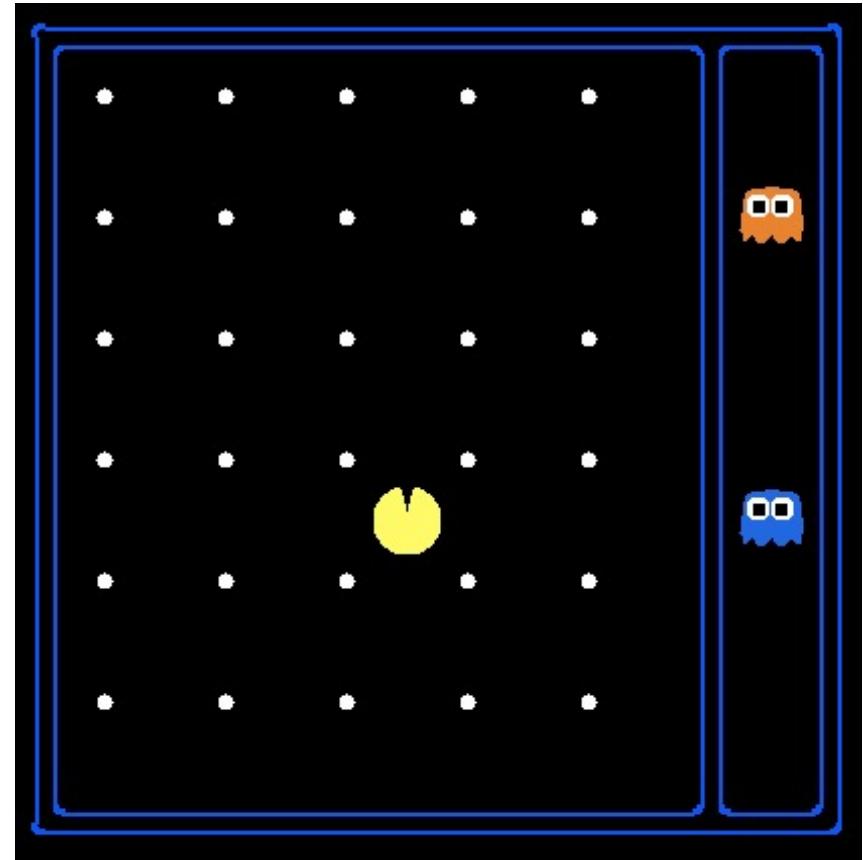
- States: (x,y) location
- Actions: NSEW
- Successor: update location only
- Goal test: is (x,y)=END

❑ Problem: Eat-All-Dots

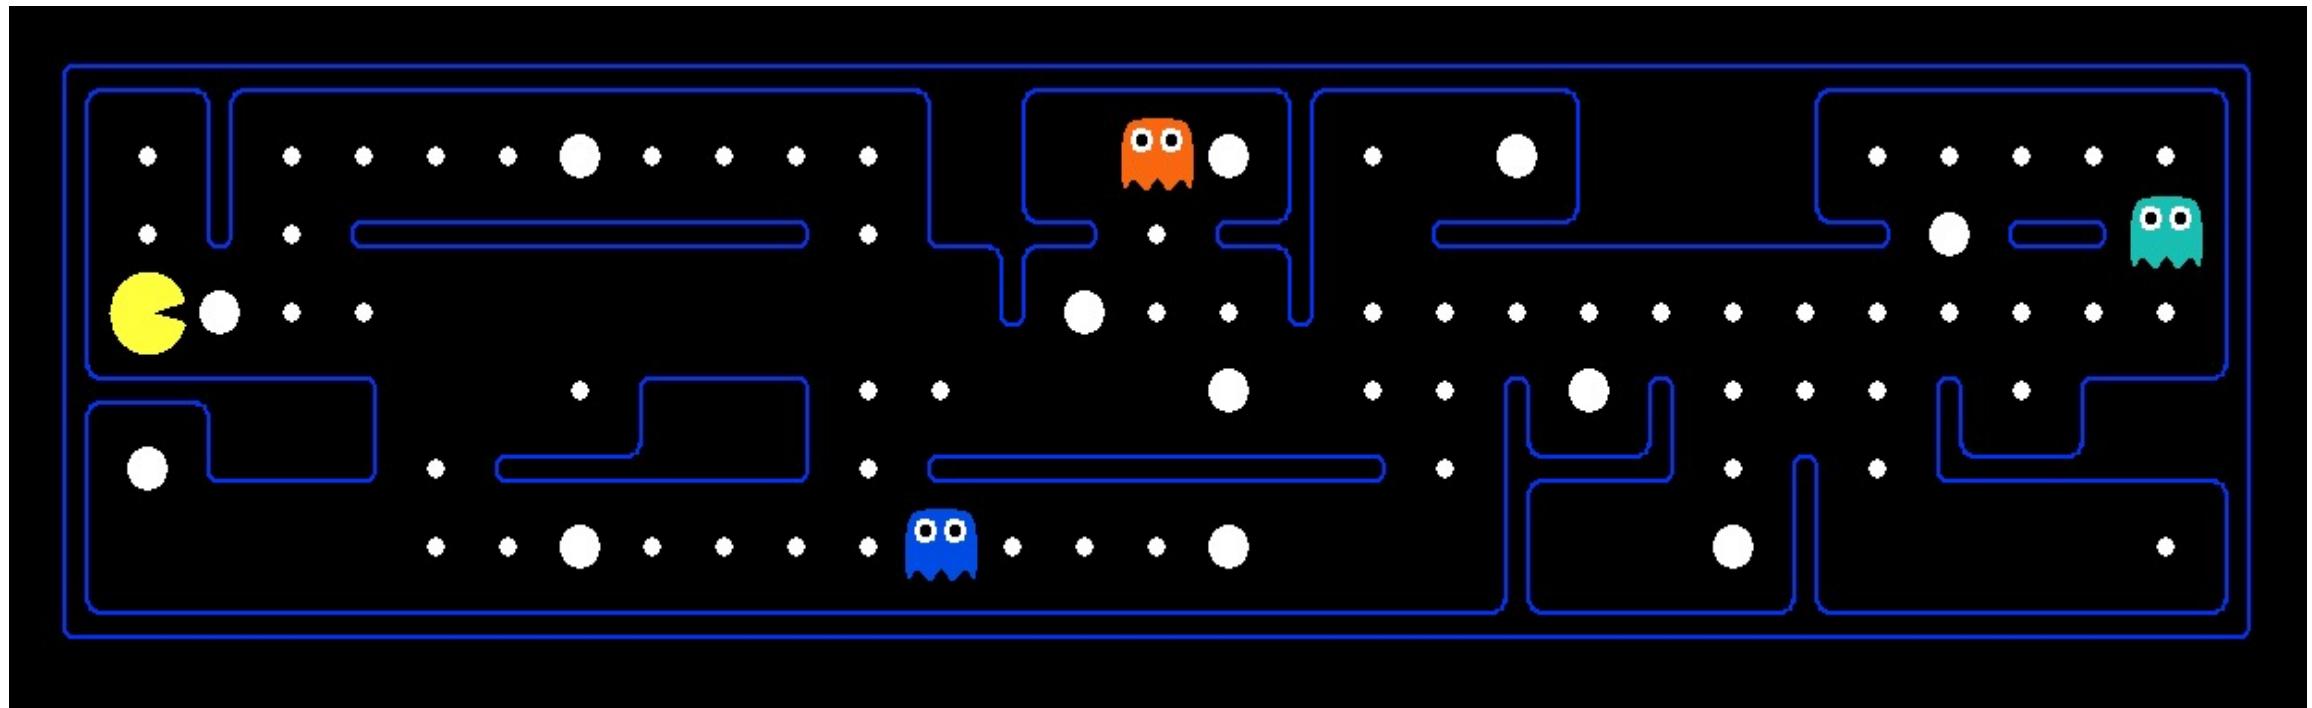
- States: {(x,y), dot booleans}
- Actions: NSEW
- Successor: update location and possibly a dot boolean
- Goal test: dots all false

State Space Sizes?

- ❑ World state:
 - Agent positions: 120
 - Food count: 30
 - Ghost positions: 12
 - Agent facing: NSEW
- ❑ How many
 - World states?
 $120 \times (2^{30}) \times (12^2) \times 4$
 - States for pathing?
120
 - States for eat-all-dots?
 $120 \times (2^{30})$

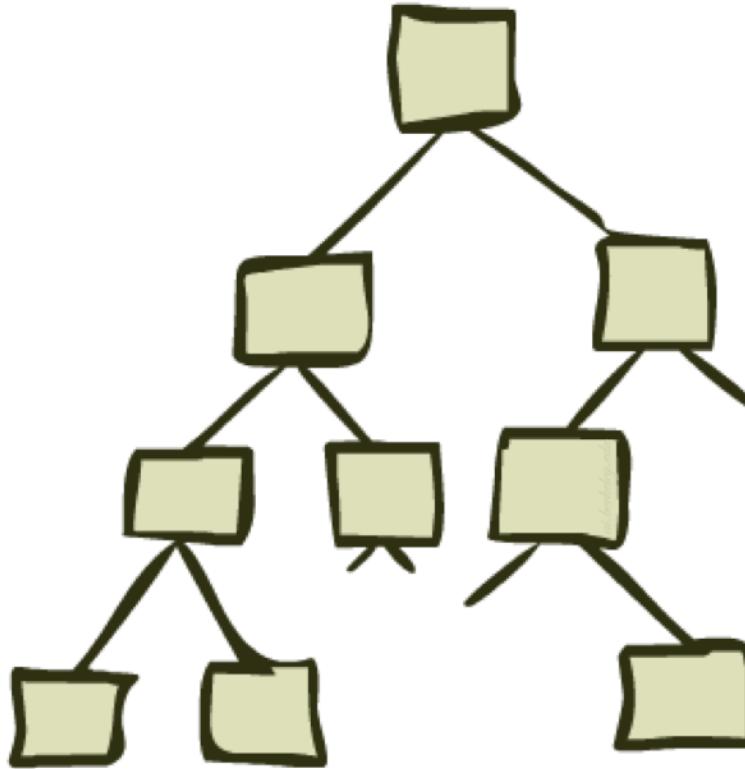


Quiz: Safe Passage



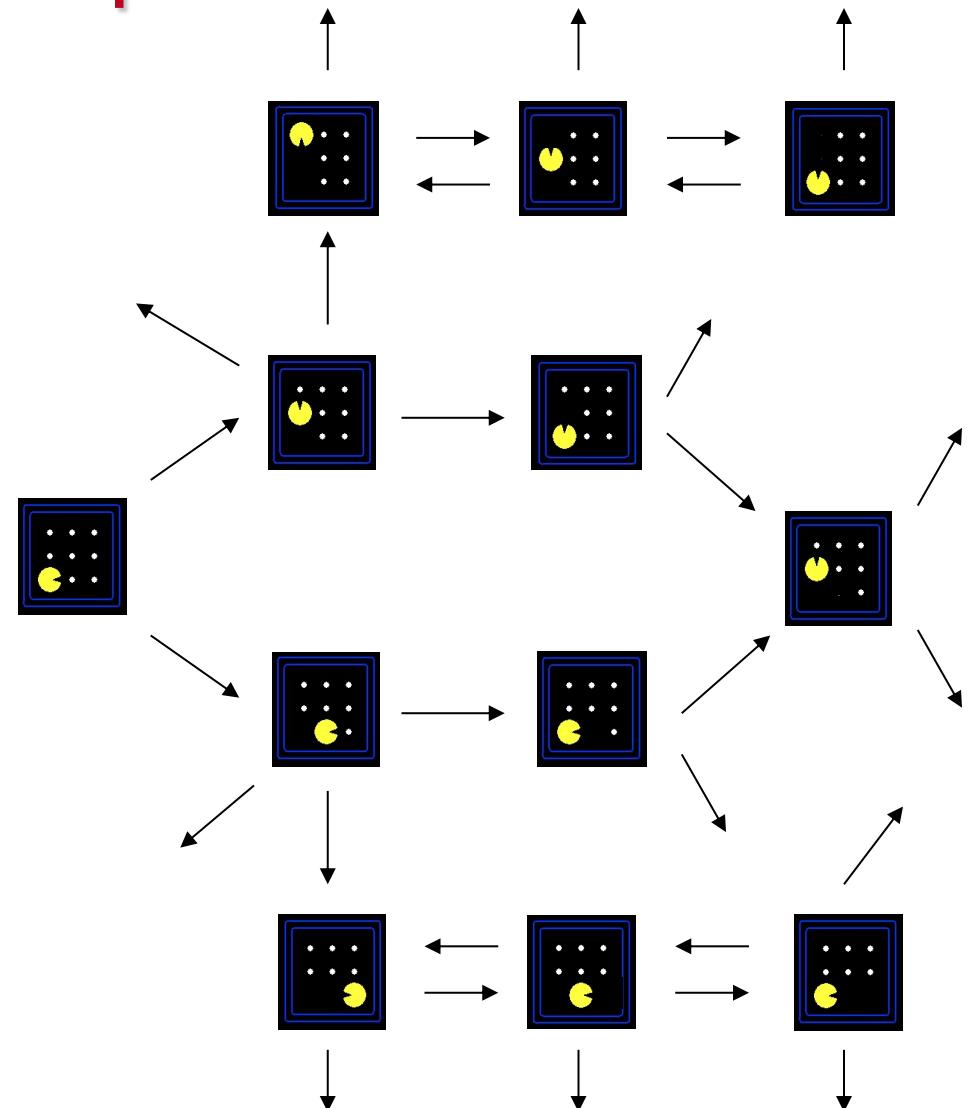
- ❑ Problem: eat all dots while keeping the ghosts perma-scared
- ❑ What does the state space have to specify?
 - (agent position, dot booleans, power pellet booleans, remaining scared time)

State Space Graphs and Search Trees



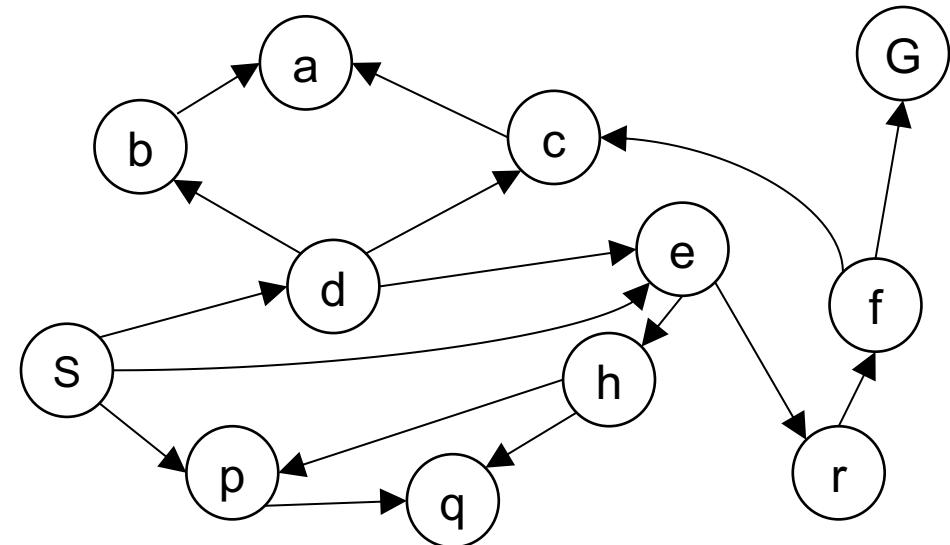
State Space Graphs

- ❑ State space graph: A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal nodes (maybe only one)
- ❑ In a state space graph, each state occurs only once!
- ❑ We can rarely build this full graph in memory (it's too big), but it's a useful idea



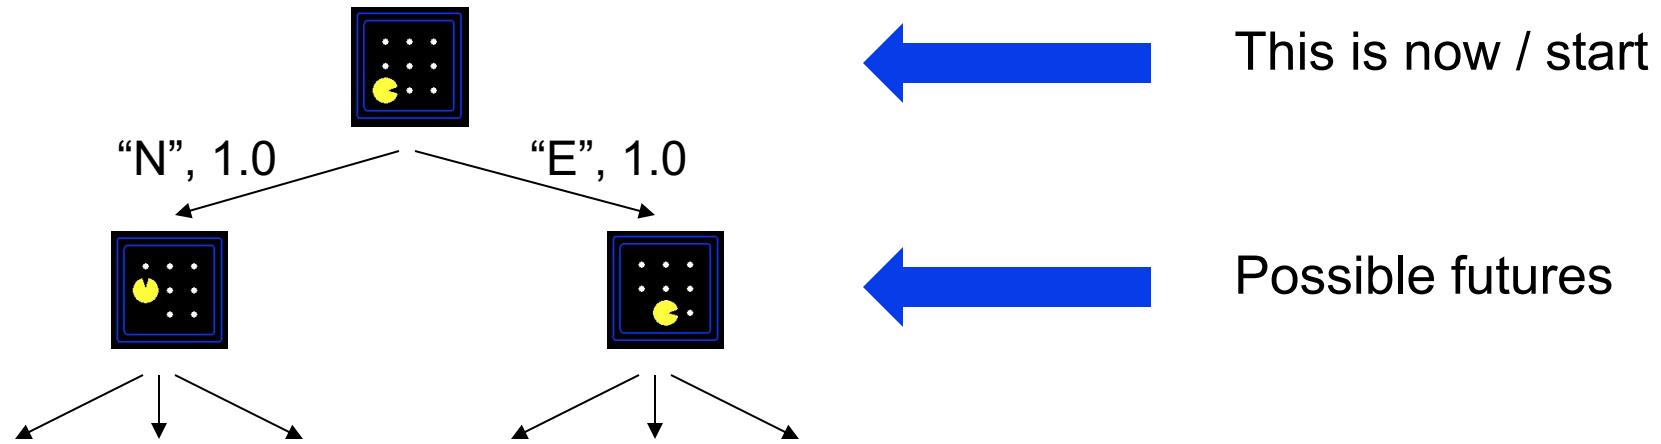
State Space Graphs

- ❑ State space graph: A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal nodes (maybe only one)
- ❑ In a state space graph, each state occurs only once!
- ❑ We can rarely build this full graph in memory (it's too big), but it's a useful idea



Tiny state space graph for a tiny search problem

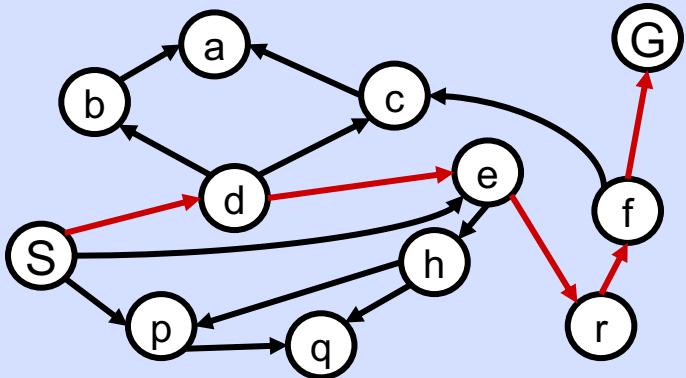
Search Trees



- A search tree:
 - A “what if” tree of plans and their outcomes
 - The start state is the root node
 - Children correspond to successors
 - Nodes show states, but correspond to PLANS that achieve those states
 - **For most problems, we can never actually build the whole tree**

State Space Graphs vs. Search Trees

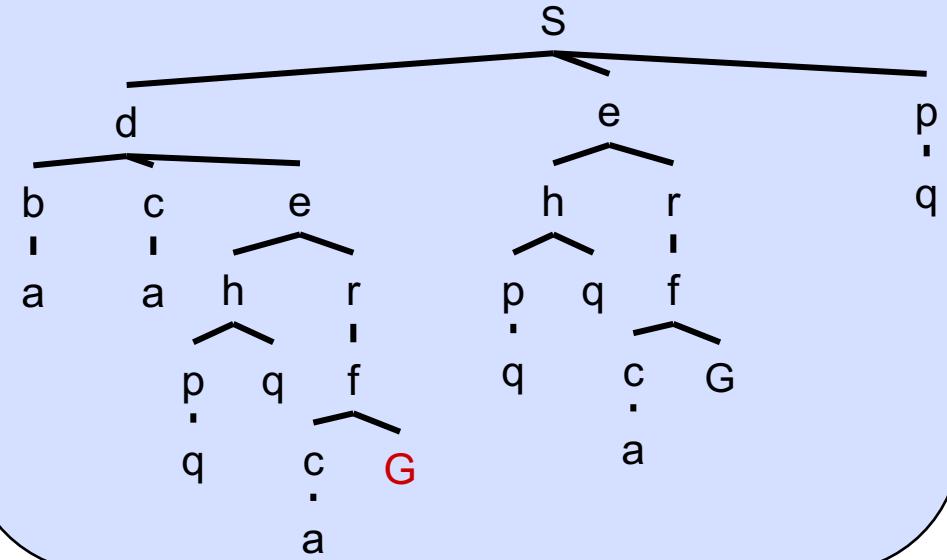
State Space Graph



Each NODE in in the search tree is an entire PATH in the state space graph.

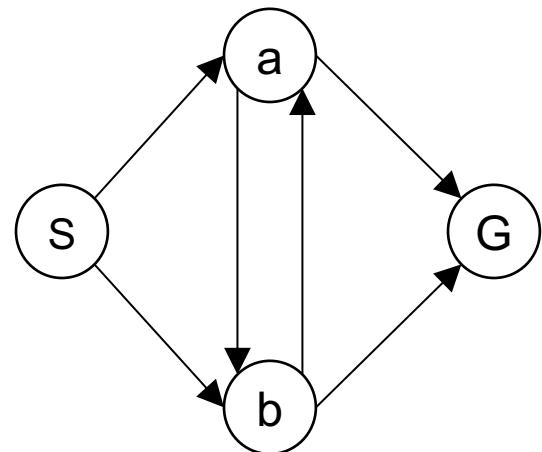
We construct both on demand – and we construct as little as possible.

Search Tree



Quiz: State Space Graphs vs. Search Trees

Consider this 4-state graph:

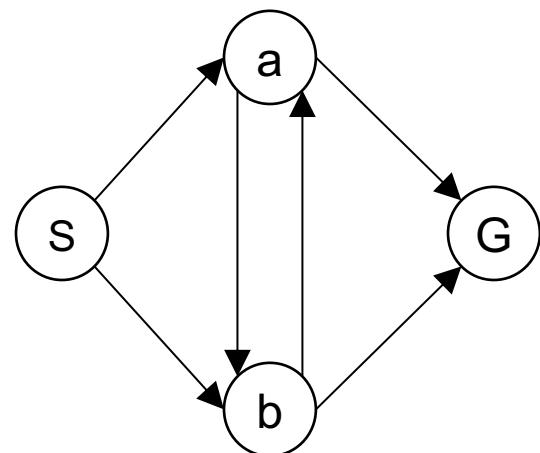


How big is its search tree (from S)?

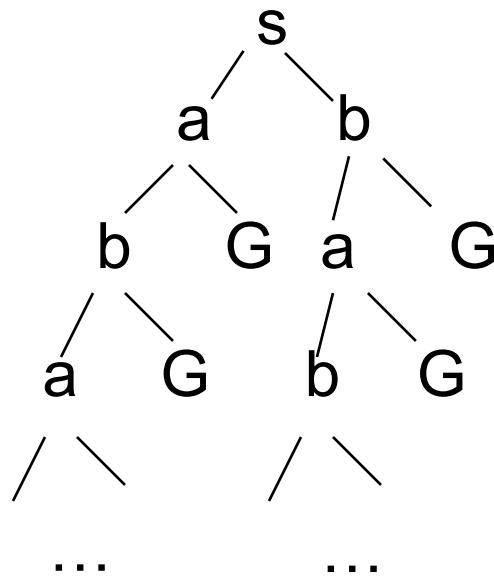


Quiz: State Space Graphs vs. Search Trees

Consider this 4-state graph:

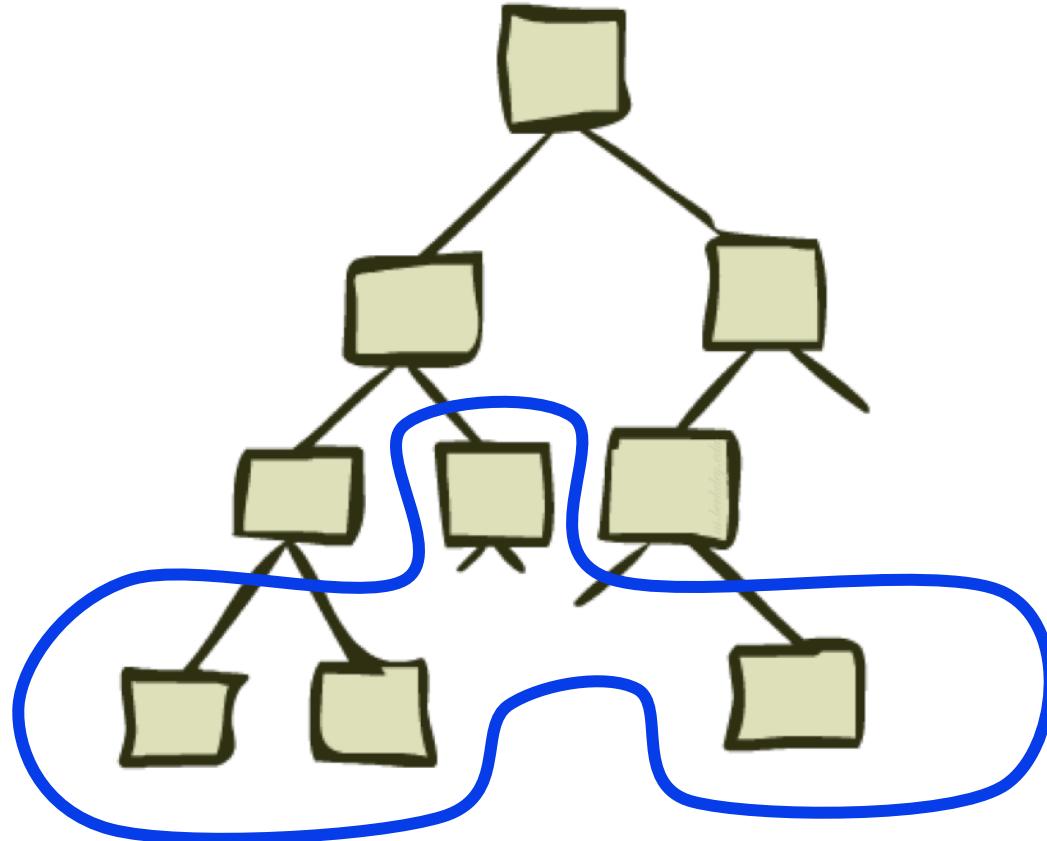


How big is its search tree (from S)?

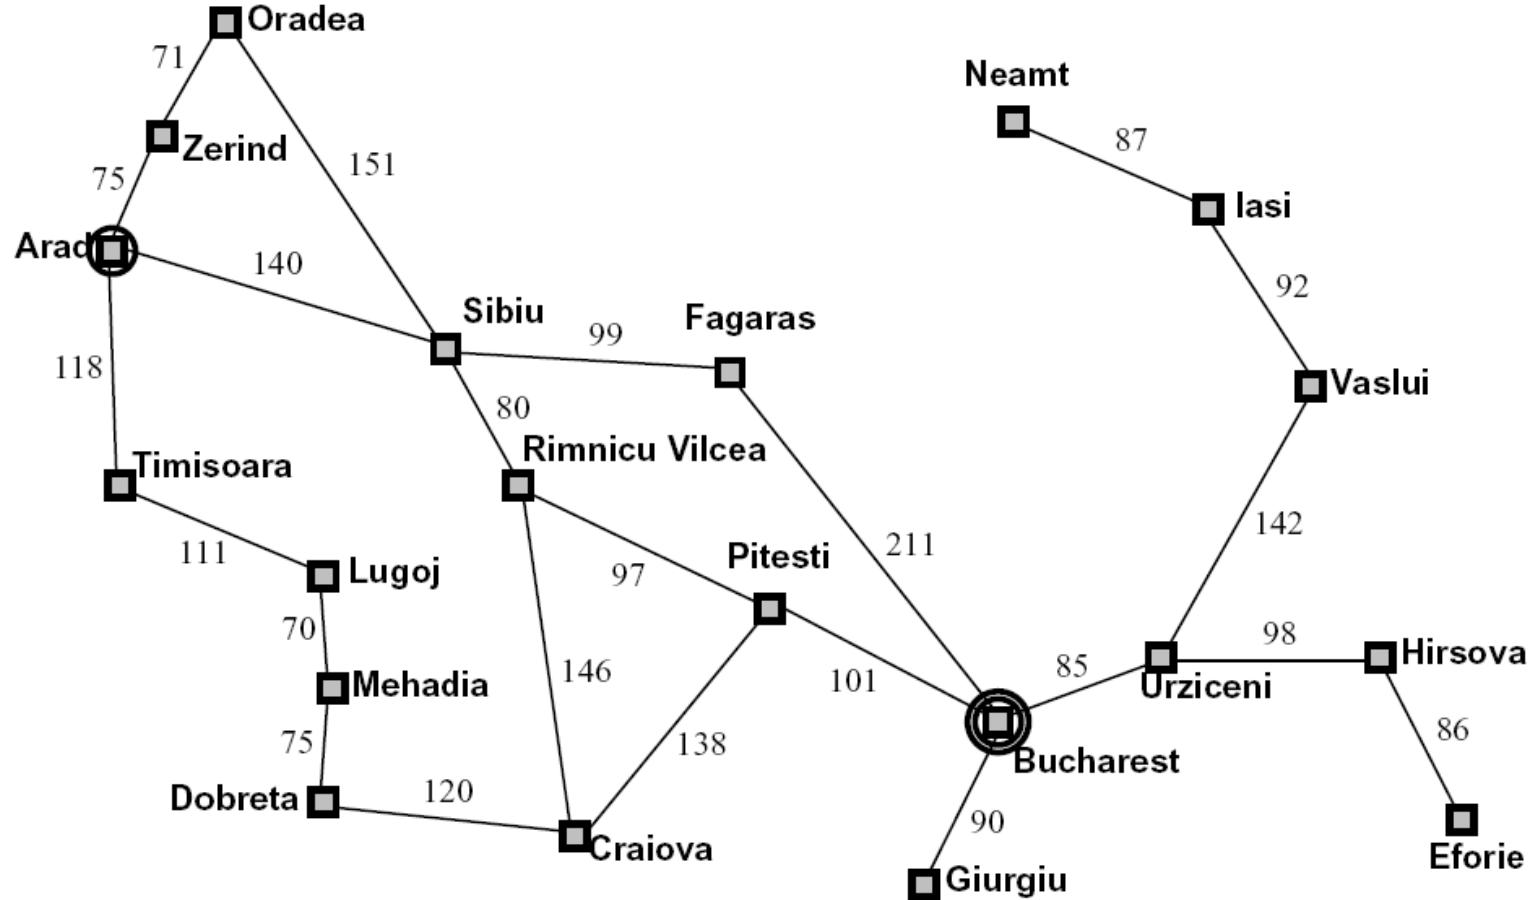


Important: Lots of repeated structure in the search tree!

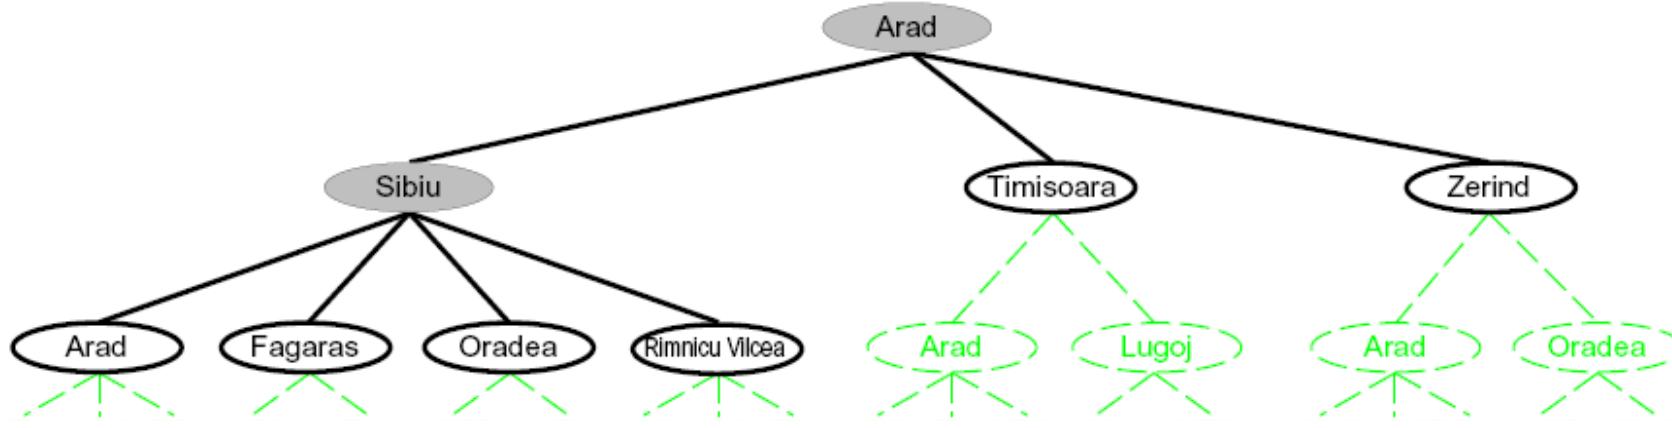
Tree Search



Search Example: Romania



Searching with a Search Tree



- ❑ Search:

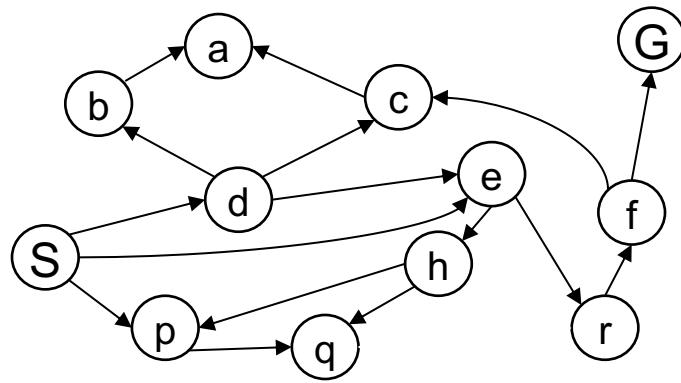
- Expand out potential plans (tree nodes)
- Maintain a **fringe** of partial plans under consideration
- Try to expand as few tree nodes as possible

General Tree Search

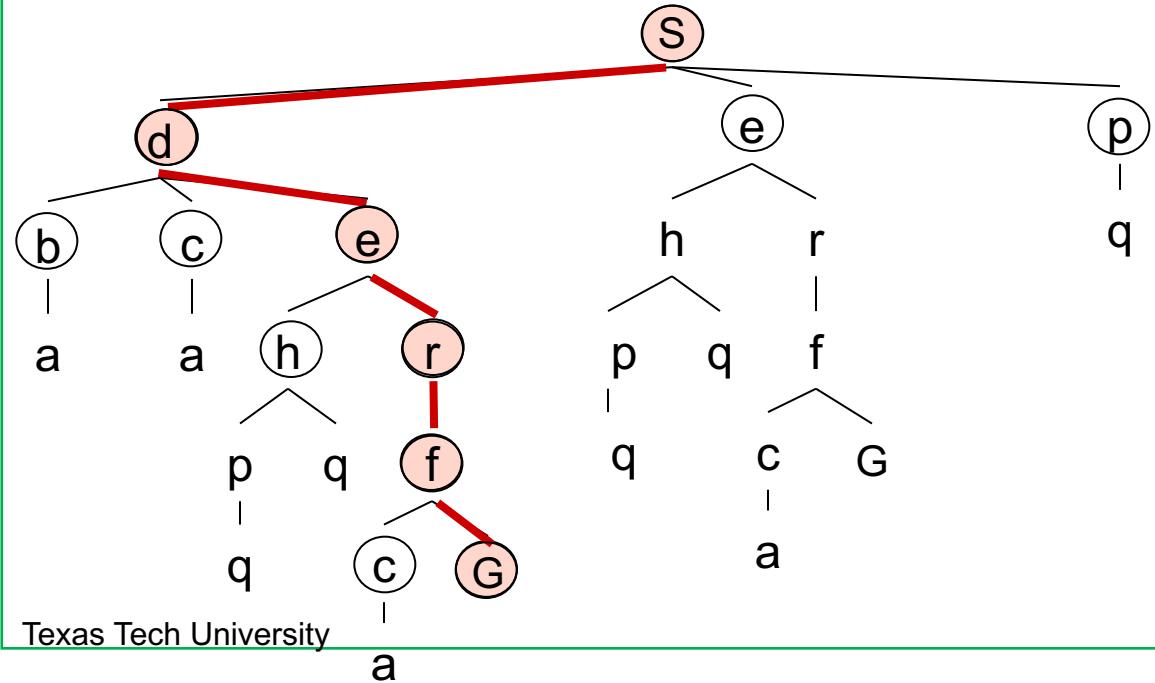
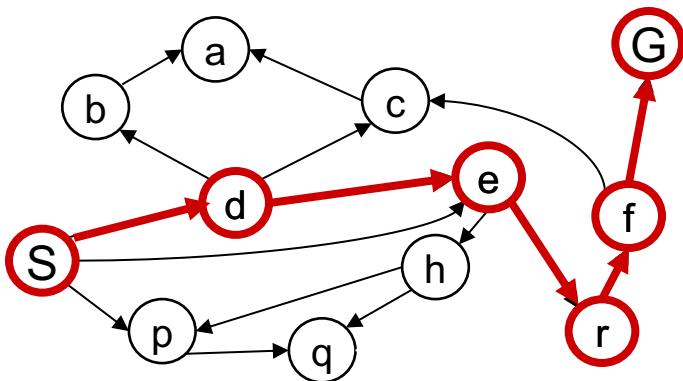
```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
    initialize the search tree using the initial state of problem
    loop do
        if there are no candidates for expansion then return failure
        choose a leaf node for expansion according to strategy
        if the node contains a goal state then return the corresponding solution
        else expand the node and add the resulting nodes to the search tree
    end
```

- Important ideas:
 - Fringe
 - Expansion
 - Exploration strategy
- Main question: which fringe nodes to explore?

Example: Tree Search



Example: Tree Search



s
s → d
s → e
s → p
s → d → b
s → d → c
s → d → e
s → d → e → h
s → d → e → r
s → d → e → r → f
s → d → e → r → f → c
s → d → e → r → f → G

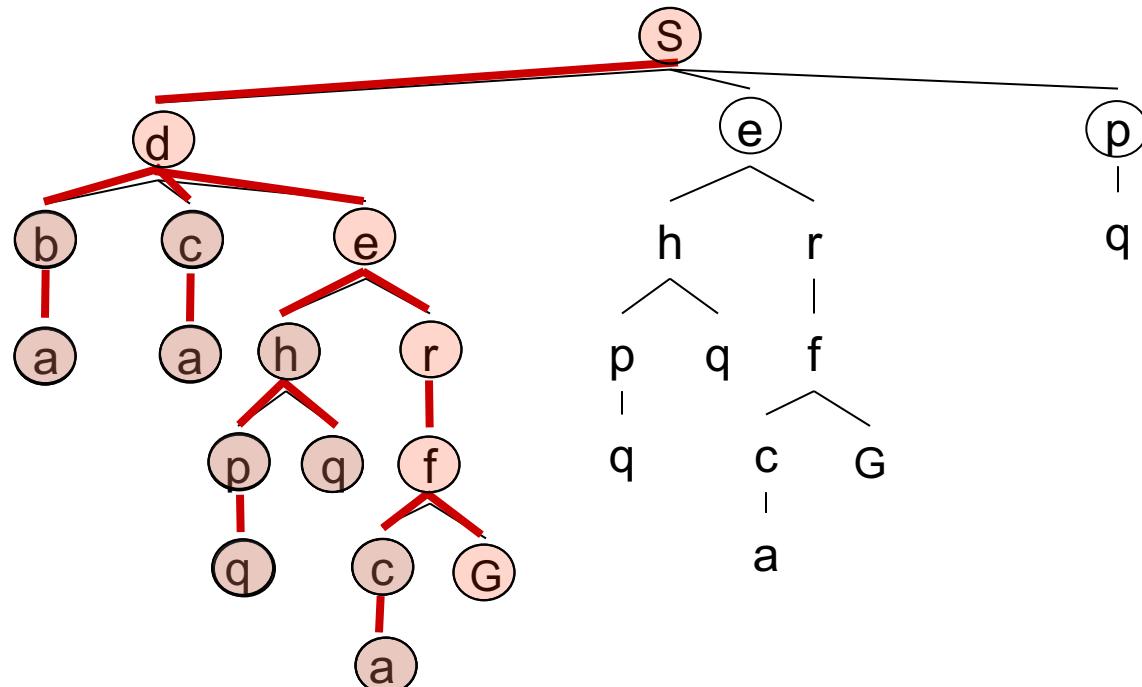
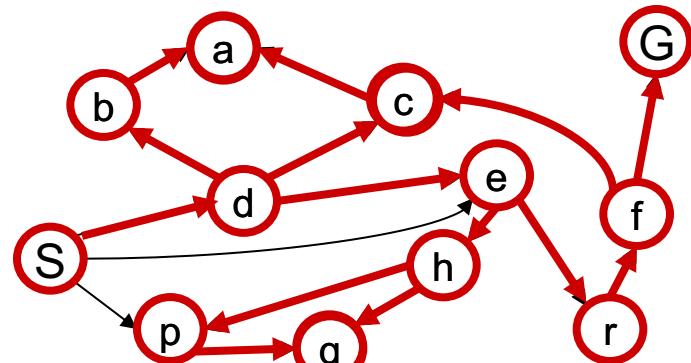
Depth-First Search



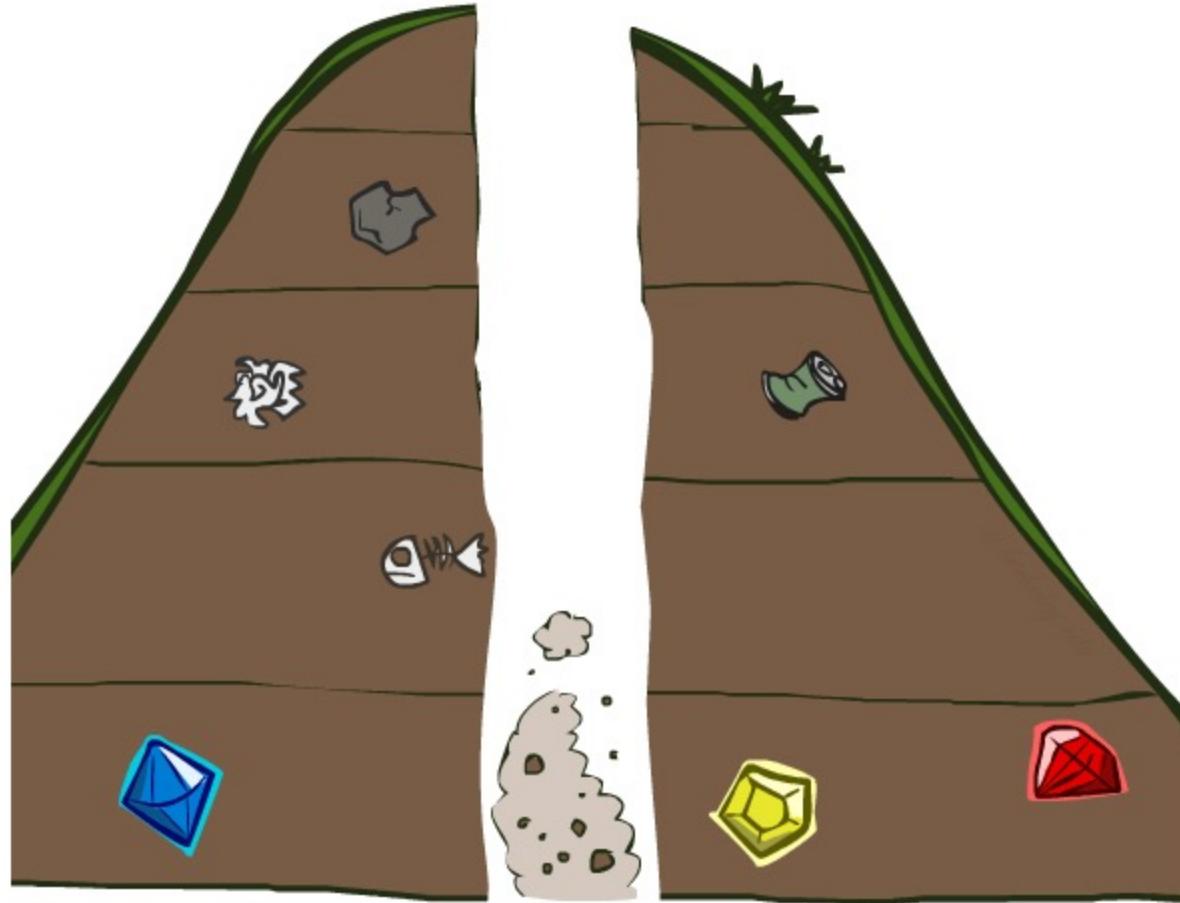
Depth-First Search

Strategy: expand a deepest node first

Implementation: Fringe is a LIFO stack

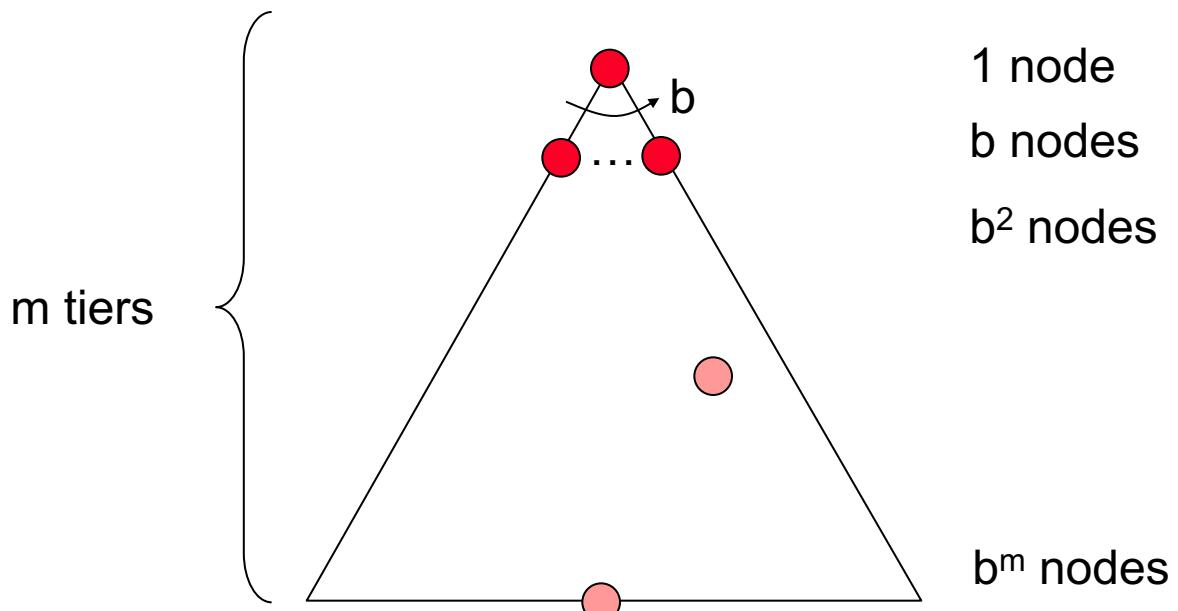


Search Algorithm Properties



Search Algorithm Properties

- ❑ Complete: Guaranteed to find a solution if one exists?
- ❑ Optimal: Guaranteed to find the least cost path?
- ❑ Time complexity?
- ❑ Space complexity?
- ❑ Cartoon of search tree:
 - b is the branching factor
 - m is the maximum depth
 - solutions at various depths
- ❑ Number of nodes in entire tree?
 - $1 + b + b^2 + \dots + b^m = O(b^m)$



Depth-First Search (DFS) Properties

□ What nodes DFS expand?

- Some left prefix of the tree.
- Could process the whole tree!
- If m is finite, takes time $O(b^m)$

□ How much space does the fringe take? m tiers

- Only has siblings on path to root, so $O(bm)$

□ Is it complete?

- m could be infinite, so only if we prevent cycles
(more later)

□ Is it optimal?

- No, it finds the “leftmost” solution, regardless of depth or cost

