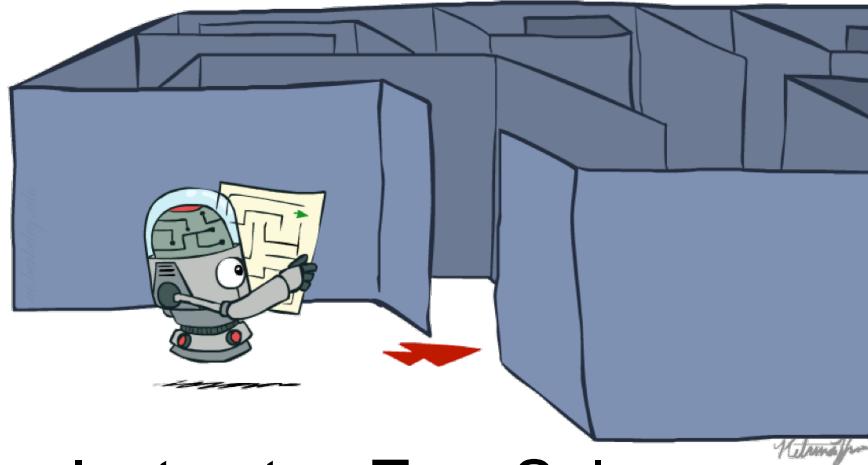


# Announcements

- ❑ For those who have recently joined
  - Welcome!
  - Last week lectures are recorded and can be viewed online  
(Blackboard or course webpage)
- ❑ TA: Chang Liu
  - Office Hour: Monday 3-5 pm EC202
- ❑ Homework 0 and Project 0
  - Have been released!
  - Project 0 submission via blackboard
  - Homework 0 via gradescope
  - Due **Tuesday, Sep 7th, at 9:30am**

# CS 3568: Intelligent Systems

## Search (Part 2)



Instructor: Tara Salman

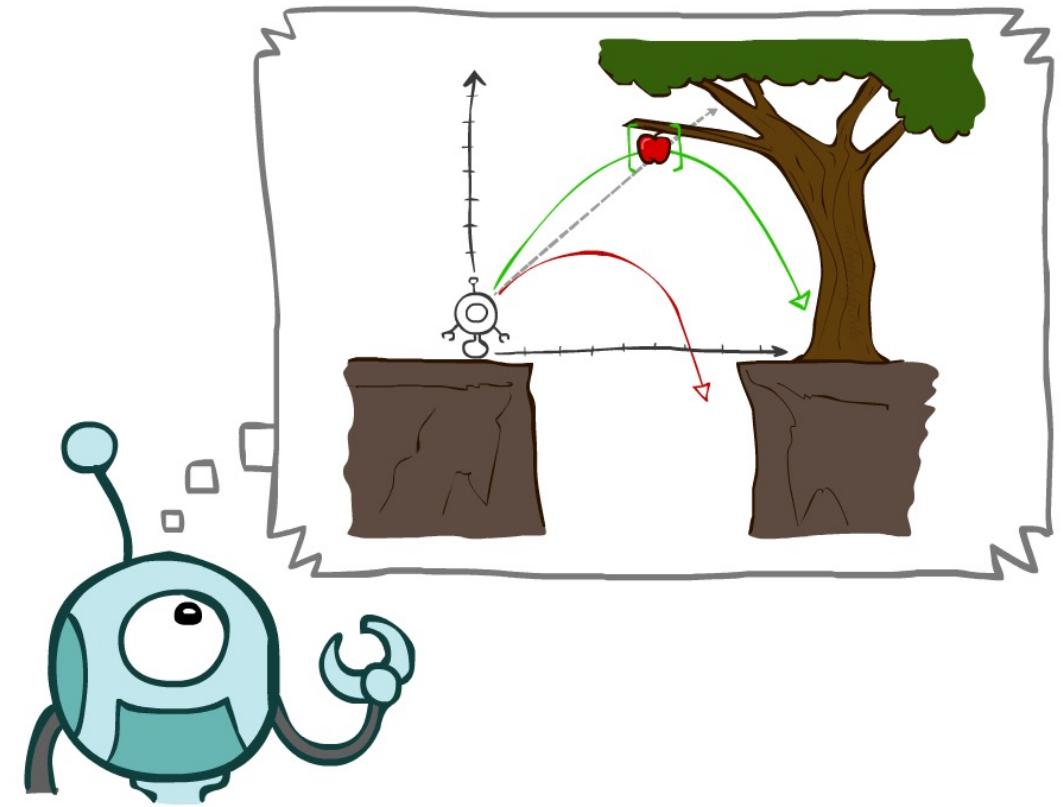
Texas Tech University

Computer Science Department

[These slides were created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley ([ai.berkeley.edu](http://ai.berkeley.edu)).]

# Last Lecture

- ❑ Agents that Plan Ahead
- ❑ Search Problems
- ❑ Uninformed Search Methods
  - Depth-First Search



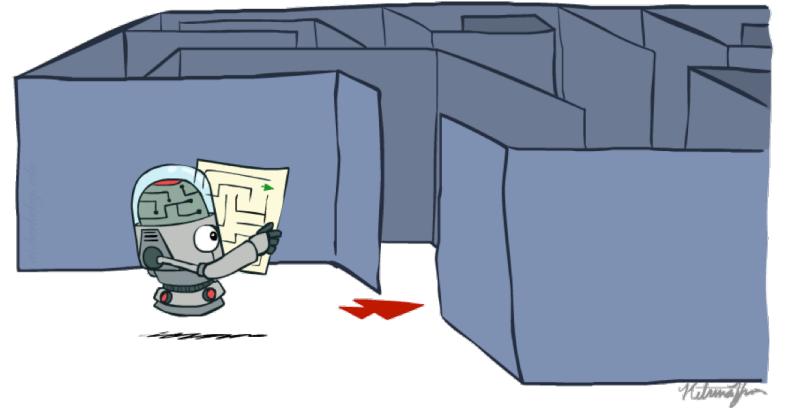
# Today

- ❑ Uniformed Search
  - BFS
  - Cost Sensitive search
  
- ❑ Informed Search
  - Heuristics
  - Greedy Search
  - A\* Search



# Recap: Search

- Search problem:
  - States (configurations of the world)
  - Actions and costs
  - Successor function (world dynamics)
  - Start state and goal test
- Search tree:
  - Nodes: represent plans for reaching states
  - Plans have costs (sum of action costs)
- Search algorithm:
  - Systematically builds a search tree
  - Chooses an ordering of the fringe (unexplored nodes)



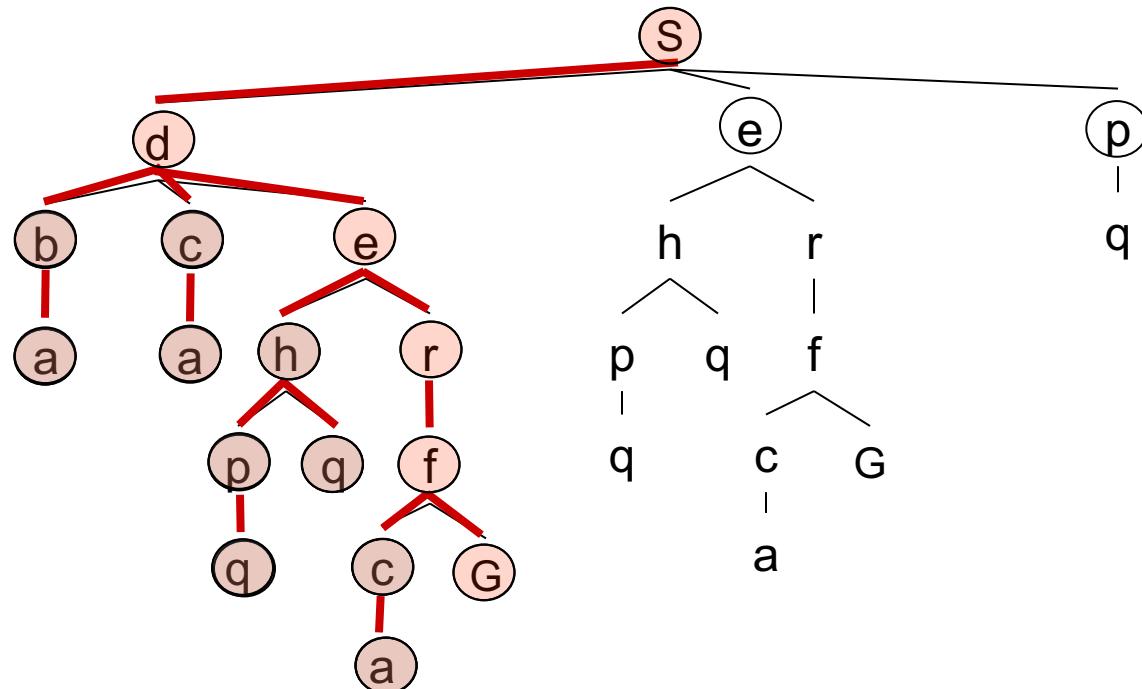
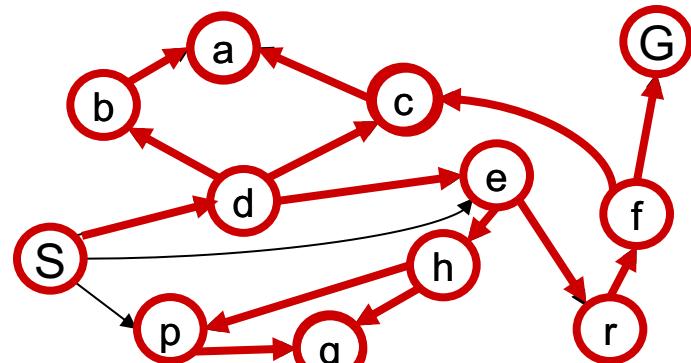
# Depth-First Search



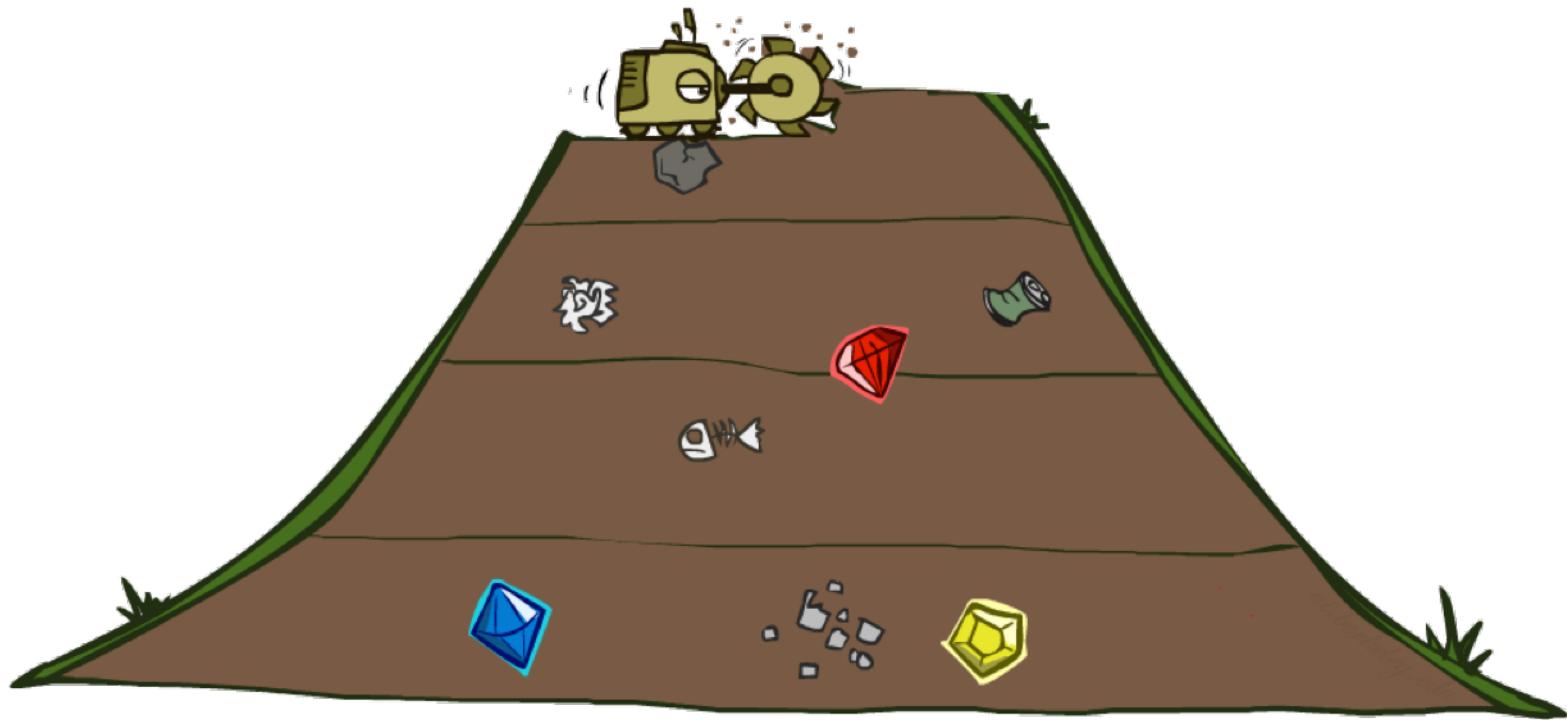
# Depth-First Search

Strategy: expand a deepest node first

Implementation: Fringe is a LIFO stack



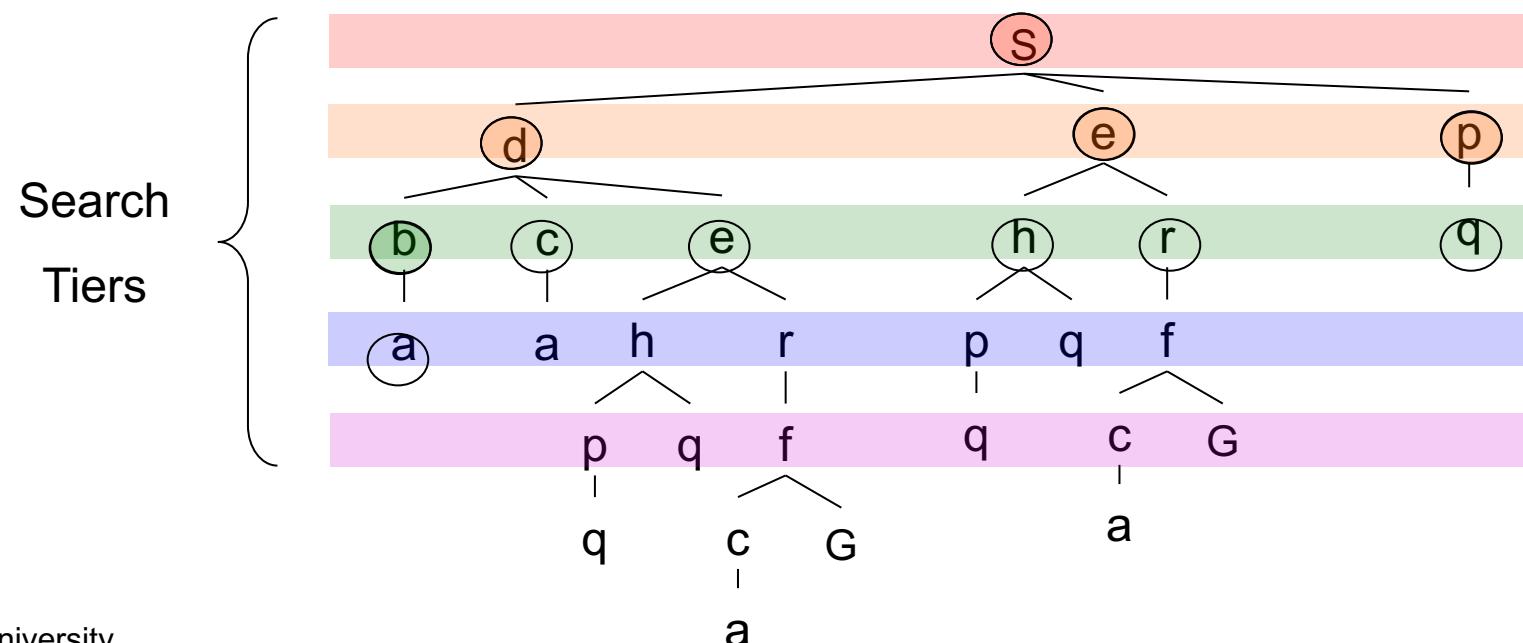
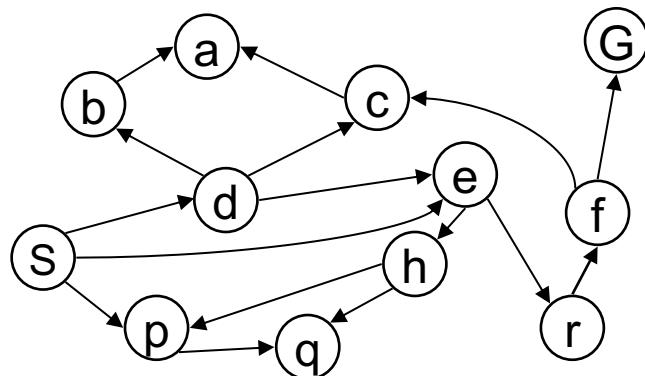
# Breadth-First Search



# Breadth-First Search

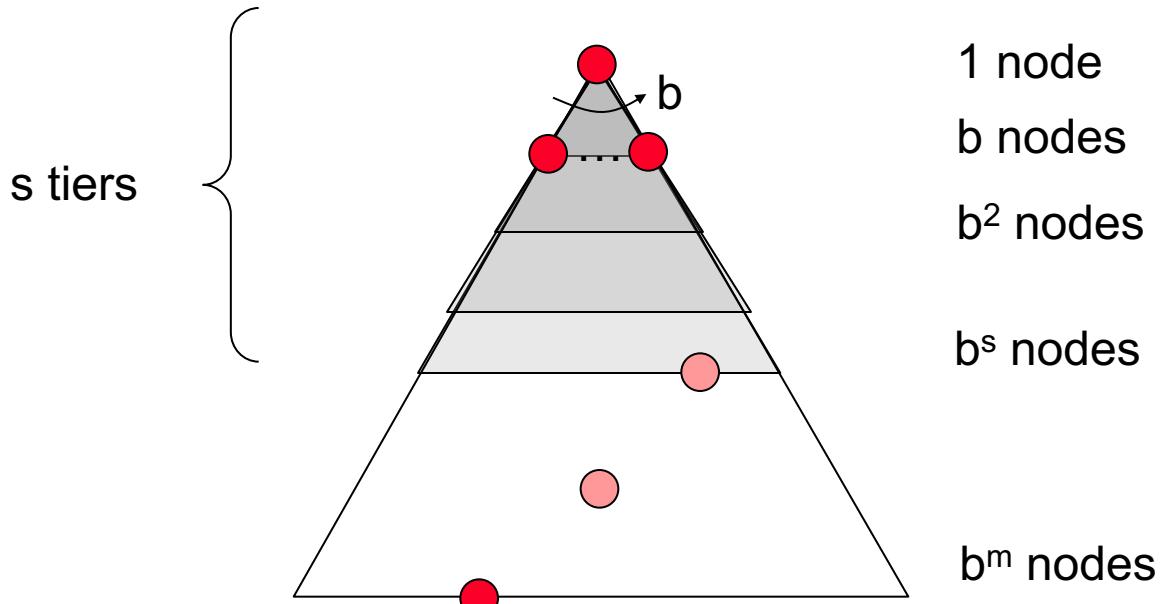
Strategy: expand a shallowest node first

Implementation: Fringe is a FIFO queue

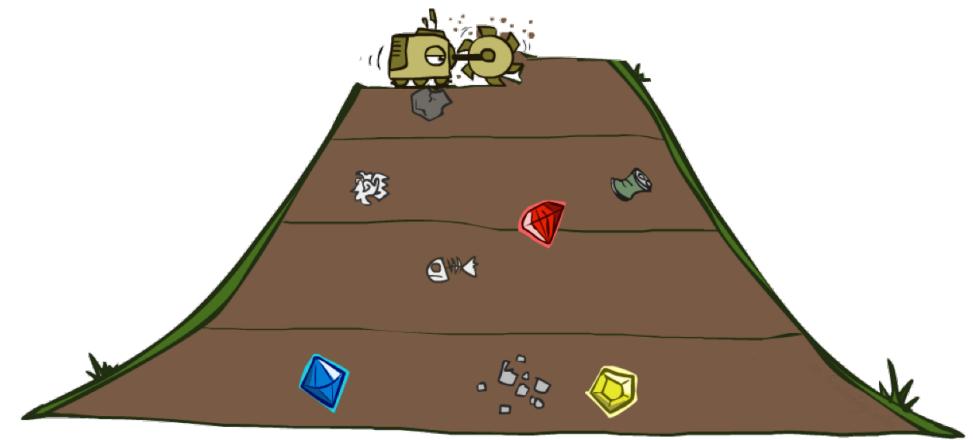


# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Processes all nodes above shallowest solution
  - Let depth of shallowest solution be  $s$
  - Search takes time  $O(b^s)$
  
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^s)$
  
- Is it complete?
  - $s$  must be finite if a solution exists, so yes!
  
- Is it optimal?
  - Only if costs are all 1 (more on costs later)



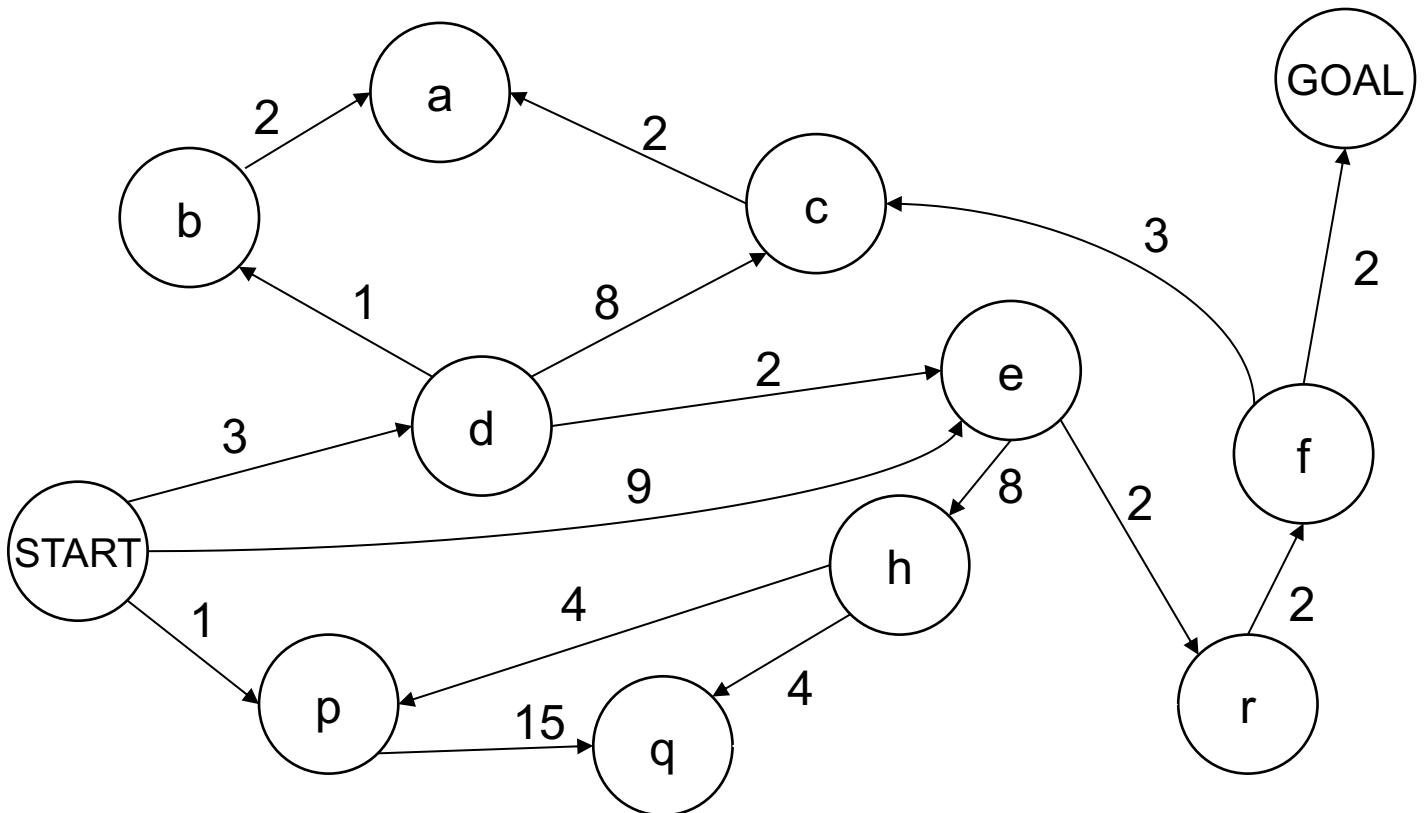
# Quiz: DFS vs BFS



# Quiz: DFS vs BFS

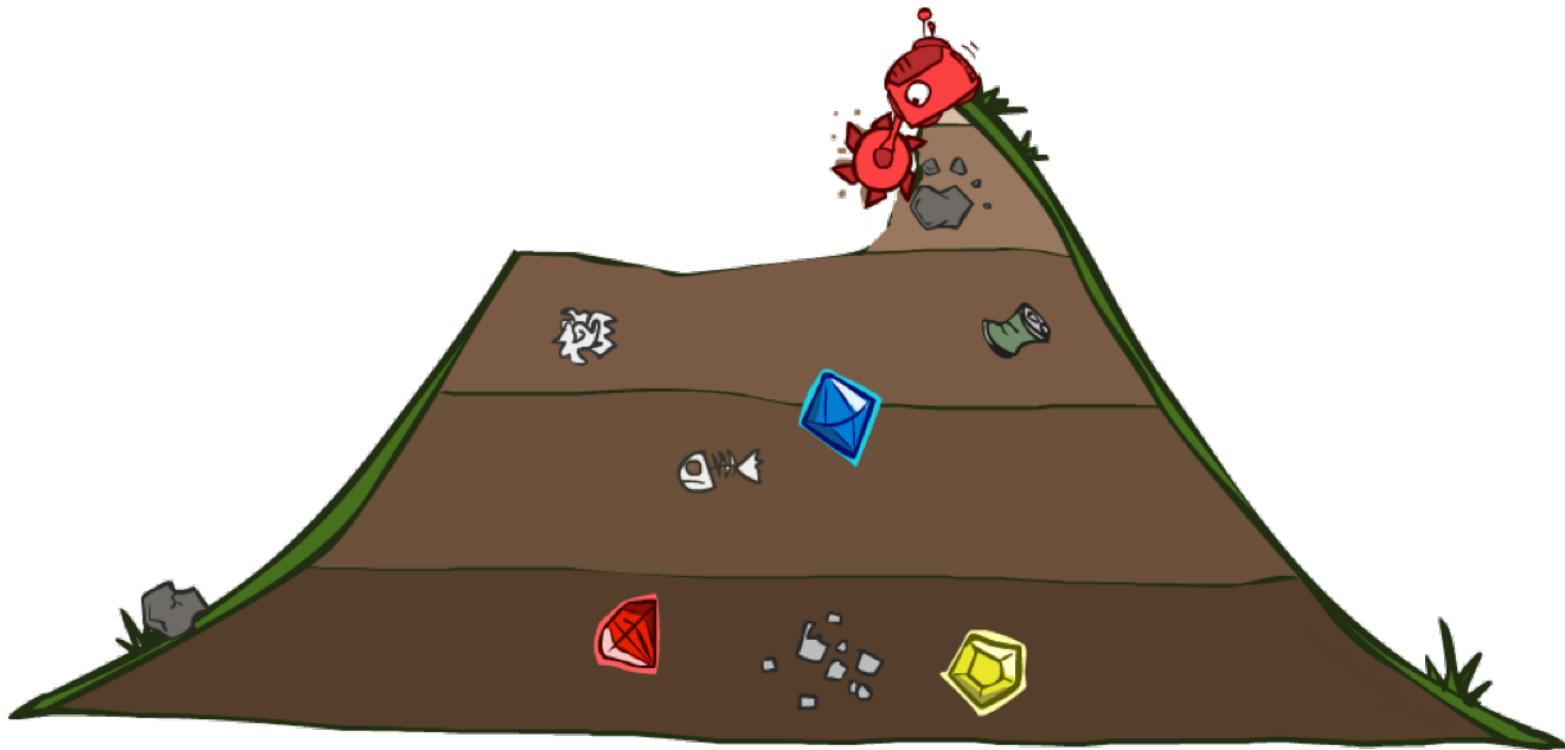
- When will BFS outperform DFS?
  
- When will DFS outperform BFS?

# Cost-Sensitive Search



BFS finds the shortest path in terms of number of actions.  
It does not find the least-cost path. We will now cover  
a similar algorithm which does find the least-cost path.

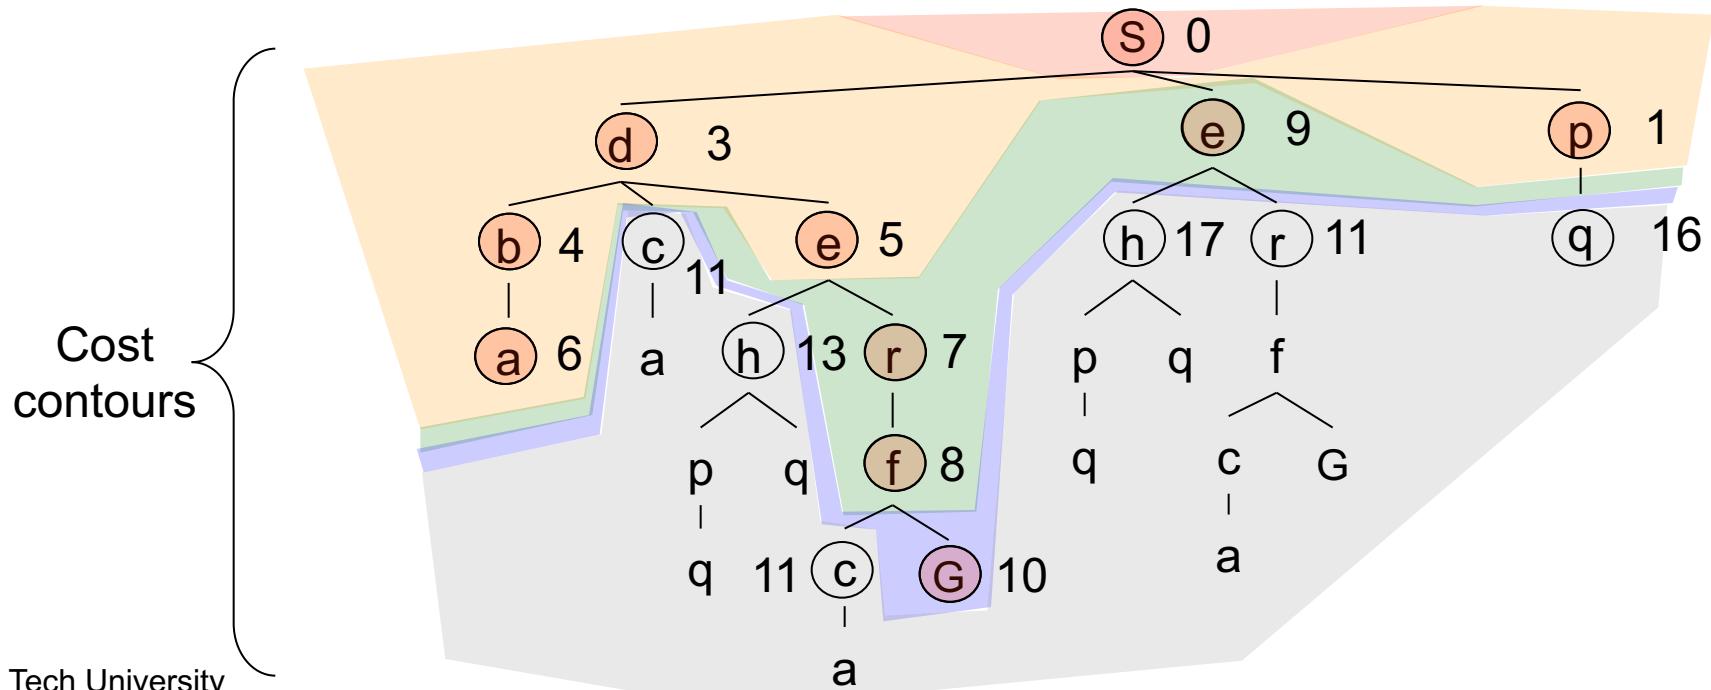
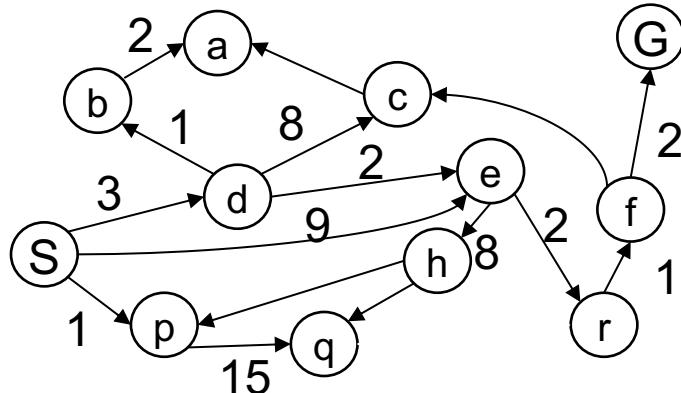
# Uniform Cost Search



# Uniform Cost Search

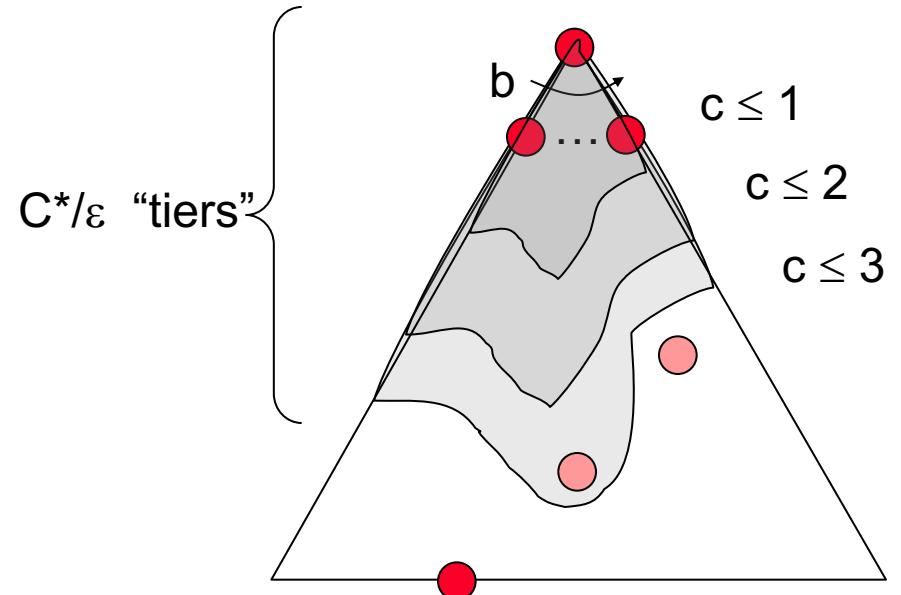
Strategy: expand a cheapest node first:

Fringe is a priority queue  
(priority: cumulative cost)



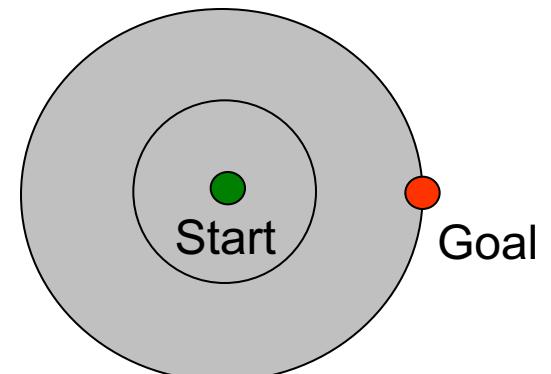
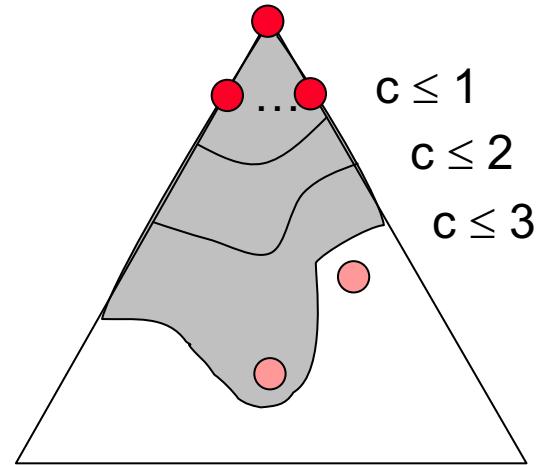
# Uniform Cost Search (UCS) Properties

- ❑ What nodes does UCS expand?
  - Processes all nodes with cost less than cheapest solution!
  - If that solution costs  $C^*$  and arcs cost at least  $\varepsilon$ , then the “effective depth” is roughly  $C^*/\varepsilon$
  - Takes time  $O(b^{C^*/\varepsilon})$  (exponential in effective depth)
- ❑ How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^{C^*/\varepsilon})$
- ❑ Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- ❑ Is it optimal?
  - Yes! (Proof via A\*)



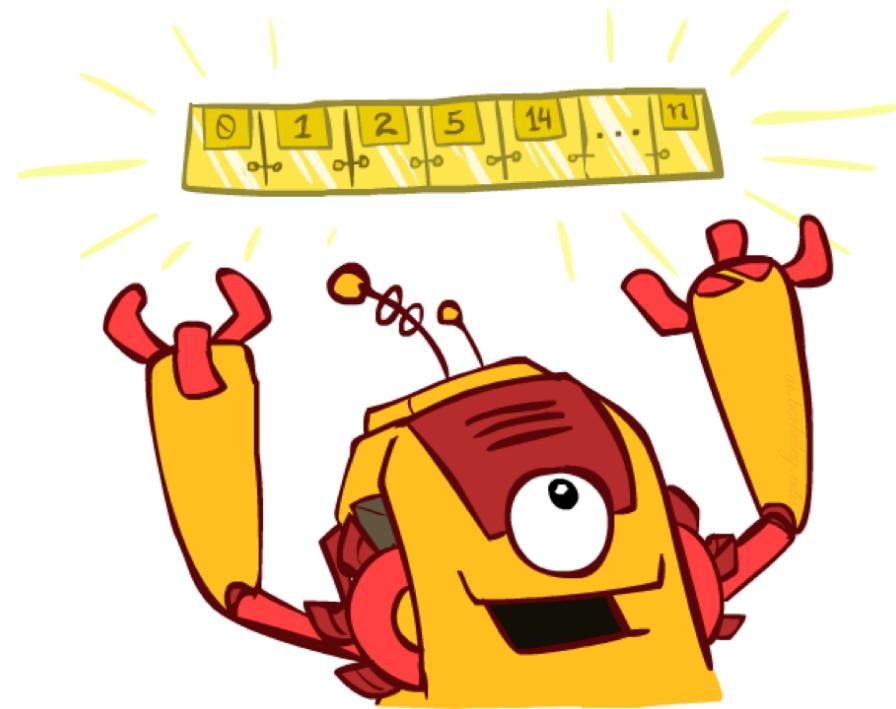
# Uniform Cost Issues

- ❑ Remember: UCS explores increasing cost contours
- ❑ The good: UCS is complete and optimal!
- ❑ The bad:
  - Explores options in every “direction”
  - No information about goal location
- ❑ We'll fix that soon!

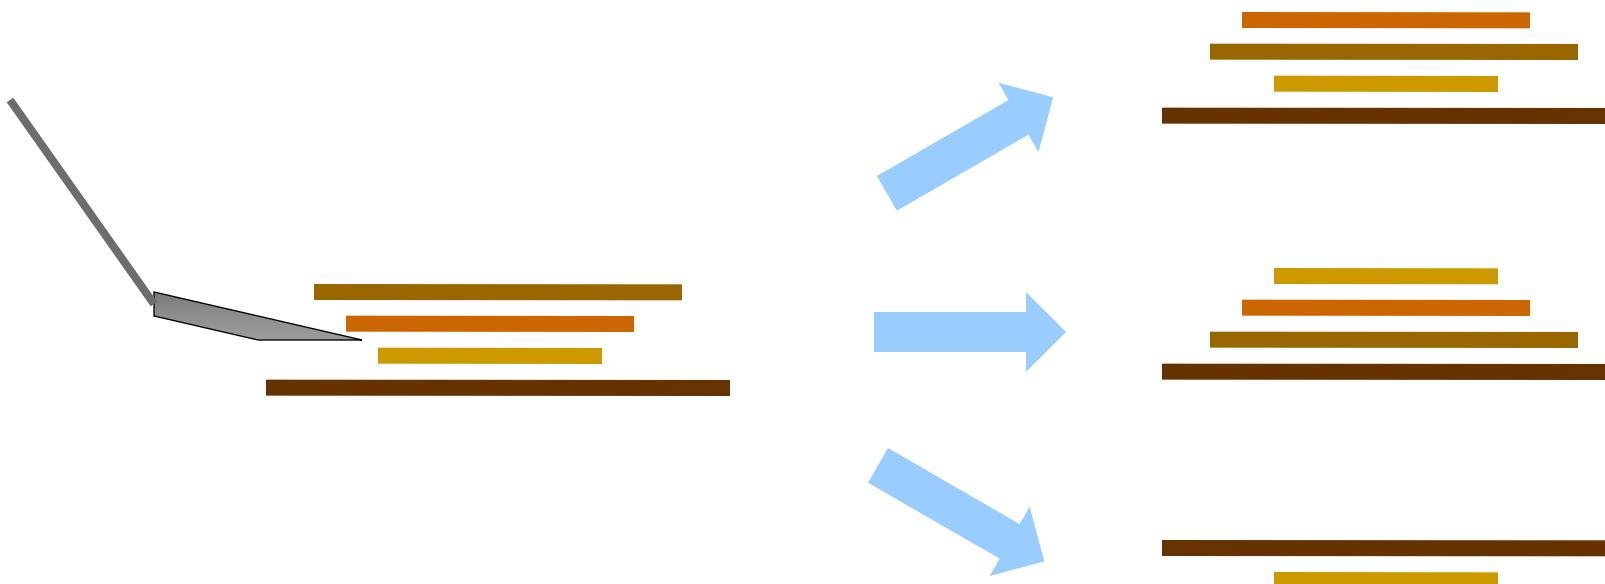


# The One Queue

- All these search algorithms are the same except for fringe strategies
  - Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
  - Practically, for DFS and BFS, you can avoid the  $\log(n)$  overhead from an actual priority queue, by using stacks and queues
  - Can even code one implementation that takes a variable queuing object



# Example: Pancake Problem



Cost: Number of pancakes flipped

# Example: Pancake Problem

## BOUNDS FOR SORTING BY PREFIX REVERSAL

William H. GATES

*Microsoft, Albuquerque, New Mexico*

Christos H. PAPADIMITRIOU\*†

*Department of Electrical Engineering, University of California, Berkeley, CA 94720, U.S.A.*

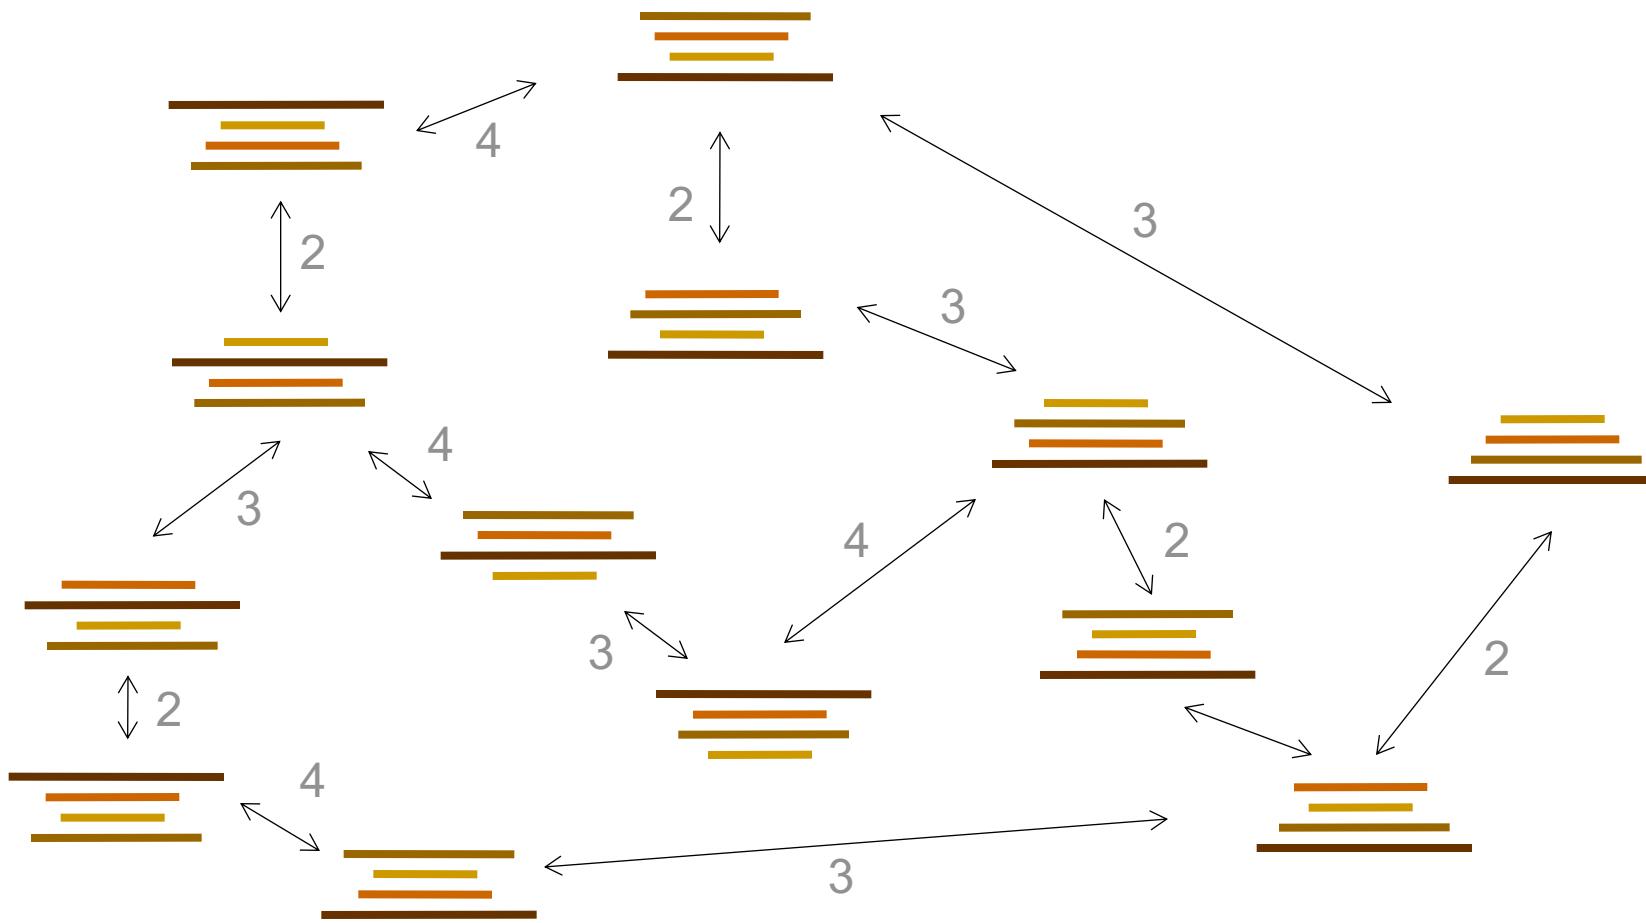
Received 18 January 1978

Revised 28 August 1978

For a permutation  $\sigma$  of the integers from 1 to  $n$ , let  $f(\sigma)$  be the smallest number of prefix reversals that will transform  $\sigma$  to the identity permutation, and let  $f(n)$  be the largest such  $f(\sigma)$  for all  $\sigma$  in (the symmetric group)  $S_n$ . We show that  $f(n) \leq (5n + 5)/3$ , and that  $f(n) \geq 17n/16$  for  $n$  a multiple of 16. If, furthermore, each integer is required to participate in an even number of reversed prefixes, the corresponding function  $g(n)$  is shown to obey  $3n/2 - 1 \leq g(n) \leq 2n + 3$ .

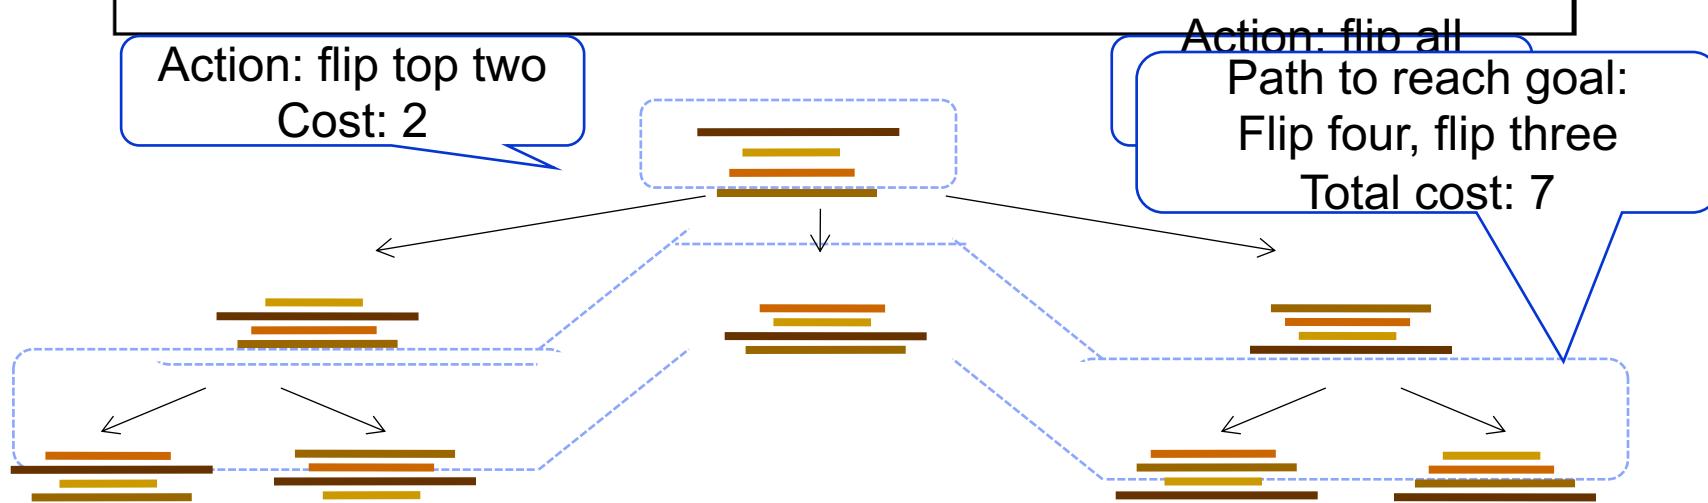
# Example: Pancake Problem

State space graph with costs as weights

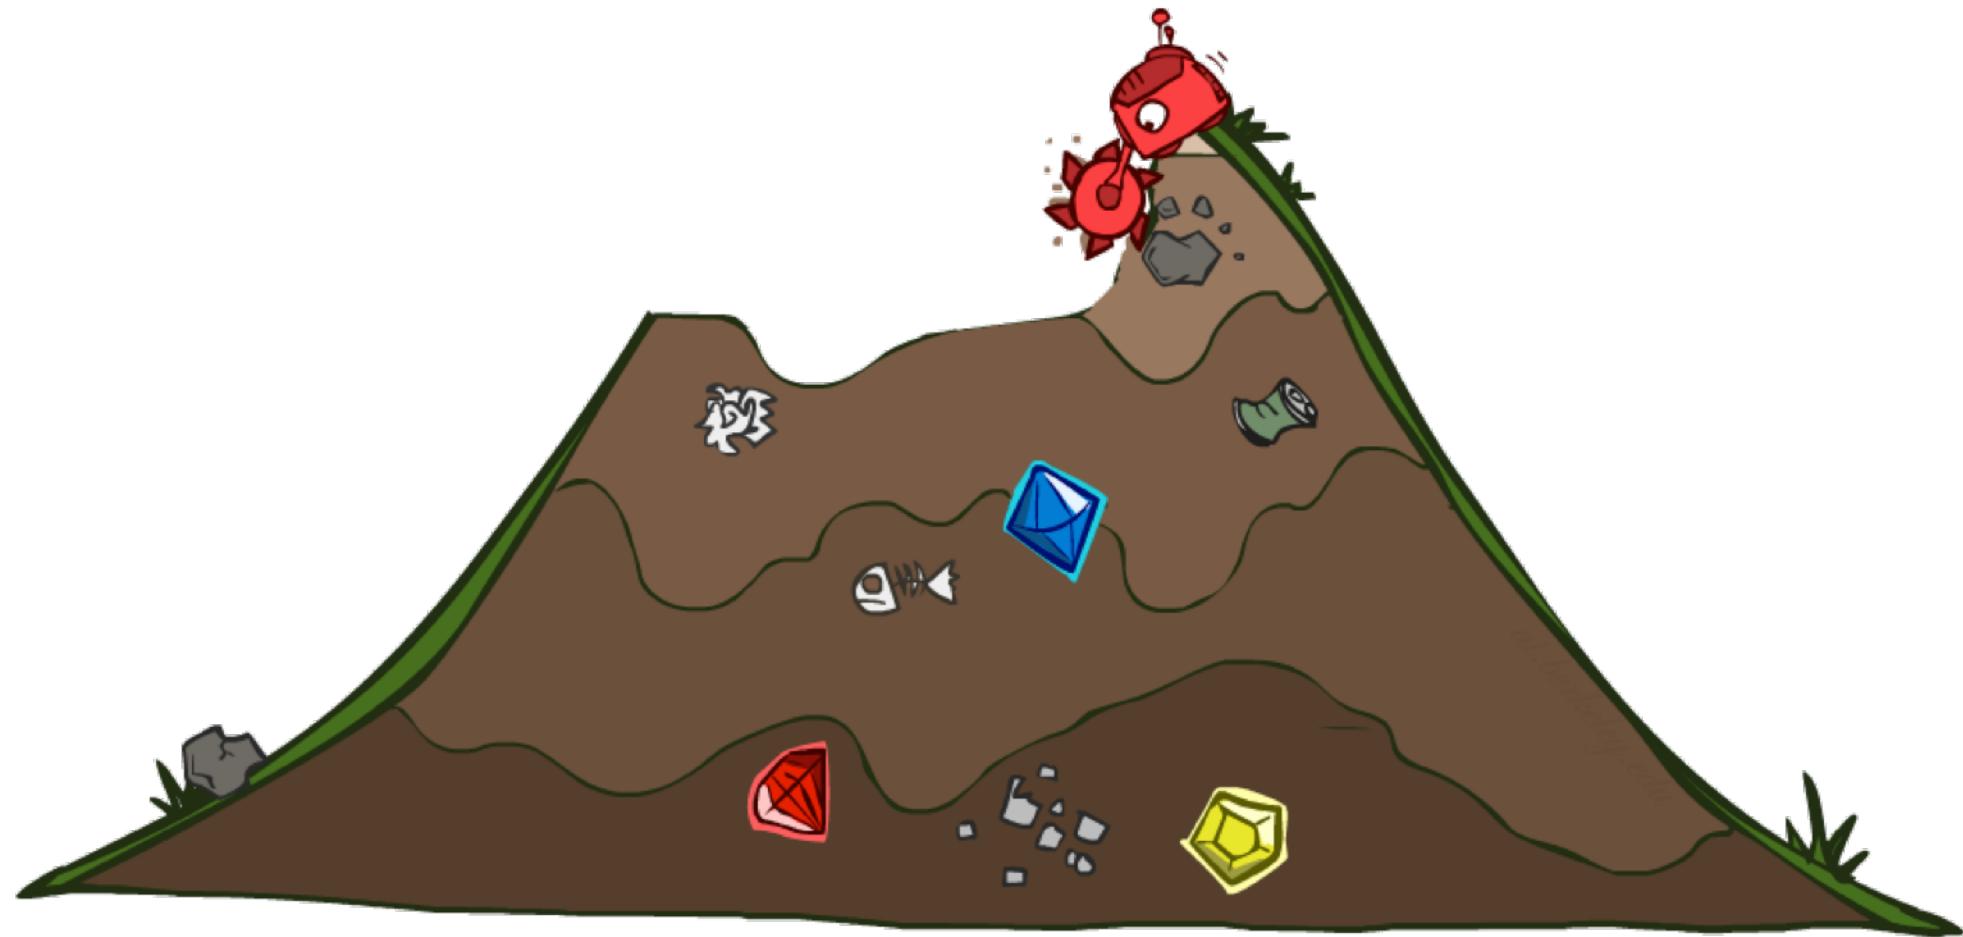


# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
    initialize the search tree using the initial state of problem
    loop do
        if there are no candidates for expansion then return failure
        choose a leaf node for expansion according to strategy
        if the node contains a goal state then return the corresponding solution
        else expand the node and add the resulting nodes to the search tree
    end
```

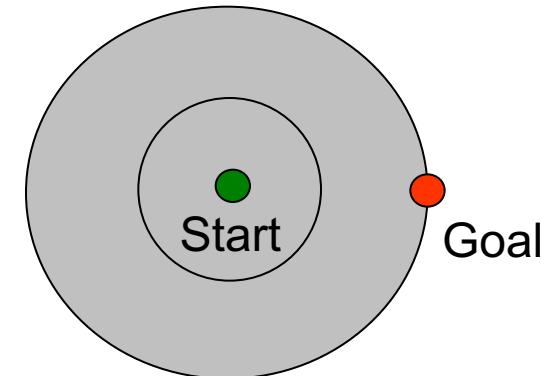
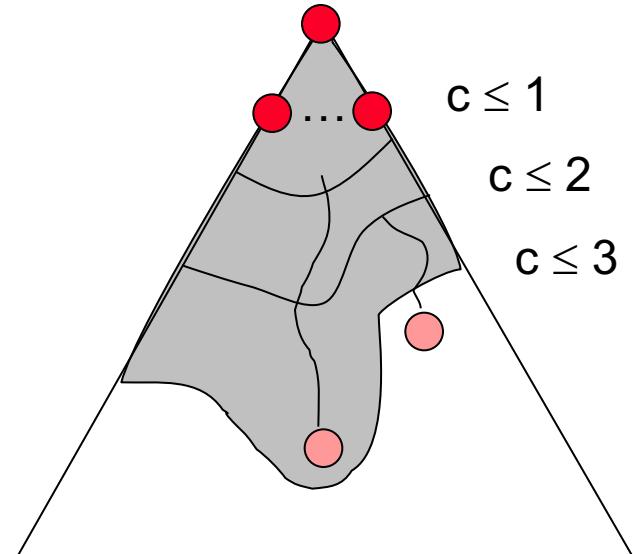


# Uninformed Search



# Uniform Cost Search

- ❑ Strategy: expand lowest path cost
- ❑ The good: UCS is complete and optimal!
- ❑ The bad:
  - Explores options in every “direction”
  - No information about goal location

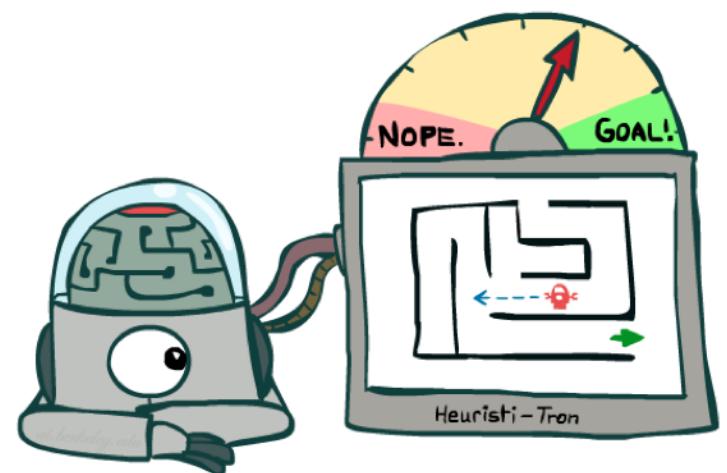
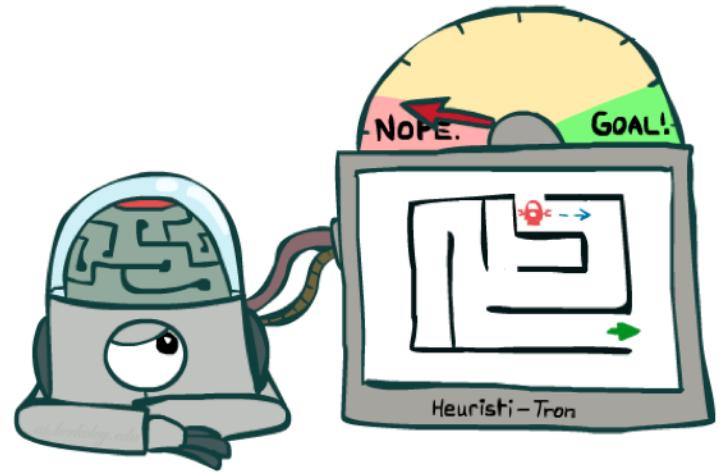
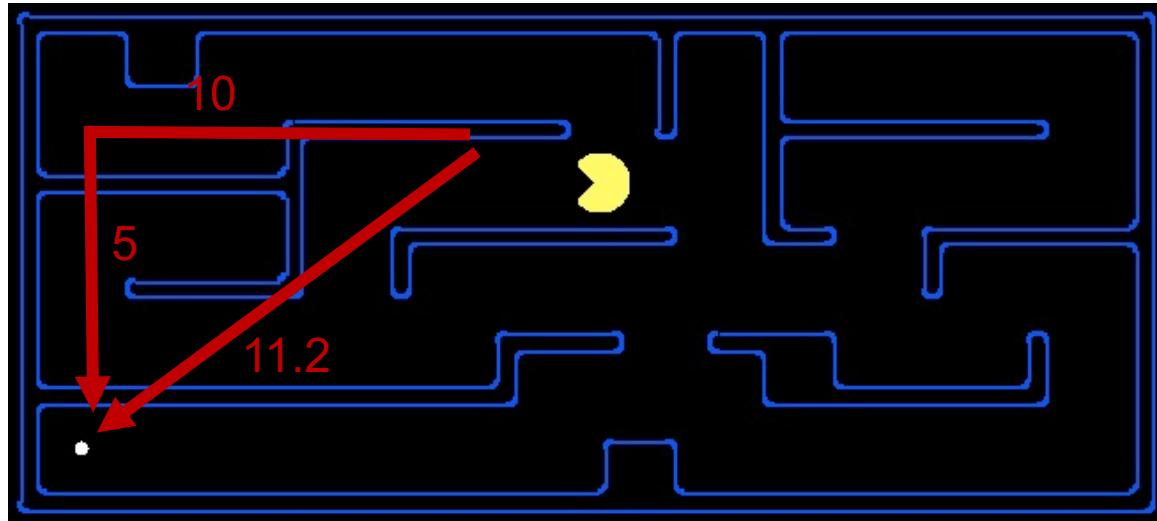


# Informed Search

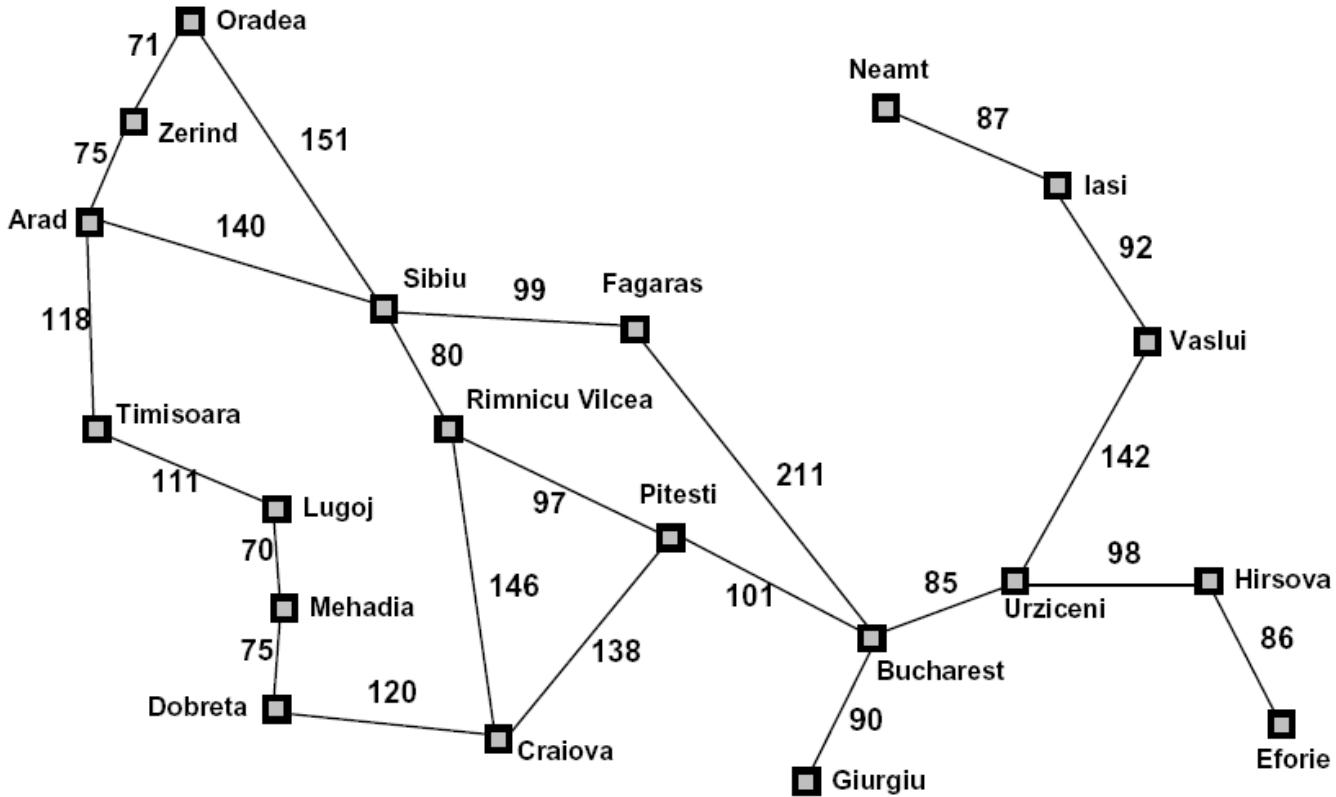


# Search Heuristics

- A heuristic is:
  - A function that estimates how close a state is to a goal
  - Designed for a particular search problem
  - Examples: Manhattan distance, Euclidean distance for pathing



# Example: Heuristic Function



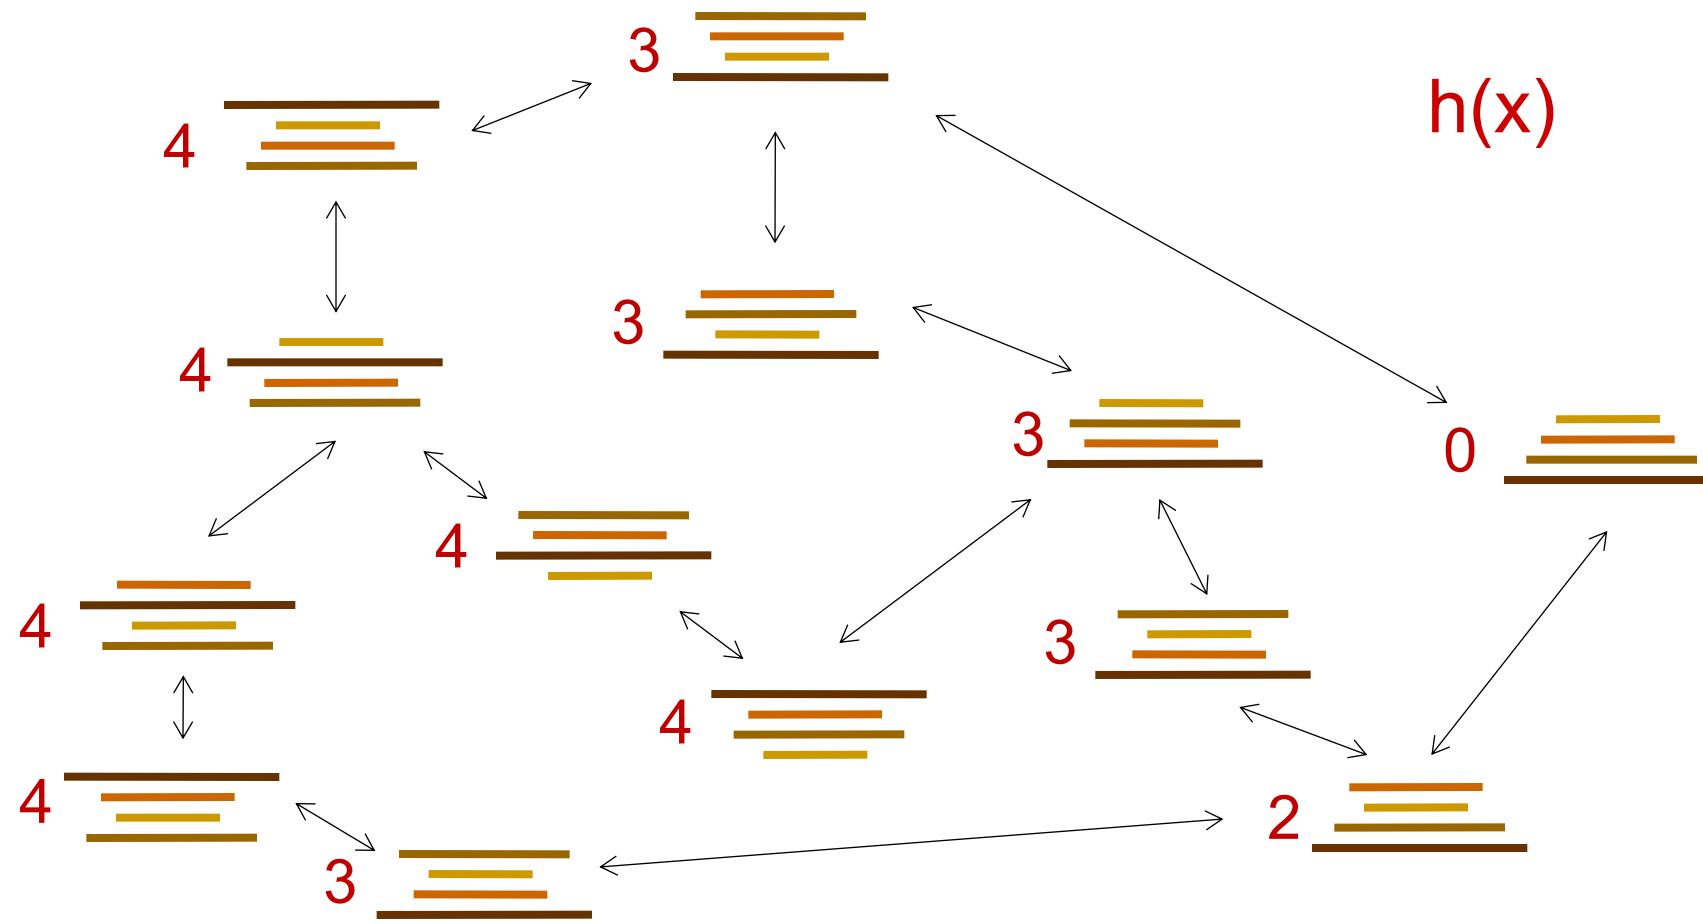
Straight-line distance to Bucharest

City	Distance to Bucharest
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

$h(x)$

# Example: Heuristic Function

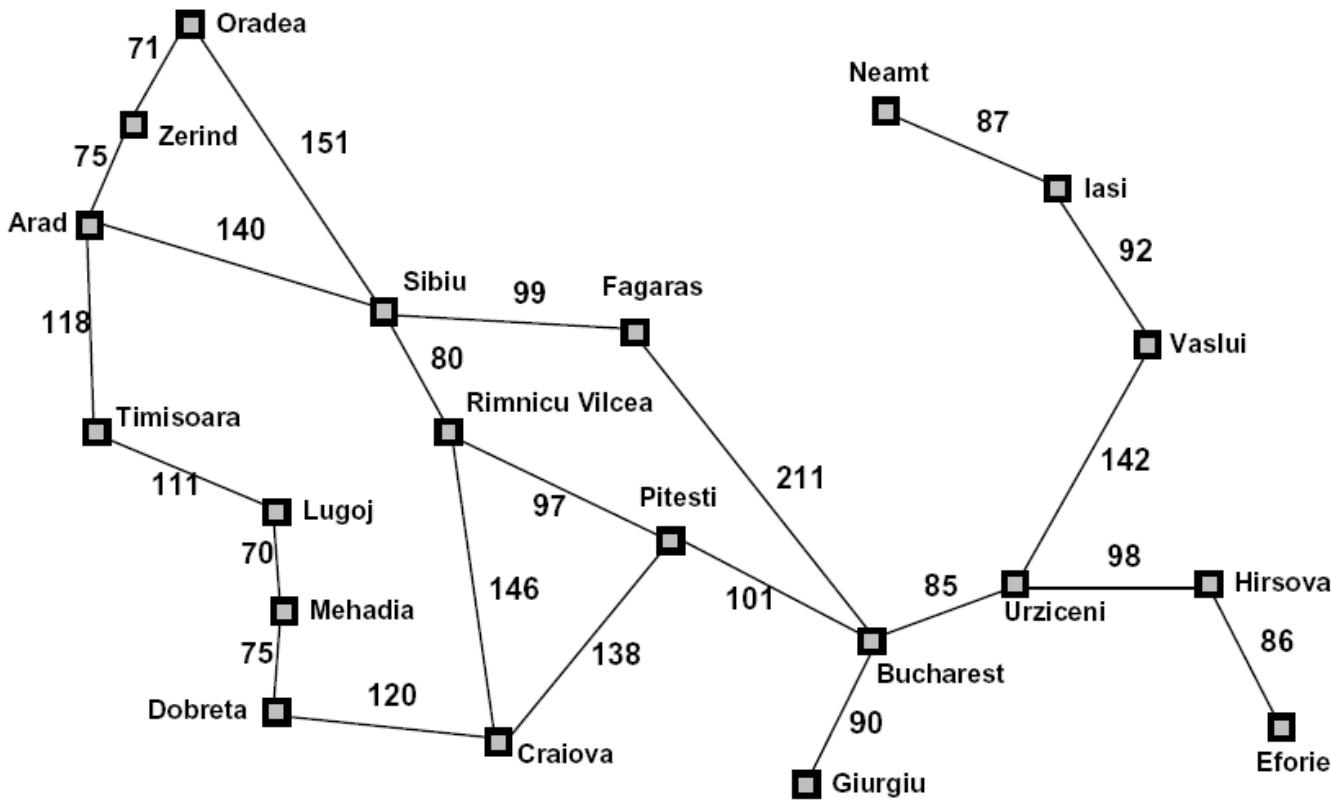
Heuristic: the number of the largest pancake that is still out of place



# Greedy Search



# Example: Heuristic Function

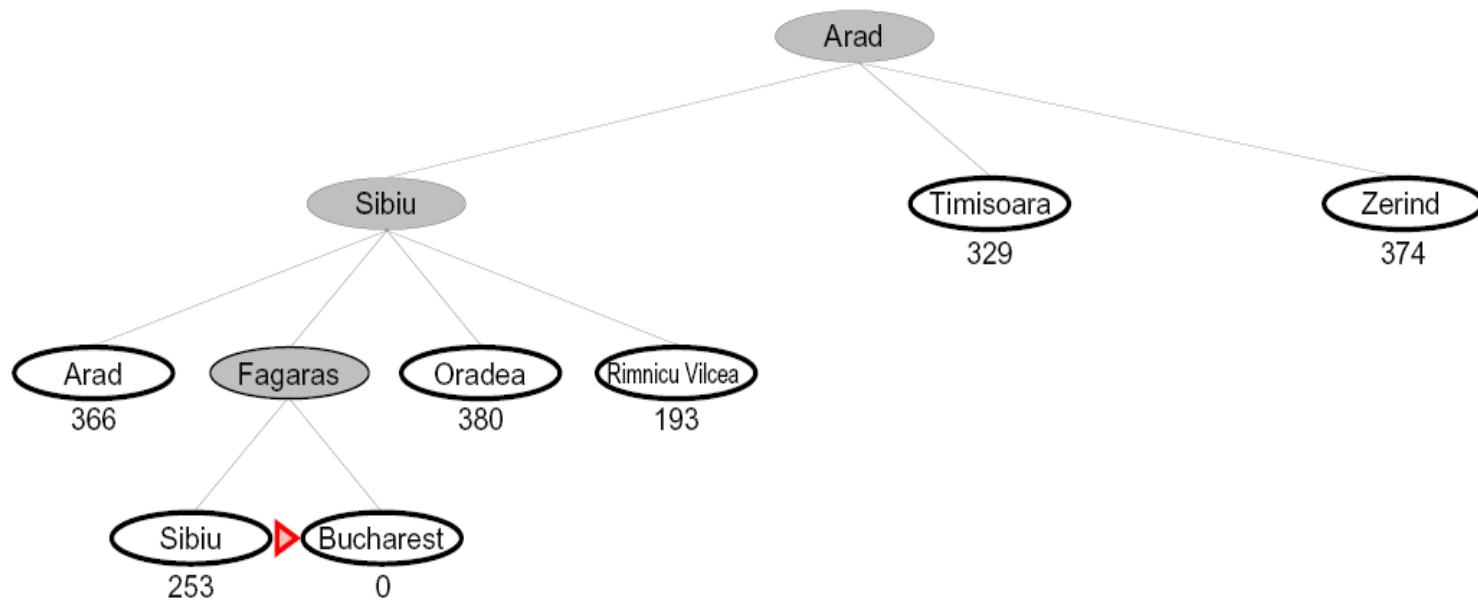


Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

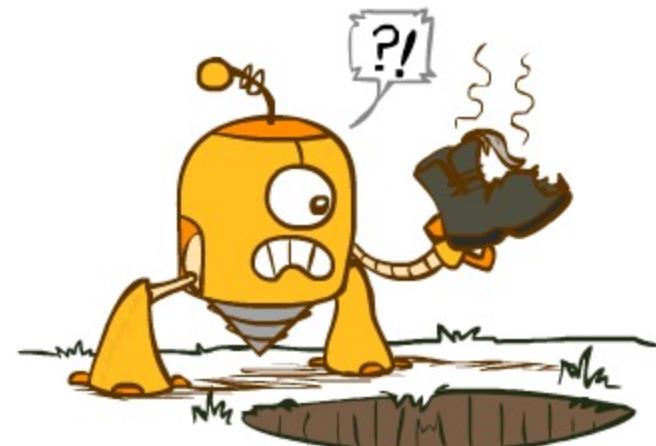
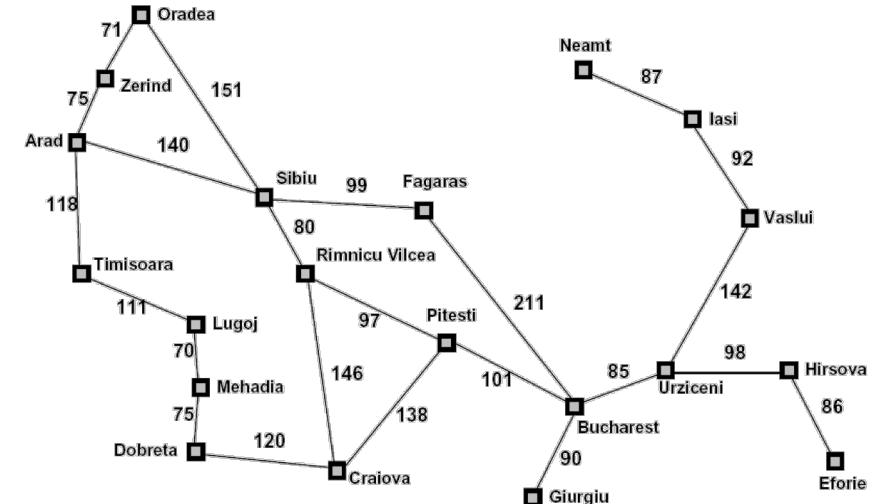
$h(x)$

# Greedy Search

- Expand the node that seems closest...

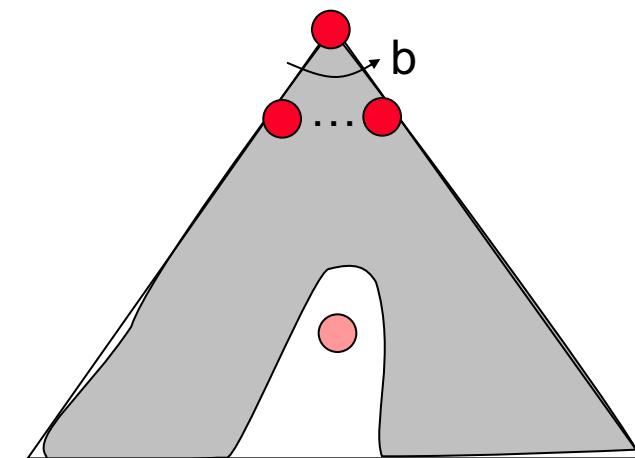
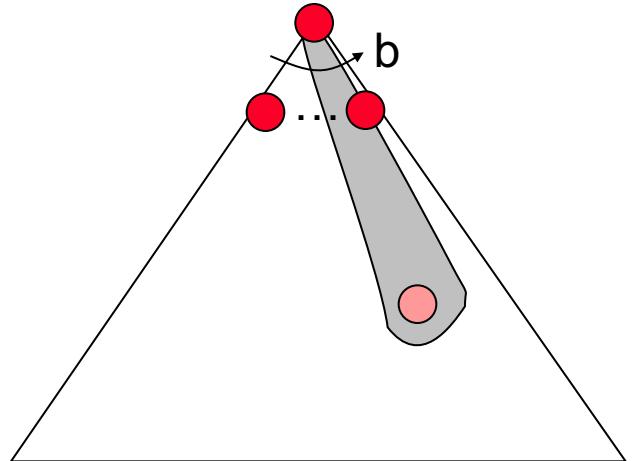


- What can go wrong?



# Greedy Search

- Strategy: expand a node that you think is closest to a goal state
  - Heuristic: estimate of distance to nearest goal for each state
- A common case:
  - Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS

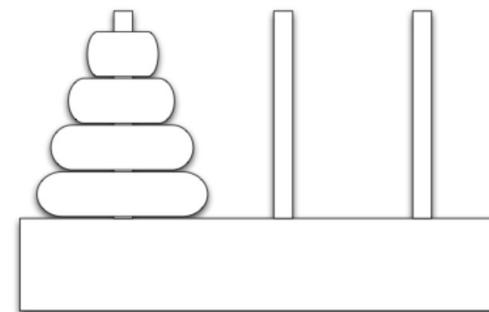


# **Review Slides for all search problems**

# Review (Search Problem formulation)

- The Towers of Hanoi is a famous problem for studying recursion in computer science and recurrence equations in discrete mathematics. We start with  $N$  discs of varying sizes on a peg (stacked in order according to size), and two empty pegs. We are allowed to move a disc from one peg to another, but we are never allowed to move a larger disc on top of a smaller disc. The goal is to move all the discs to the rightmost peg (see figure).

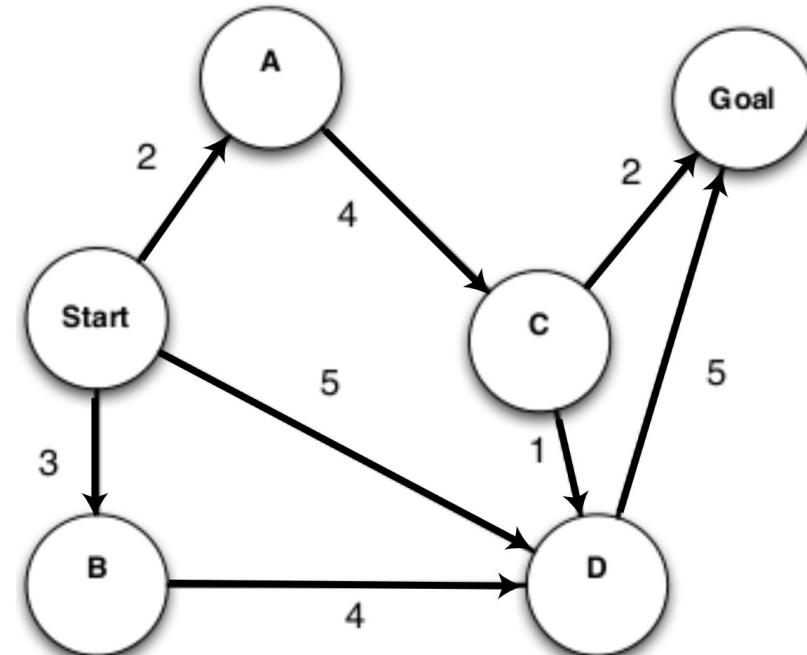
- Propose a state representation for the problem
- What is the start state?
- From a given state, what actions are legal?
- What is the goal test?



# Review (Search Algorithms)

For each of the following search strategies, work out the path returned by the search on the graph shown above. In all cases, assume ties resolve in such a way that states with earlier alphabetical order are expanded first. The start and goal state are S and G, respectively.

- a) Depth-first search.
- b) Breadth-first search.
- c) Uniform cost search.



# Review (Search Algorithms)

- ❑ Which of the following are true and which are false? Explain your answers.
1. Depth-first search always expands at least as many nodes as A\* search with an admissible heuristic.
  2.  $h(n) = 0$  is an admissible heuristic for the 8-puzzle.
  3. Breadth-first search is complete even if zero step costs are allowed.
  4. Assume that a rook can move on a chessboard any number of squares in a straight line, vertically or horizontally, but cannot jump over other pieces. Manhattan distance is an admissible heuristic for the problem of moving the rook from square A to square B in the smallest number of moves.

# Review (Search Algorithms)

- The heuristic path algorithm is a best-first search in which the objective function is  $f(n) = (2-w)g(n) + w h(n)$ . For what values of  $w$  is this algorithm guaranteed to be optimal? (You may assume that this is admissible.) What kind of search does this perform when  $w=0$ ? When  $w=1$ ? When  $w=2$ ?