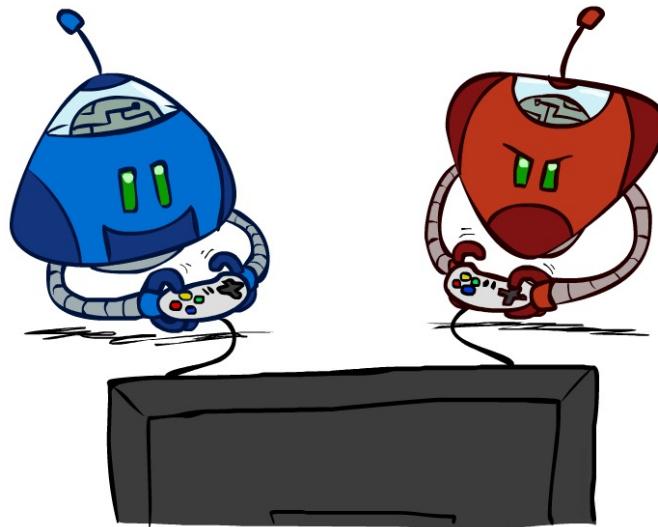


CS 3568: Intelligent Systems

Adversarial Search (Part 2)



Instructor: Tara Salman

Texas Tech University

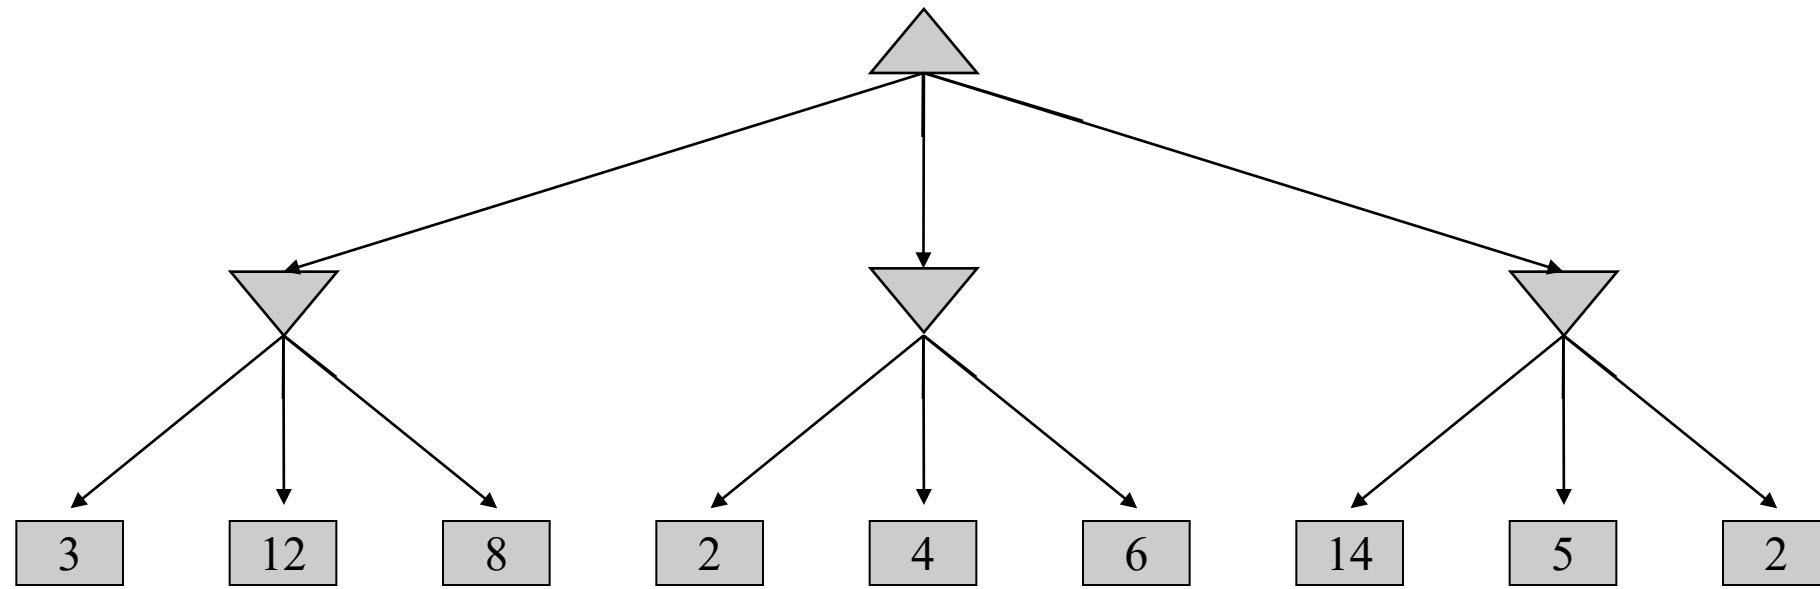
Computer Science Department

[These slides were created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley (ai.berkeley.edu).]

Texas Tech University

Tara Salman

Minimax Example



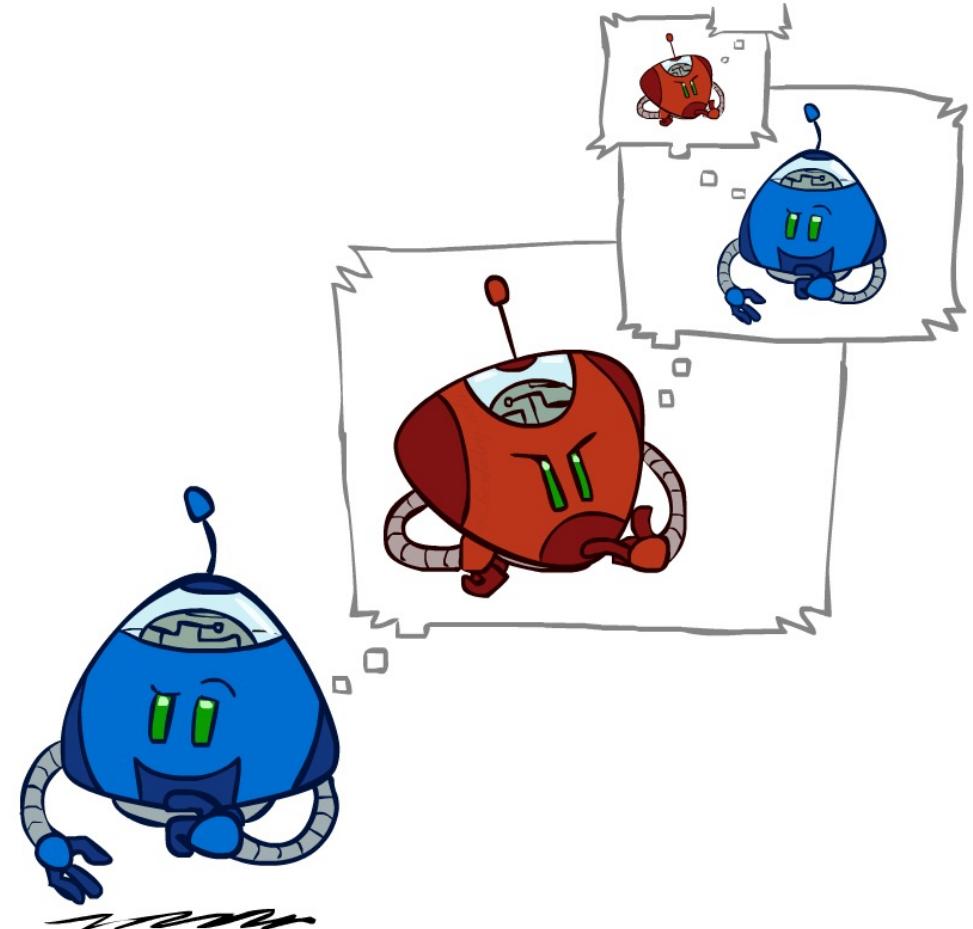
Minimax Efficiency

❑ How efficient is minimax?

- Just like (exhaustive) DFS
- Time: $O(b^m)$
- Space: $O(bm)$

❑ Example: For chess, $b \approx 35$, $m \approx 100$

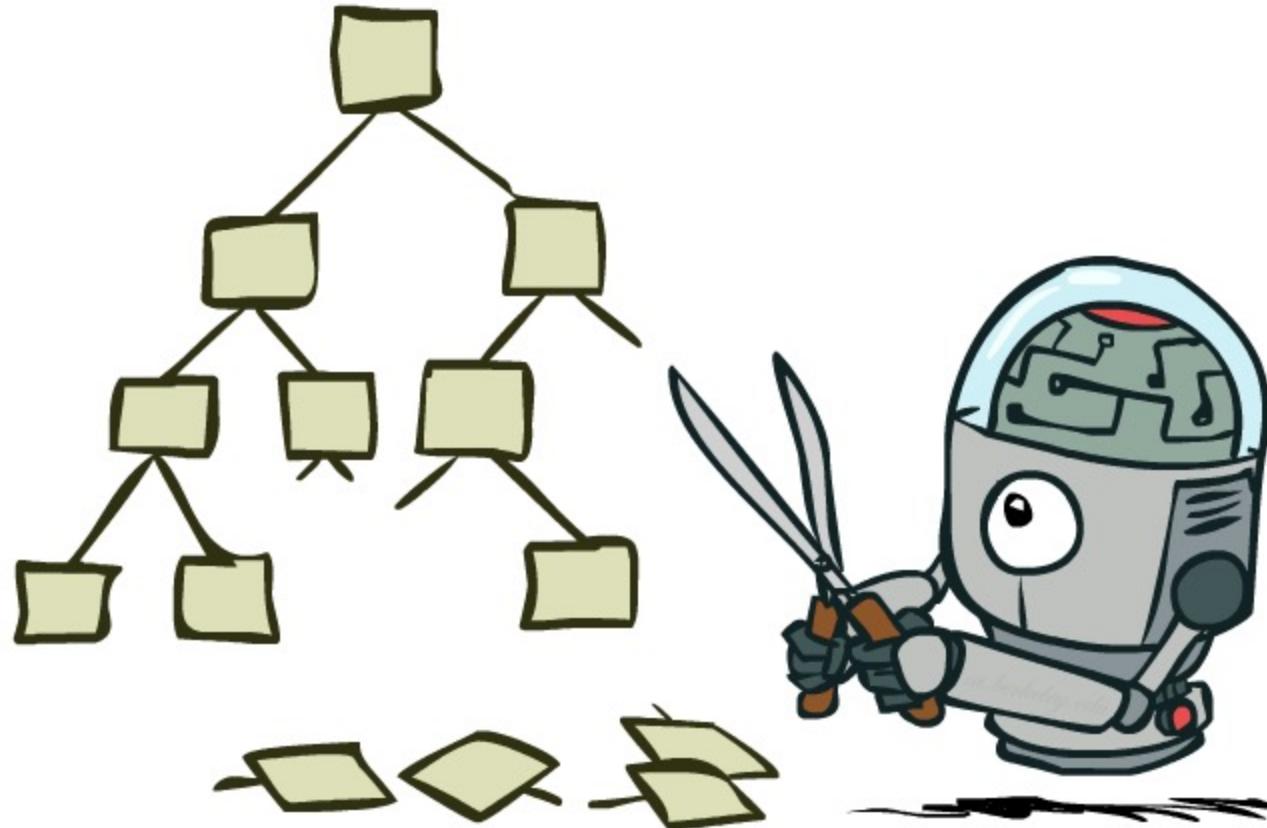
- Exact solution is completely infeasible
- But, do we need to explore the whole tree?



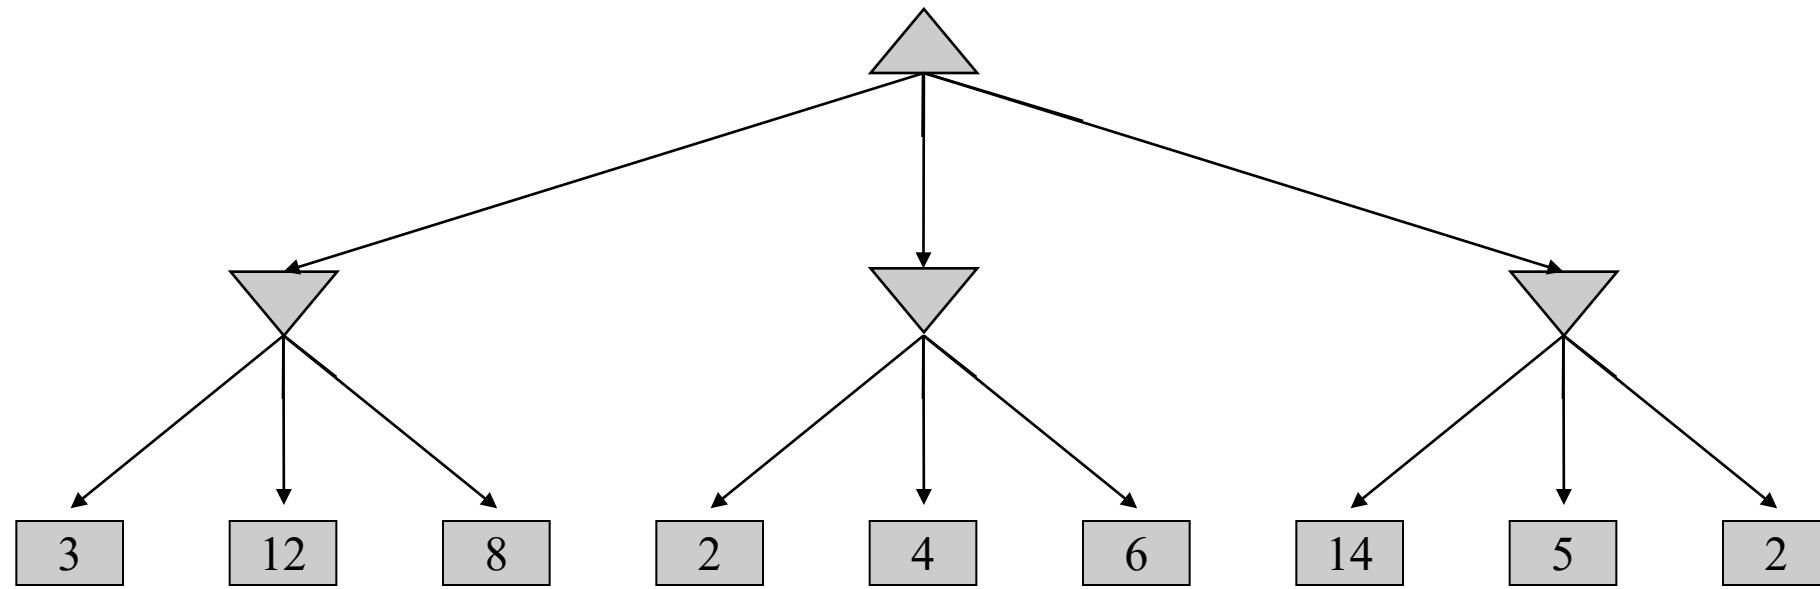
Resource Limits



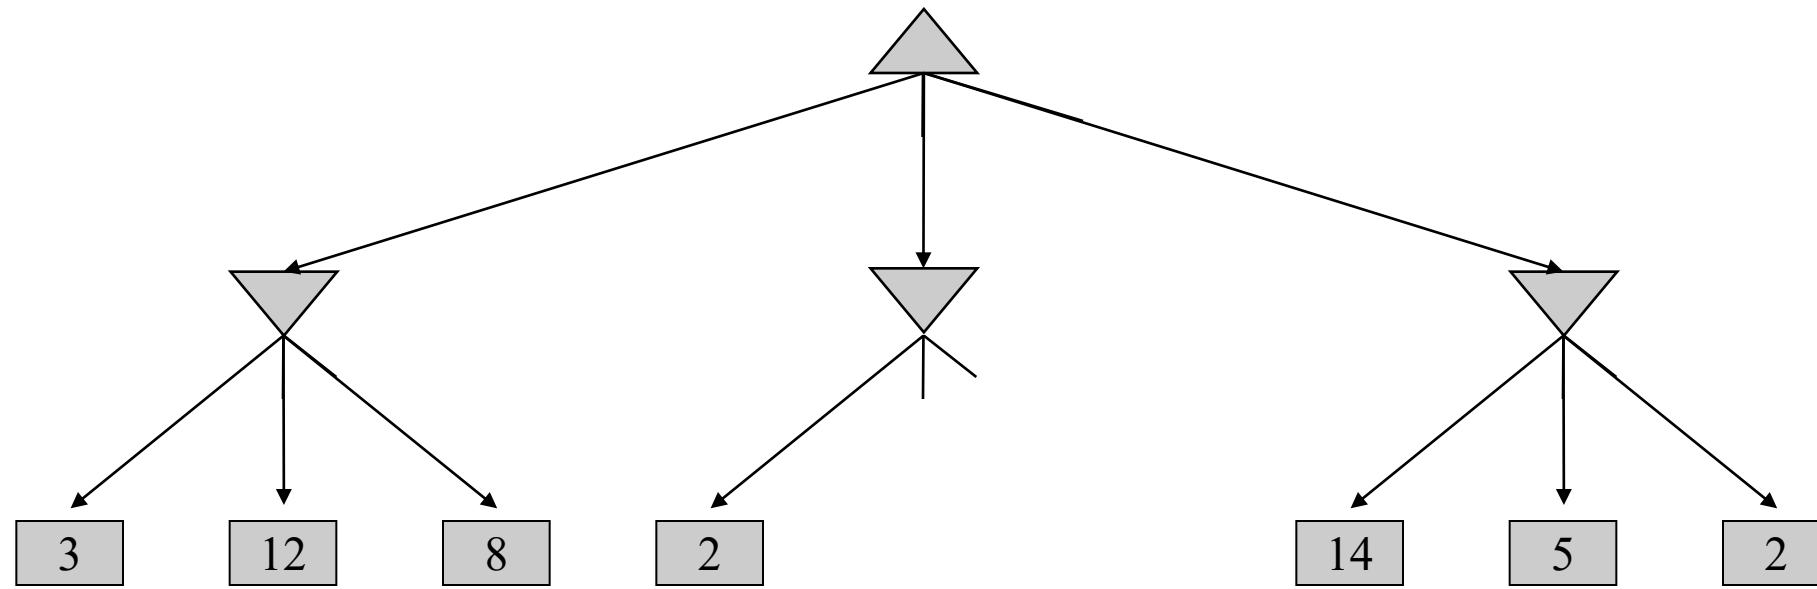
Game Tree Pruning



Minimax Example



Minimax Pruning

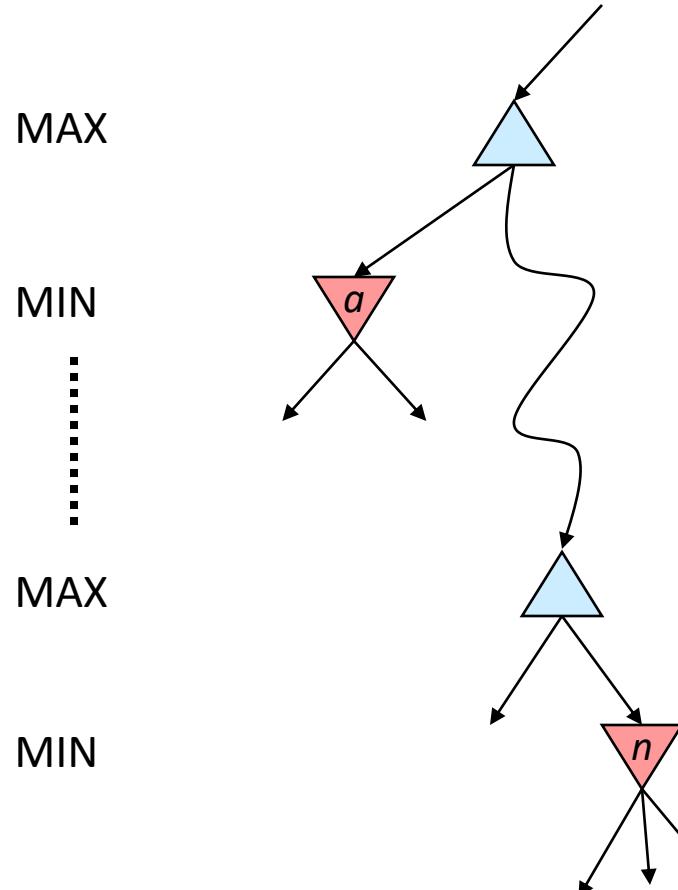


Alpha-Beta Pruning

□ General configuration (MIN version)

- We're computing the MIN-VALUE at some node n
- We're looping over n 's children
- n 's estimate of the childrens' min is dropping
- Who cares about n 's value? MAX
- Let a be the best value that MAX can get at any choice point along the current path from the root
- If n becomes worse than a , MAX will avoid it, so we can stop considering n 's other children (it's already bad enough that it won't be played)

□ MAX version is symmetric



Alpha-Beta Implementation

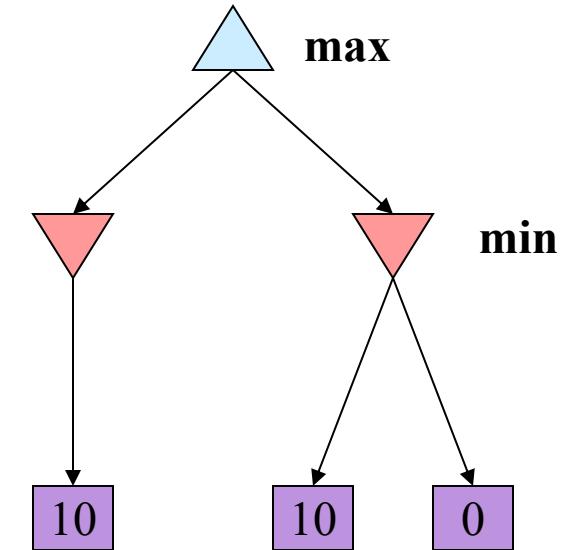
α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize v = - $\infty$   
    for each successor of state:  
        v = max(v, value(successor,  $\alpha$ ,  $\beta$ ))  
        if v  $\geq \beta$  return v  
         $\alpha$  = max( $\alpha$ , v)  
    return v
```

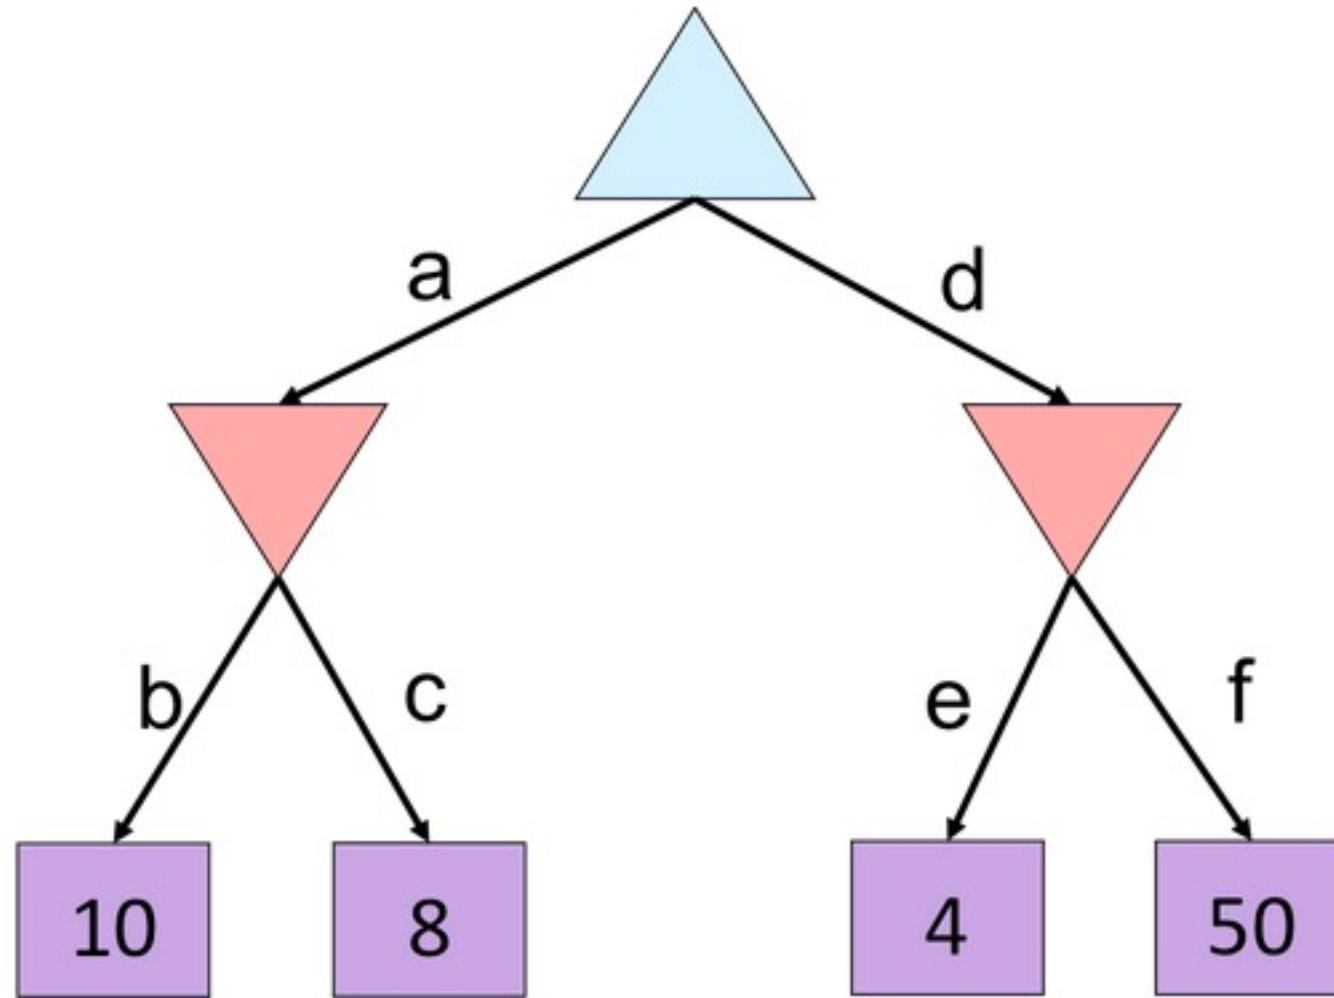
```
def min-value(state ,  $\alpha$ ,  $\beta$ ):  
    initialize v = + $\infty$   
    for each successor of state:  
        v = min(v, value(successor,  $\alpha$ ,  $\beta$ ))  
        if v  $\leq \alpha$  return v  
         $\beta$  = min( $\beta$ , v)  
    return v
```

Alpha-Beta Pruning Properties

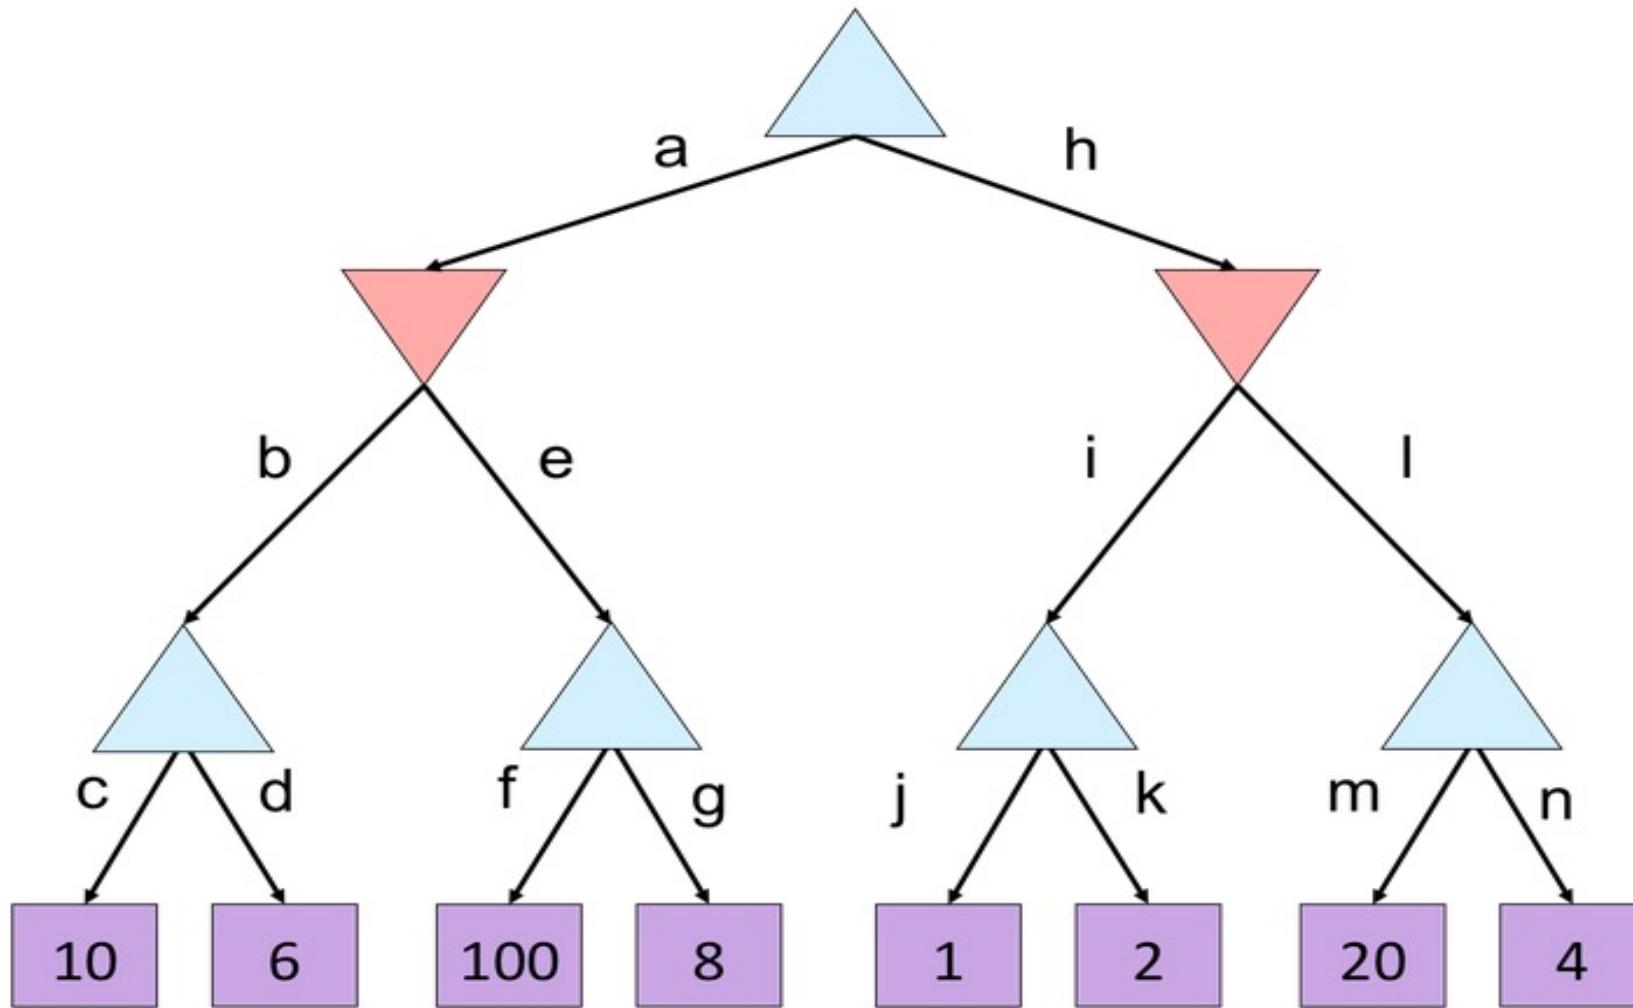
- ❑ This pruning has **no effect** on minimax value computed for the root!
- ❑ Values of intermediate nodes might be wrong
 - Important: children of the root may have the wrong value
 - So the most naïve version won't let you do action selection
- ❑ Good child ordering improves effectiveness of pruning
- ❑ With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!
 - Full search of, e.g. chess, is still hopeless...
- ❑ This is a simple example of **metareasoning** (computing about what to compute)



Alpha-Beta Quiz



Alpha-Beta Quiz 2

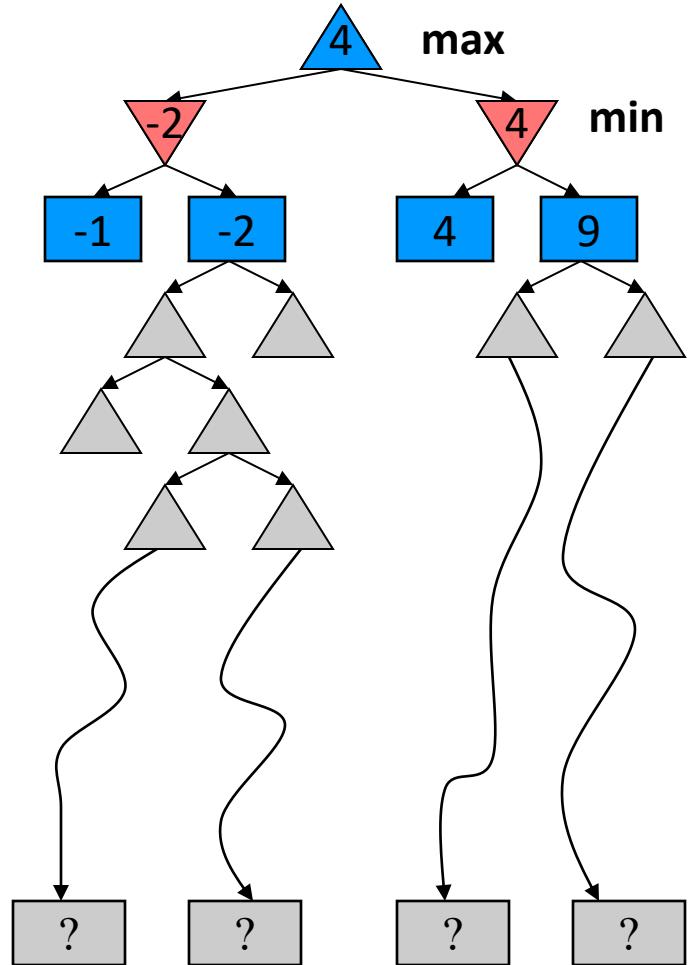


Resource Limits

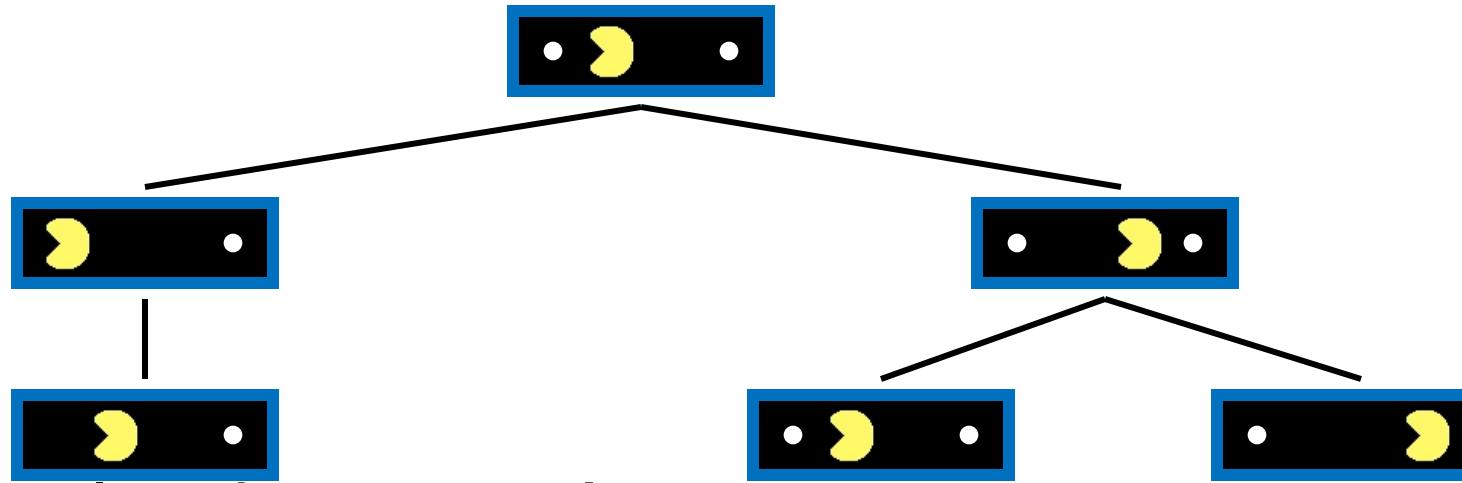


Resource Limits

- ❑ Problem: In realistic games, cannot search to leaves!
- ❑ Solution: Depth-limited search
 - Instead, search only to a limited depth in the tree
 - Replace terminal utilities with an evaluation function for non-terminal positions
- ❑ Example:
 - Suppose we have 100 seconds, can explore 10K nodes / sec
 - So can check 1M nodes per move
 - $\alpha\beta$ reaches about depth 8 – decent chess program
- ❑ Guarantee of optimal play is gone
- ❑ More plies makes a BIG difference
- ❑ Use iterative deepening for an anytime algorithm



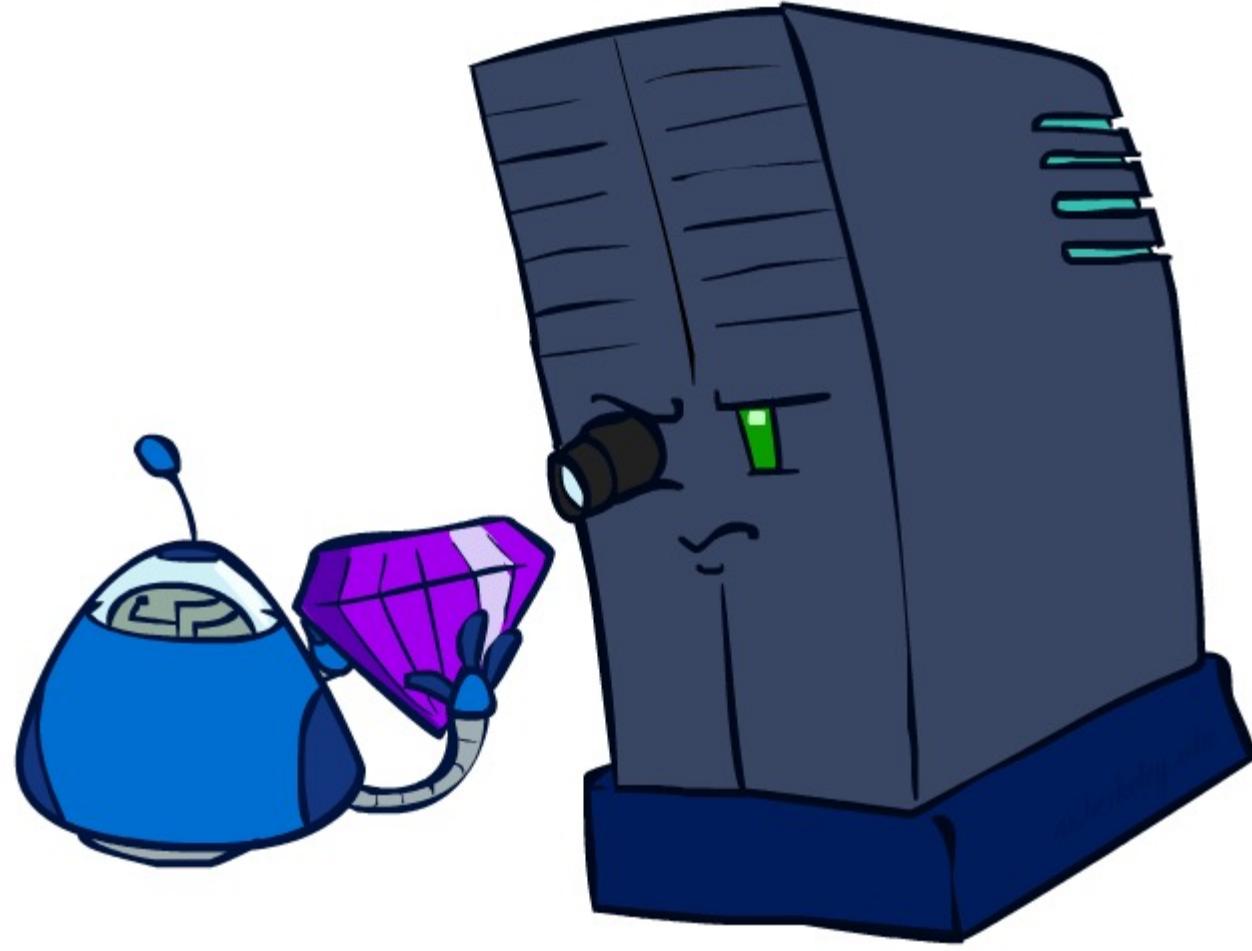
Why Pacman Starves



❑ A danger of replanning agents!

- He knows his score will go up by eating the dot now (west, east)
- He knows his score will go up just as much by eating the dot later (east, west)
- There are no point-scoring opportunities after eating the dot (within the horizon, two here)
- Therefore, waiting seems just as good as eating: he may go east, then back west in the next round of replanning!

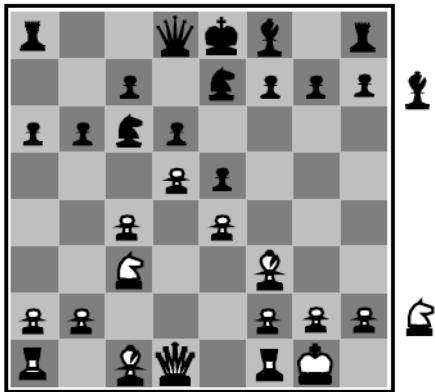
Evaluation Functions



ms

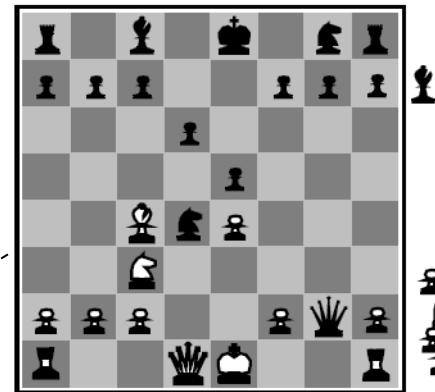
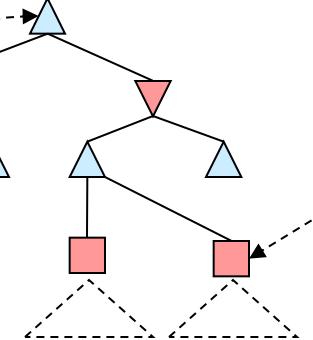
Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



Black to move

White slightly better



White to move

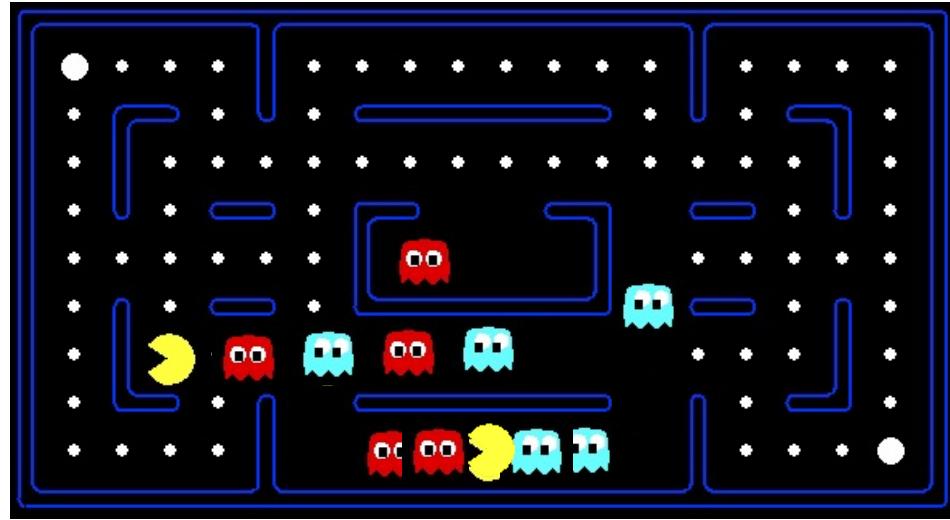
Black winning

- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

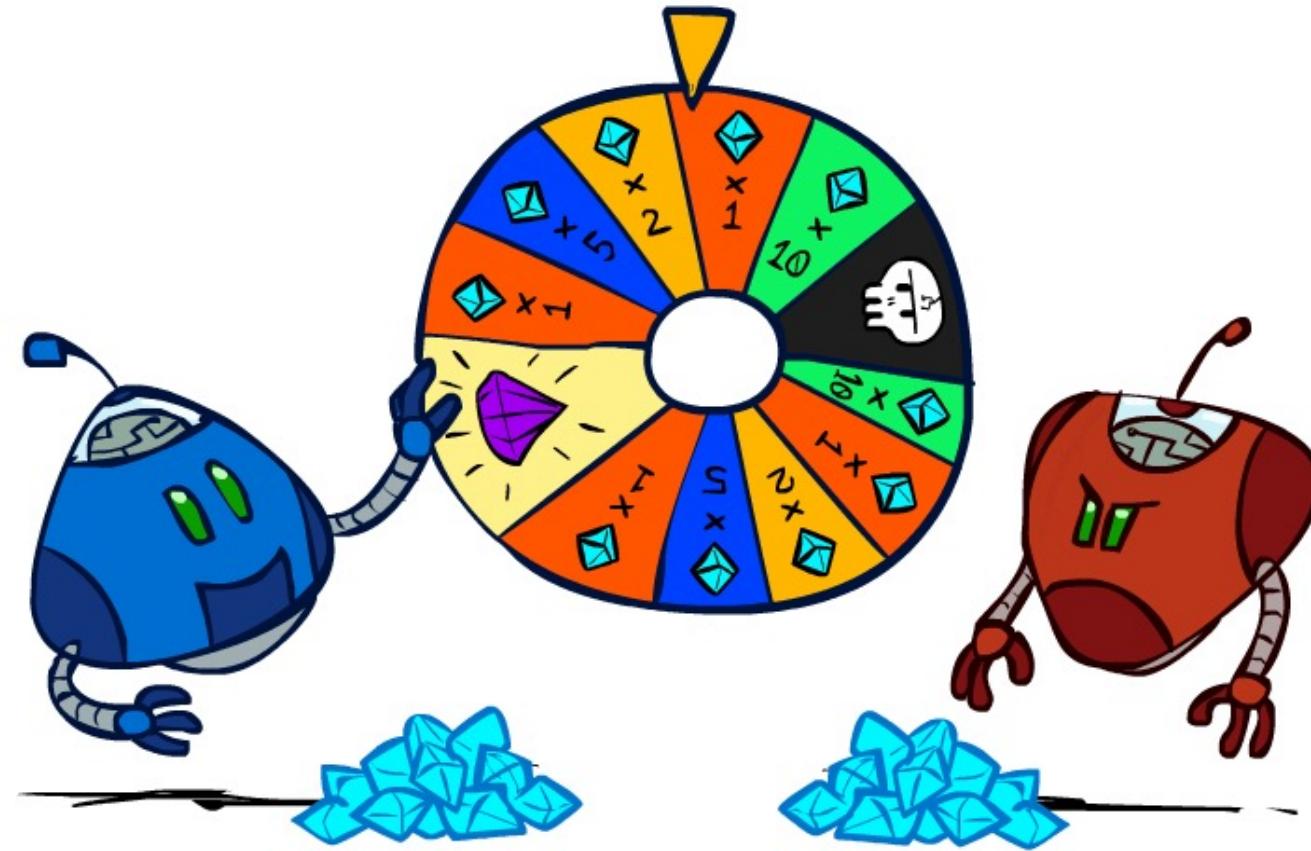
- e.g. $f_1(s) = (\text{num white queens} - \text{num black queens})$, etc.

Evaluation for Pacman

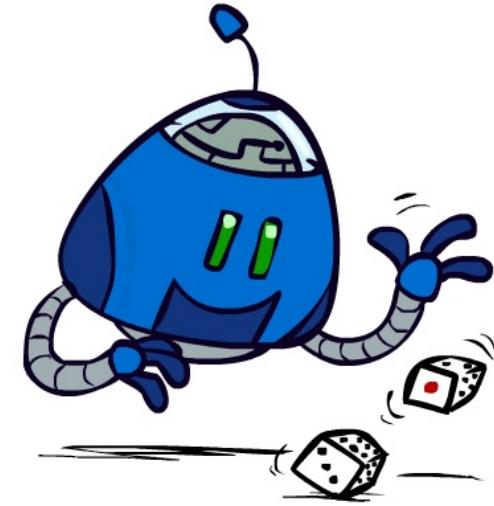
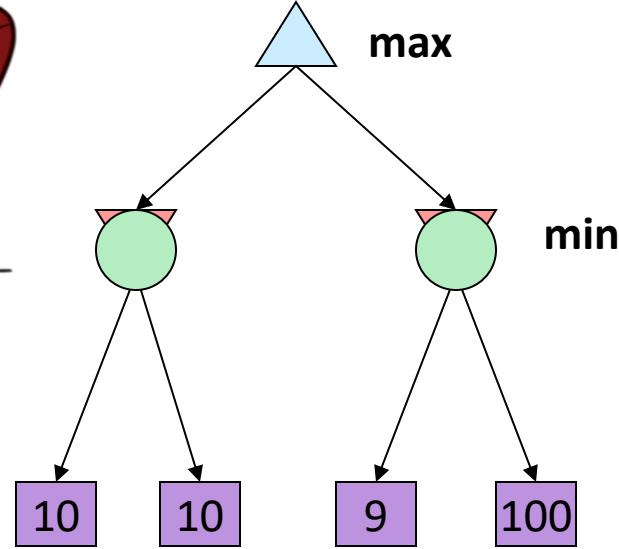
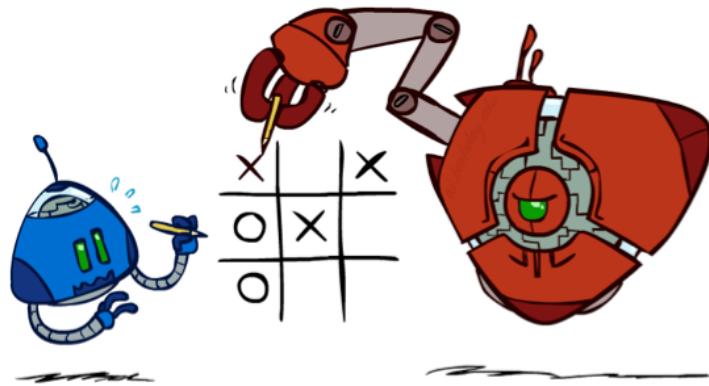


CS 3568: Intelligent Systems

Uncertainty and Utilities



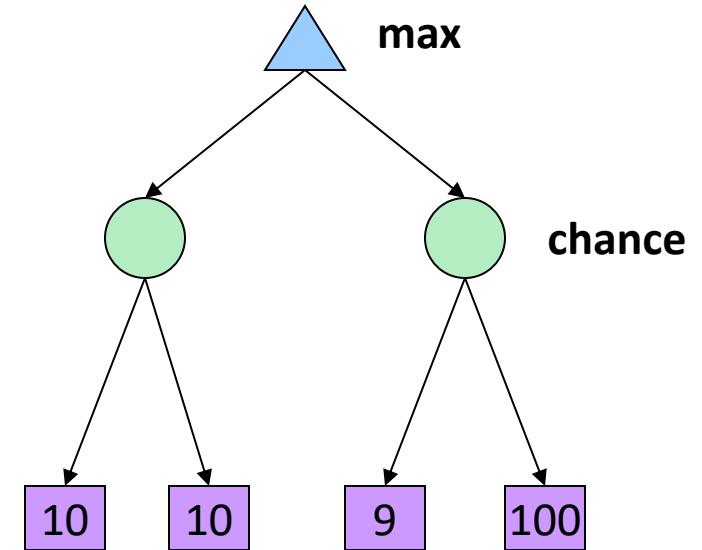
Worst-Case vs. Average Case



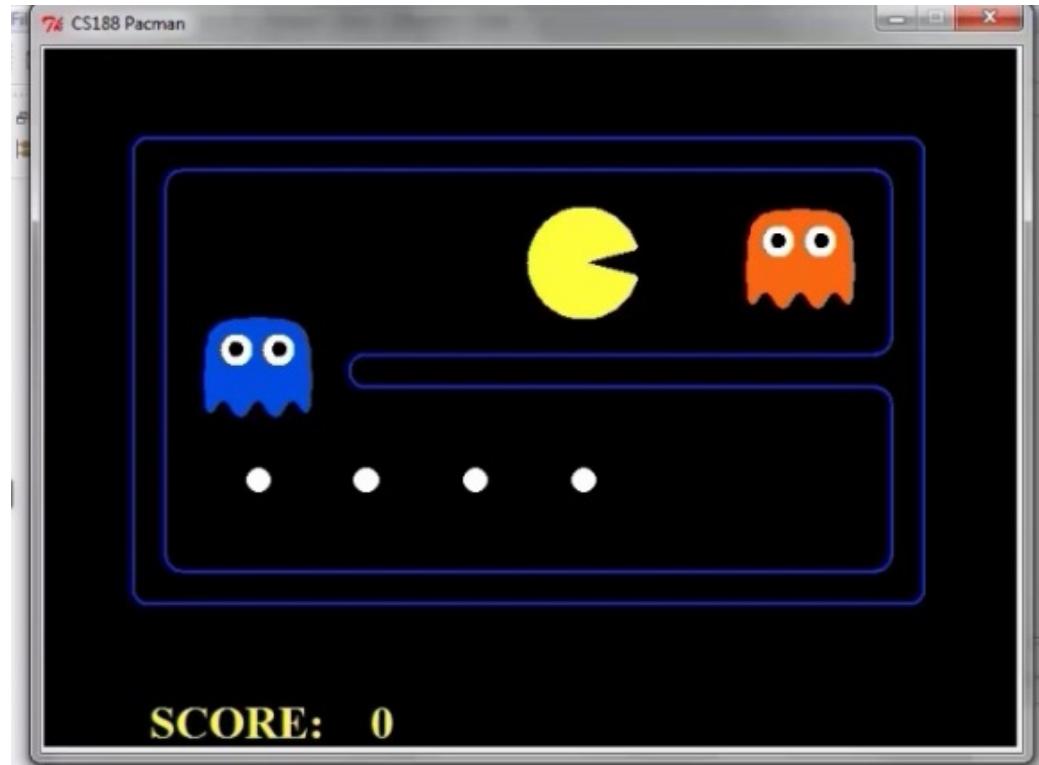
Idea: Uncertain outcomes controlled by chance, not an adversary!

Expectimax Search

- ❑ Why wouldn't we know what the result of an action will be?
 - Explicit randomness: rolling dice
 - Unpredictable opponents: the ghosts respond randomly
 - Actions can fail: when moving a robot, wheels might slip
- ❑ Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcomes
- ❑ **Expectimax search:** compute the average score under optimal play
 - Max nodes as in minimax search
 - Chance nodes are like min nodes but the outcome is uncertain
 - Calculate their **expected utilities**
 - I.e. take weighted average (expectation) of children



Video of Demo Min vs. Exp (Min)

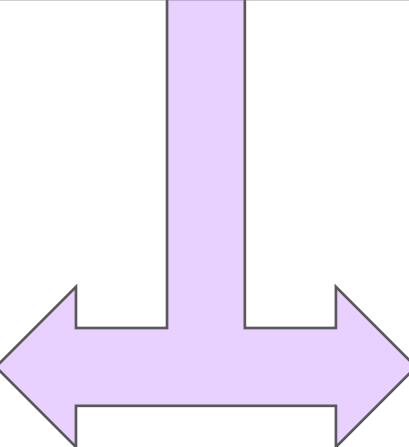


Expectimax Pseudocode

```
def value(state):
    if the state is a terminal state: return the state's utility
    if the next agent is MAX: return max-value(state)
    if the next agent is EXP: return exp-value(state)
```

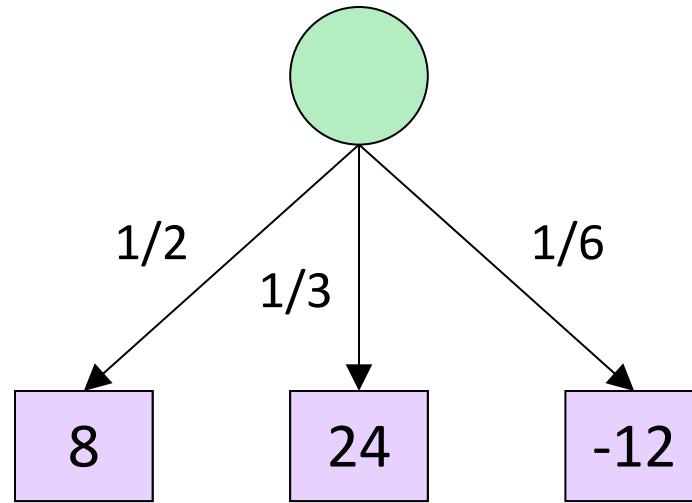
```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, value(successor))
    return v
```

```
def exp-value(state):
    initialize v = 0
    for each successor of state:
        p = probability(successor)
        v += p * value(successor)
    return v
```



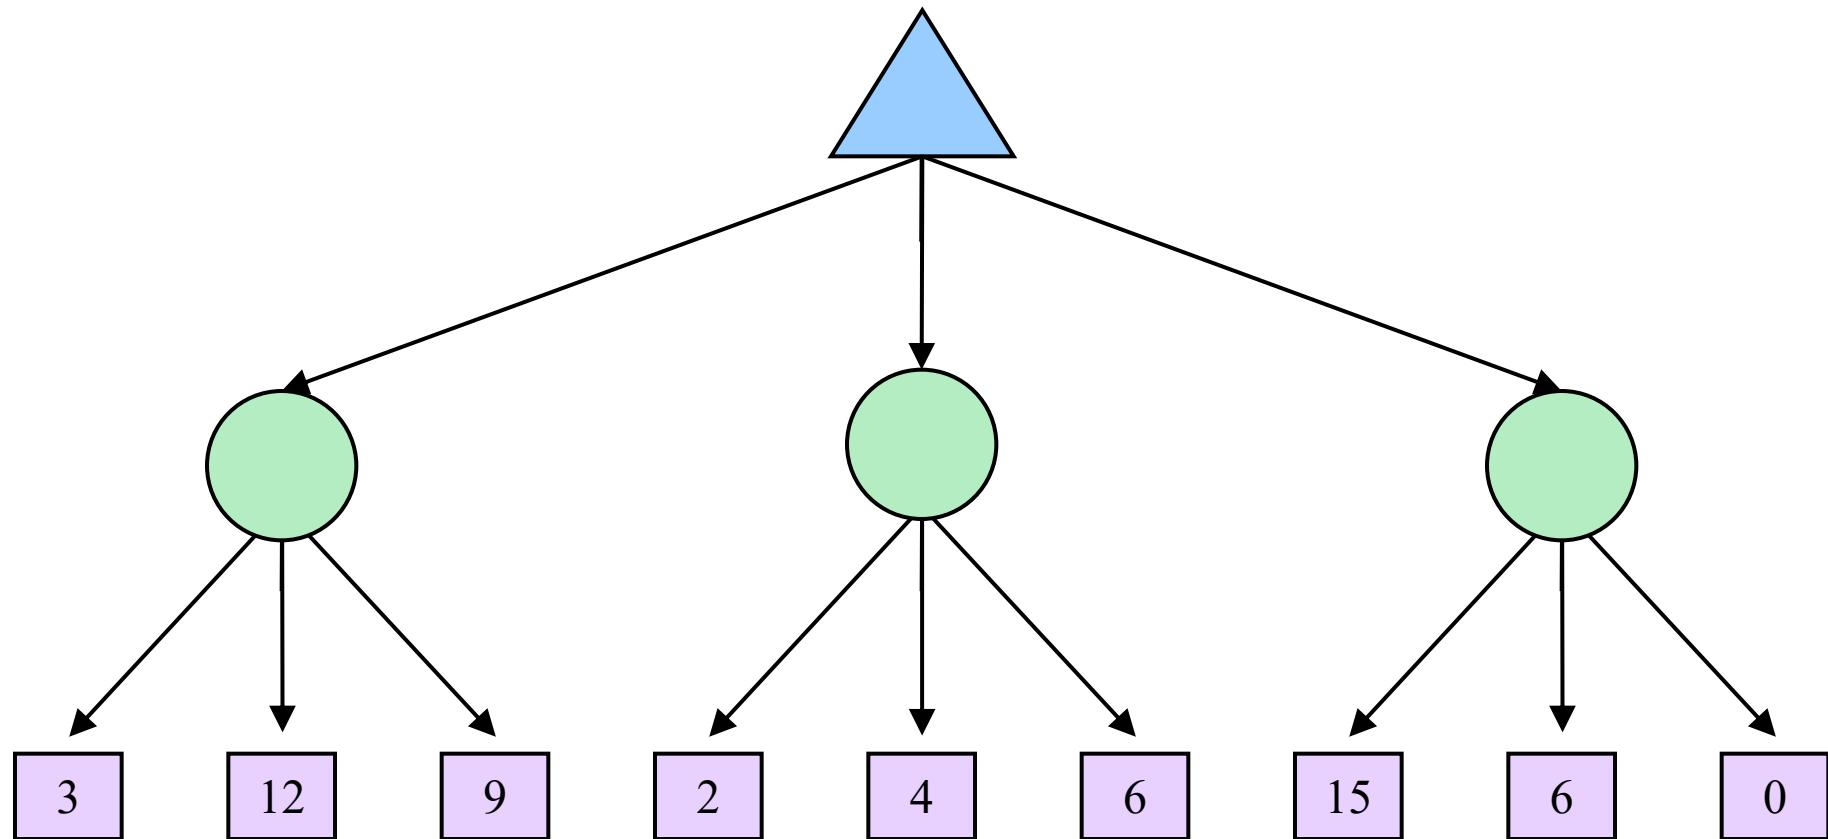
Expectimax Pseudocode

```
def exp-value(state):
    initialize v = 0
    for each successor of state:
        p = probability(successor)
        v += p * value(successor)
    return v
```

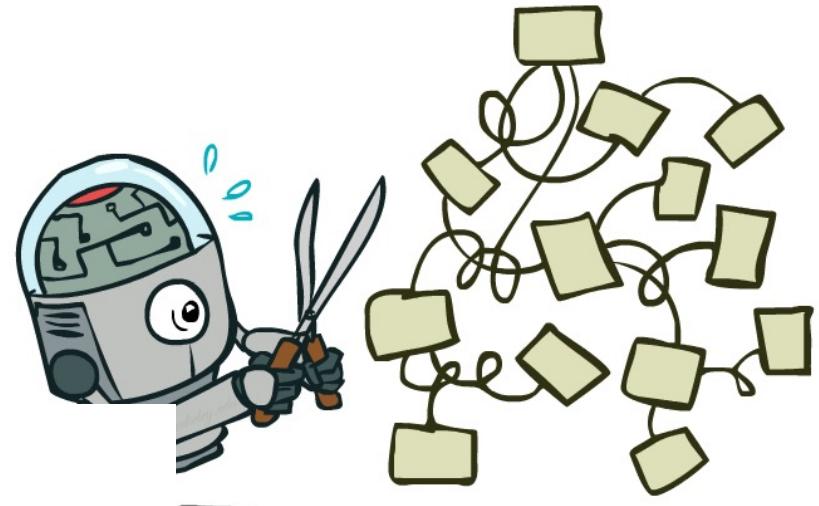
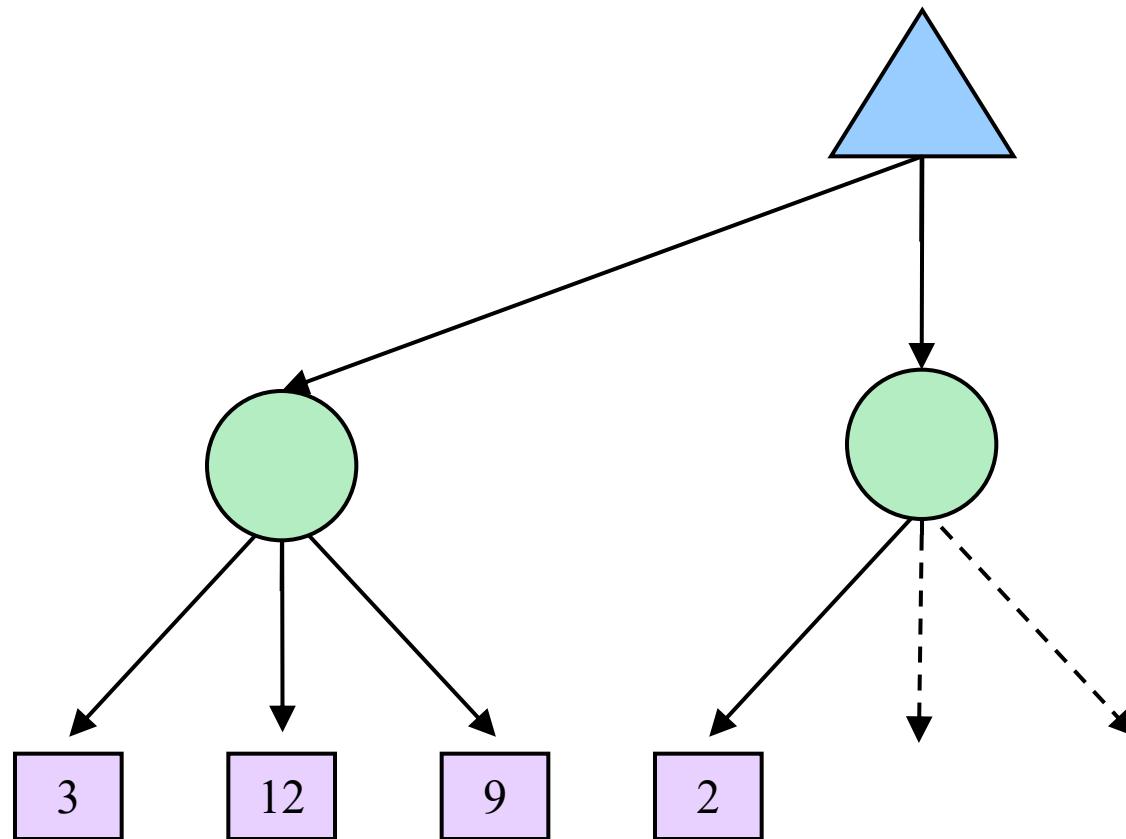


$$v = (1/2)(8) + (1/3)(24) + (1/6)(-12) = 10$$

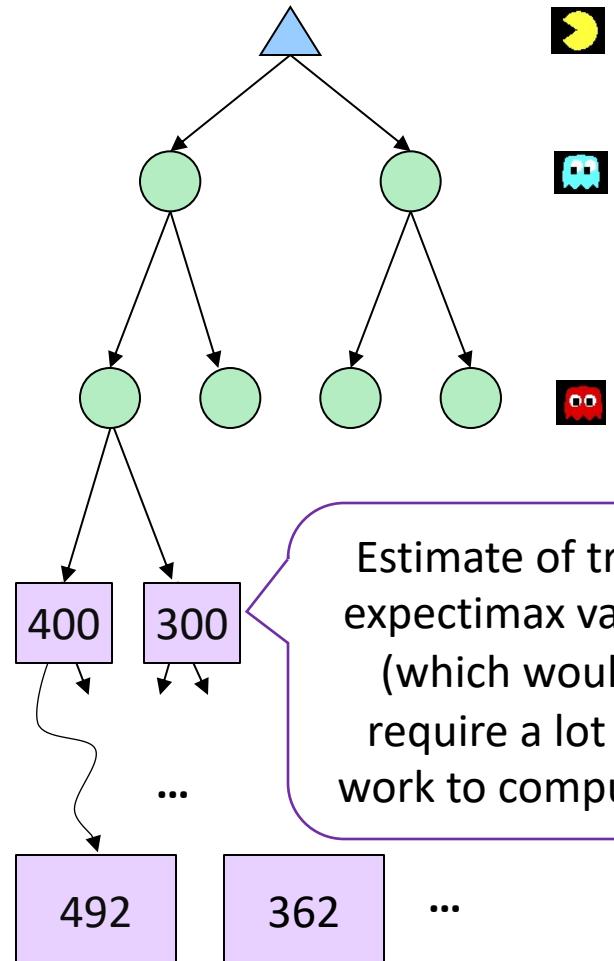
Expectimax Example



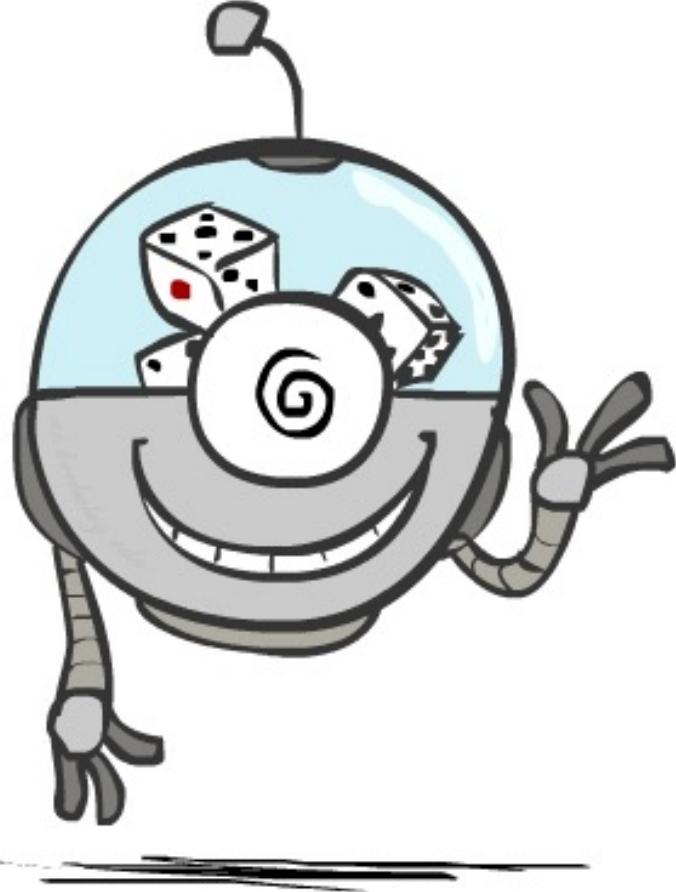
Expectimax Pruning?



Depth-Limited Expectimax

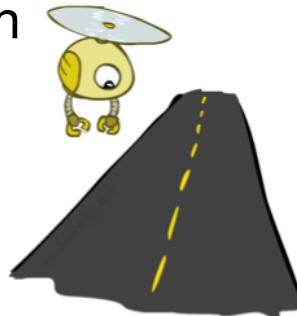


Probabilities



Reminder: Probabilities

- ❑ A **random variable** represents an event whose outcome is unknown
- ❑ A **probability distribution** is an assignment of weights to outcomes
- ❑ Example: Traffic on freeway
 - Random variable: T = whether there's traffic
 - Outcomes: T in {none, light, heavy}
 - Distribution: $P(T=\text{none}) = 0.25$, $P(T=\text{light}) = 0.50$, $P(T=\text{heavy}) = 0.25$
- ❑ Some laws of probability (more later):
 - Probabilities are always non-negative
 - Probabilities over all possible outcomes sum to one
- ❑ As we get more evidence, probabilities may change:
 - $P(T=\text{heavy}) = 0.25$, $P(T=\text{heavy} \mid \text{Hour}=8\text{am}) = 0.60$
 - We'll talk about methods for reasoning and updating probabilities later



0.25



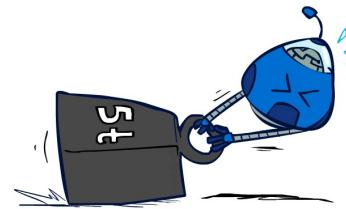
0.50



0.25

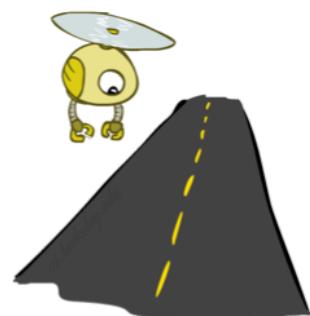
Reminder: Expectations

- The expected value of a function of a random variable is the average, weighted by the probability distribution over outcomes
- Example: How long to get to the airport?



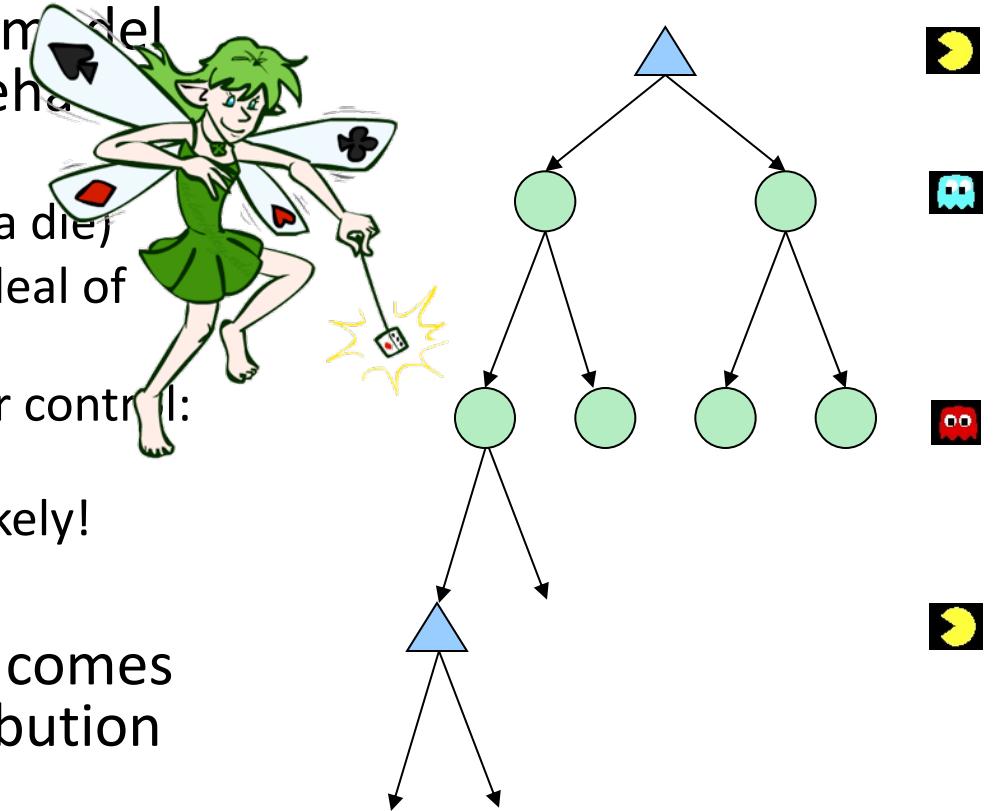
Time:	20 min	x	+	30 min	x	+	60 min	x	35 min
Probability:	0.25			0.50			0.25		

A large blue arrow points from the table to the right.



What Probabilities to Use?

- In expectimax search, we have a probabilistic model of how the opponent (or environment) will behave in any state
 - Model could be a simple uniform distribution (roll a die)
 - Model could be sophisticated and require a great deal of computation
 - We have a chance node for any outcome out of our control: opponent or environment
 - The model might say that adversarial actions are likely!
- For now, assume each chance node magically comes along with probabilities that specify the distribution over its outcomes



Having a probabilistic belief about another agent's action does not mean that the agent is flipping any coins!

Modeling Assumptions



The Dangers of Optimism and Pessimism

Dangerous Optimism

Assuming chance when the world is adversarial

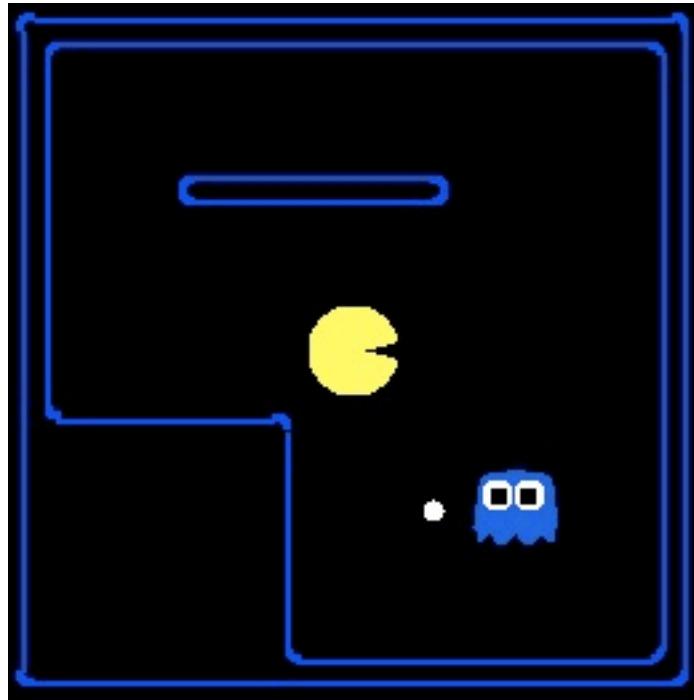


Dangerous Pessimism

Assuming the worst case when it's not likely



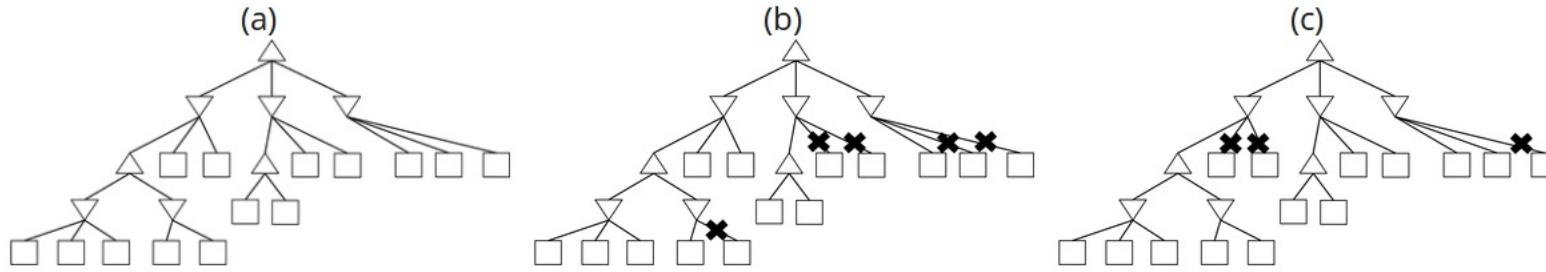
Assumptions vs. Reality



	Adversarial Ghost	Random Ghost
Minimax Pacman	Won 5/5 Avg. Score: 483	Won 5/5 Avg. Score: 493
Expectimax Pacman	Won 1/5 Avg. Score: -303	Won 5/5 Avg. Score: 503

Results from playing 5 games

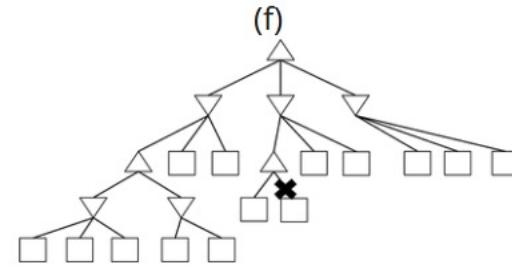
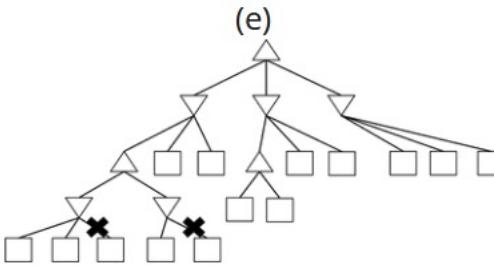
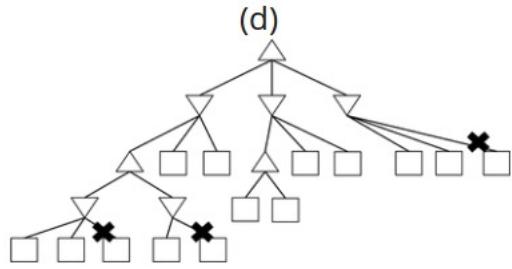
Review Question-Alpha beta pruning



Assume we run α - β pruning, expanding successors from left to right, on a game with tree as shown in figures. Which of the following statements are true?

1. There exists an assignment of utilities to the terminal nodes such that no pruning will be achieved (shown in Figure (a)).
2. There exists an assignment of utilities to the terminal nodes such that the pruning shown in Figure (b) will be achieved.
3. There exists an assignment of utilities to the terminal nodes such that the pruning shown in Figure (c) will be achieved.
4. None of the above.

Review Question-Alpha beta pruning



Assume we run α - β pruning, expanding successors from left to right, on a game with tree as shown in figures. Which of the following statements are true?

1. There exists an assignment of utilities to the terminal nodes such that the pruning shown in Figure (d) will be achieved.
2. There exists an assignment of utilities to the terminal nodes such that the pruning shown in Figure (e) will be achieved.
3. There exists an assignment of utilities to the terminal nodes such that the pruning shown in Figure (f) will be achieved.
4. None