

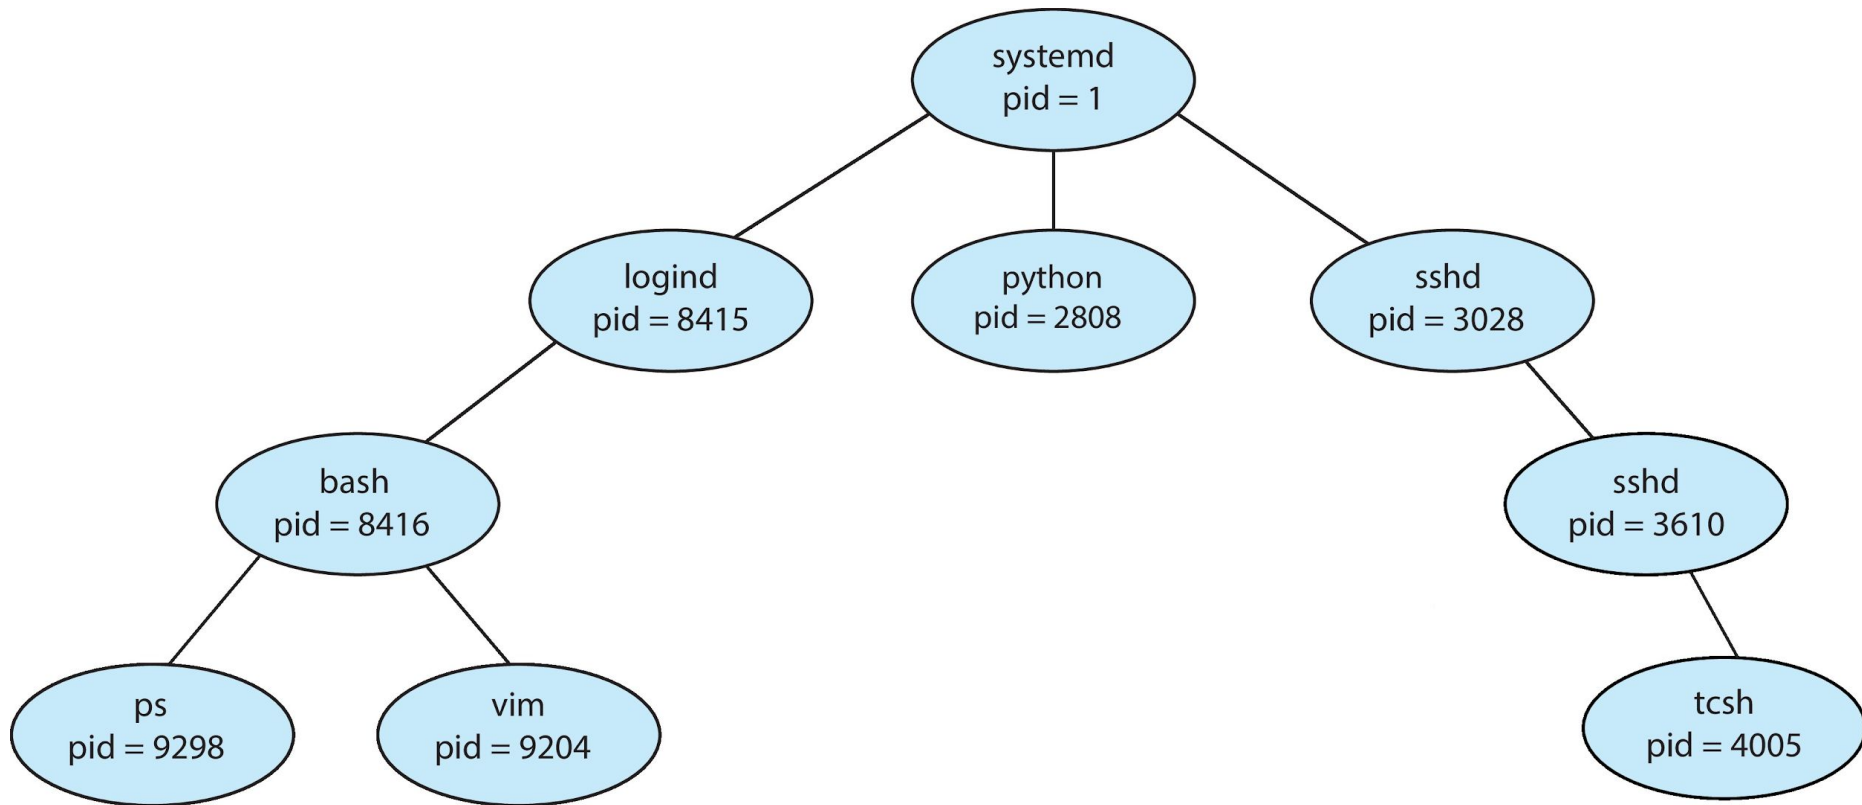
04. Process API

CS 4352 Operating Systems

Process Creation

- Parent process creates child processes, which, in turn create other processes, forming **a tree of processes**
 - Generally, process identified and managed via a process identifier (**PID**)
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and children share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

A Tree of Processes in Linux



Obtaining Process IDs

- `pid_t getpid(void)`
 - Returns PID of current process
- `pid_t getppid(void)`
 - Returns PID of parent process
- The `pid_t` data type is a signed integer type which is capable of representing a process ID
 - In the GNU C Library, this is an `int`

fork() System Call

- In Unix, processes are created using **fork()**
- **int fork(void)**
 - Returns 0 to the child process, child's PID to parent process
 - Child is almost identical to parent:
 - Child get an identical (but separate) copy of the parent's virtual address space
 - Child gets identical copies of the parent's open file descriptors
 - Child has a different PID than the parent
- fork() is interesting (and often confusing) because it is called once but returns twice
 - Again, remember this! It returns the child's PID to the parent, 0 to the child

fork() Example

- Call once, return twice
- Concurrent execution
 - Can't predict execution order of parent and child
- Duplicate but separate address space
 - x has a value of 1 when fork returns in parent and child
 - Subsequent changes to x are independent
- Shared open files
 - stdout is the same in both parent and child

```
#include <stdio.h>
#include <unistd.h>
```

```
int main()
{
    pid_t pid;
    int x = 1;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        printf("child: x = %d\n", ++x);
    }
    else { /* parent process */
        printf("parent: x = %d\n", --x);
    }

    return 0;
}
```

```
zzk@isis:~/courses/CS4352/demo$ ./a.out
parent: x = 0
child: x = 2
```

Duplicating Address Spaces


PC →

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;
    int x = 1;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        printf("child: x = %d\n", ++x);
    }
    else { /* parent process */
        printf("parent: x = %d\n", --x);
    }

    return 0;
}
```

pid = 1234



PC →

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;
    int x = 1;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        printf("child: x = %d\n", ++x);
    }
    else { /* parent process */
        printf("parent: x = %d\n", --x);
    }

    return 0;
}
```

pid = 0

Parent process

Child process

Divergence

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;
    int x = 1;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        printf("child: x = %d\n", ++x);
    }
    else { /* parent process */
        printf("parent: x = %d\n", --x);
    }

    return 0;
}
```

PC →

pid = 1234

Parent process



```
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;
    int x = 1;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        printf("child: x = %d\n", ++x);
    }
    else { /* parent process */
        printf("parent: x = %d\n", --x);
    }

    return 0;
}
```

PC →

pid = 0

Child process

Execution Order


```
zzk@isis:~/courses/CS4352/demo$ ./a.out  
parent: x = 0  
child: x = 2
```

```
zzk@isis:~/courses/CS4352/demo$ ./a.out  
child: x = 2  
parent: x = 0
```

Modeling fork() with Process Graphs

- A process graph is a useful tool for capturing the partial ordering of statements in a concurrent program:
 - Each vertex is the execution of a statement
 - $a \rightarrow b$ means a happens before b
 - Edges can be labeled with current value of variables
 - `printf()` vertices can be labeled with output
 - Each graph begins with a vertex with no in-edges
- Any topological sort of the graph corresponds to a feasible total ordering
 - Total ordering of vertices where all edges point **from left to right**

Process Graph Example




```
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;
    int x = 1;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        printf("child: x = %d\n", ++x);
    }
    else { /* parent process */
        printf("parent: x = %d\n", --x);
    }


    return 0;
}
```

Interpreting Process Graphs


- Original graph:




- Relabeled graph:



Feasible total ordering:

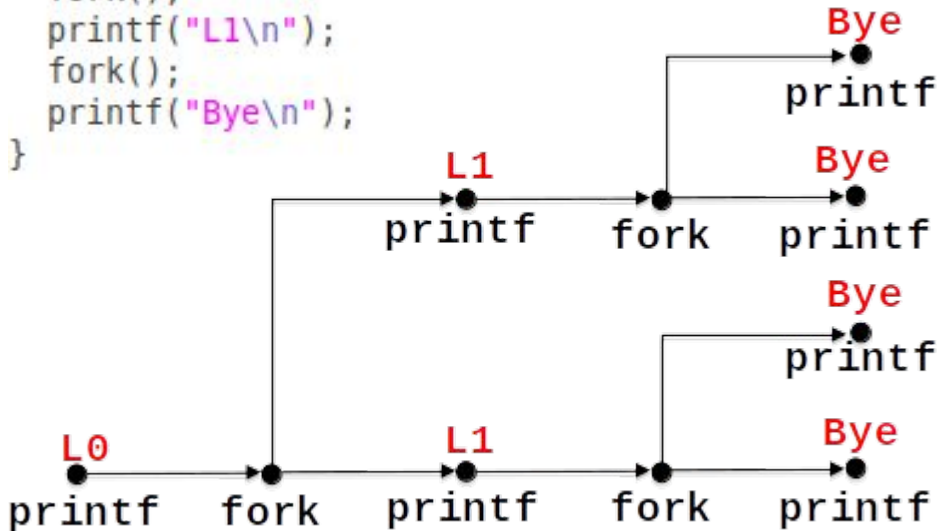


Infeasible total ordering:



fork() Example: Two Consecutive forks

```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```



Feasible output:

L0
L1
Bye
Bye
L1
Bye
Bye

Infeasible output:

L0
Bye
L1
Bye
L1
Bye
Bye

fork() Example: Nested forks in Parent


```
void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

Feasible output:

L0
L1
Bye
Bye
L2
Bye

Infeasible output:

L0
Bye
L1
Bye
Bye
L2



exec() System Call

- Wait a second. How do we actually start a new program?
 - We need to use the exec() system call!
- The exec() system call loads a new program
 - The existing address space is blown away and loaded with the data and instructions of the new program
 - The exec() causes the OS to:
 - Destroy the address space of the calling process
 - Load the new program in memory, creating a new stack and heap
 - Run the new program from its entry point
 - However, things like the PID and file descriptors remain the same

Different Wrapper Flavors

```
int execl(char const *path, char const *arg0, ...);  
int execlp(char const *path, char const *arg0, ..., char const *envp[]);  
int execlp(char const *file, char const *arg0, ...);  
int execv(char const *path, char const *argv[]);  
int execve(char const *path, char const *argv[], char const *envp[]);  
int execvp(char const *file, char const *argv[]);
```

- The base of each is `exec` (execute), followed by one or more letters:
 - `e` – an array of pointers to environment variables is explicitly passed to the new process image
 - `l` – command-line arguments are passed individually (a list) to the function.
 - `p` – uses the `PATH` environment variable to find the file named in the file argument to be executed.
 - `v` – command-line arguments are passed to the function as an array (vector) of pointers.

execve(): Loading and Running Programs

- `int execve(char const *path, char const *argv[], char const *envp[])`
- Loads and runs in the current process:
 - Executable file filename
 - Can be object file or script file beginning with `#!/interpreter` (e.g., `#!/bin/bash`)
 - Argument list `argv`
 - By convention `argv[0] == filename`
 - Environment variable list `envp`
 - “name=value” strings (e.g., `USER=zzk`)
- Overwrites code, data, and stack
 - Retains PID, open files and signal context
- Called once and never returns
 - ...except if there is an error

execve() Example

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    pid_t pid;
    char *myargv[] = {"/bin/ls", "-al", "/", NULL};
    char *environ[] = {NULL};
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        if (execve(myargv[0], myargv, environ) < 0) {
            printf("%s: Command not found.\n", myargv[0]);
            exit(1);
        }
    }
    printf("how many times am I printed?\n");
}
```

Why Are `fork()` and `exec()` Separate?

- Most calls to `fork()` followed by `exec()`
 - Why not combine them together?
 - Not exactly the same as the combination, but `CreateProcess()` on Windows creates a new process and loads a program
- This separation is useful when the child ...
 - Is cooperating with the parent
 - Relies upon the parent's data to accomplish its task
- It has been thought as one part of Unix philosophy

Process Termination

- Process becomes terminated for one of three reasons:
 - Receiving a signal whose default action is to terminate
 - Returning from the `main()` routine
 - Calling the `exit()` function
- `void exit(int status)`
 - Terminates with an exit status of `status`
 - Convention: normal return status is 0, nonzero on error
 - Another way to explicitly set the exit status is to return an integer value from the main routine
 - The wrapper `exit()` function runs the functions registered by `atexit()`
 - `atexit()` function appends a list of routines that should be called when the application terminates normally
 - The `_exit()` after `main()` actually returns control to the operating system
- `exit()` is called **once** but **never** returns

Reaping Child Processes

- When a process terminates for any reason, the kernel does not remove it from the system immediately
 - Instead, the process is kept around in a terminated state until it is reaped by its parent
- When the parent reaps the terminated child, the kernel passes the child's exit status to the parent, and then discards the terminated process, at which point it ceases to exist
 - A terminated process that has not yet been reaped is called a **zombie**
 - Living corpse, half alive and half dead

How to reap the terminated child processes?

wait(): Synchronizing with Children

- Parent can reap a child by calling the wait() function
- `int wait(int *child_status)`
 - Suspends current process until one of its children terminates
 - Return value is the PID of the child process that terminated
 - If `child_status != NULL`, then the integer it points to will be set to a value that indicates reason the child terminated and the exit status:
 - The `<sys/wait.h>` include file defines several macros for interpreting the status argument
 - E.g., `WIFEXITED(status)`: returns true if the child terminated normally via a call to `exit()` or a return
 - `WEXITSTATUS(status)`: returns the exit status of a normally terminated child
 - This status is only defined if `WIFEXITED` returned true

waitpid(): Waiting for a Specific Process

- `pid_t waitpid(pid_t pid, int &status, int options)`
 - Suspends current process until specific process terminates
 - Many options that can be found by typing “man waitpid”

```
for (i = 0; i < N; ++i) {
    if ((pid[i] = fork()) == 0)
        exit(100 + i); /* Child */
}

for (i = 0; i < N; ++i) {
    pid_t wpid = waitpid(pid[i], &child_status, 0);
    if (WIFEXITED(child_status))
        printf("Child %d terminated with exit status %d\n",
            wpid, WEXITSTATUS(child_status));
    else
        printf("Child %d terminate abnormally\n", wpid);
}
```

What If ...

- What happens if a parent process terminates before its child?
 - Sadly :-(it becomes an orphan
 - An orphan process is adopted by the init process once its parent process dies
- The init process has a PID of 1 and is created by the kernel during system initialization
 - The init process will periodically execute `wait()` to reap processes
- Long-running programs such as shells or servers should always reap their zombie children
 - Again, even though zombies are not running, they still consume system memory resources

Zombie Example

```
int main()
{
    if (fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n", getpid());
        exit(0);
    } else {
        printf("Running Parent, PID = %d\n", getpid());
        while (1)
            ; /* Infinite loop */
    }
}
```

```
zzk@isis:~/courses/CS4352/demo$ ./a.out
Running Parent, PID = 18053
Terminating Child, PID = 18054
^Z
[1]+  Stopped                  ./a.out
zzk@isis:~/courses/CS4352/demo$ ps
  PID TTY          TIME CMD
 18028 pts/0        00:00:00 bash
 18053 pts/0        00:00:31 a.out
 18054 pts/0        00:00:00 a.out <defunct>
 18148 pts/0        00:00:00 ps
zzk@isis:~/courses/CS4352/demo$ fg
./a.out
^C
zzk@isis:~/courses/CS4352/demo$ ps
  PID TTY          TIME CMD
 18028 pts/0        00:00:00 bash
 18250 pts/0        00:00:00 ps
zzk@isis:~/courses/CS4352/demo$
```

Homework

- Start reading Chapter 5

Next Lecture

- We learn process scheduling



Credit: <https://xkcd.com/734/>